Explainable AI For Epileptic Seizure Prediction Using EEG Signals: Enhancing Transparency and Clinical Trust

Deepiga S	(21Z214)
Maddu Hemali Sai Pravallika	(21Z225)
Monaleka M	(21Z231)
Santhoshi R	(21Z251)
Vasudha R B	(21Z267)

Dissertation submitted in partial fulfillment of the requirements for the degree of

BACHELOR OF ENGINEERING

Branch: COMPUTER SCIENCE AND ENGINEERING

April 2025

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

PSG COLLEGE OF TECHNOLOGY

(Autonomous Institution)

COIMBATORE – 641 004

CERTIFICATE

Certified that this report titled "**EXPLAINABLE AI FOR EPILEPTIC SEIZURE PREDICTION USING EEG SIGNALS: ENHANCING TRANSPARENCY AND TRUST**", for the Project work II (19Z820) is a bonafide work of **Deepiga S (21Z214), Maddu Hemali Sai Pravallika (21Z225), Monaleka M (21Z231), Santhoshi R (21Z251), Vasudha R B (21Z267)** and they have carried out the work under my supervision for the partial fulfillment of the requirements for the award of the degree of Bachelor of Engineering in Computer Science and Engineering. Certified further that to the best of my knowledge and belief, the work reported herein does not form part of any other thesis or dissertation on the basis of which a degree or an award was conferred on an earlier occasion.

Deepiga S (21Z214)

Maddu Hemali Sai Pravallika (21Z225)

Monaleka M (21Z231)

Santhoshi R (21Z251)

Vasudha R B (21Z267)

Dr. Suriya S
Associate Professor
Department of Computer Science and Engineering
PSG College of Technology

Place: Coimbatore

Date:

COUNTERSIGNED

HEAD
Department of Computer Science and Engineering
PSG College of Technology
Coimbatore – 641 004

PSG COLLEGE OF TECHNOLOGY

(Autonomous Institution)

COIMBATORE – 641 004

EXPLAINABLE AI FOR EPILEPTIC SEIZURE PREDICTION USING EEG SIGNALS: ENHANCING TRANSPARENCY AND TRUST

Bona fide record of work done by

Deepiga S	(21Z214)
Maddu Hemali Sai Pravallika	(21Z225)
Monaleka M	(21Z231)
Santhoshi R	(21Z251)
Vasudha R B	(21Z267)

Dissertation submitted in partial fulfillment of the requirements for the degree of

BACHELOR OF ENGINEERING

Branch: COMPUTER SCIENCE AND ENGINEERING

of Anna University

April 2025

...……	……..…….
Dr. Suriya S	**Dr. G. Sudha Sadasivam**
Faculty guide	Head of the Department

Certified that the candidate was examined in the viva-voce examination held on ……

……..	……..
(Internal Examiner)	(External Examiner)

CONTENTS

CHAPTER **Page No.**

SYNOPSIS

Epileptic seizures refer to sudden bursts of abnormal electrical activity in the brain, often detected through EEG signals. The seizure detection and prediction are crucial for timely intervention and suitable treatment. This project investigates the integration of XAI techniques into a deep-learning EEG signal classification system for predicting epileptic seizures. The intention is to classify brain signals in three modes: ictal, interictal, and preictal, through ResNet-18 architecture with enhanced interpretability of the model.

Five leading XAI methods SHAP, LIME, Grad-CAM, Integrated Gradients, and DeepLIFT were used to provide visual and quantitative insights into the predictions of the model. The performance of these techniques is assessed using various metrics, such as fidelity, localization, and stability, to determine the meaningfulness of the insights provided by the deep learning model into its own decision-making process.

The dataset contained pre-recorded EEG signals classified into three seizure stages, and each sample was stored in .mat format with fixed duration. Hence no real-time signal acquisition and hardware integration exist, and data processing here is offline. A web application based on Streamlit was developed to be the frontend, which allows clinicians and users to visualize the model predictions and explanations interactively.

The project aims to bring about a reconciliatory step between accuracy and interpretability of medical AI applications, thereby assuring that the system remains high-performing and interpretable in a safety-critical domain like healthcare.

LIST OF FIGURES

LIST OF TABLES

LIST OF ABBREVIATIONS

ABBREVIATION	DEFINITION
1D	One Dimensional
2D	Two Dimensional
3D	Three Dimensional
AI	Artificial Intelligence
BCI	Brain-Computer Interface
BiLSTM	Bidirectional Long Short-Term Memory
CNN	Convolutional Neural Network
CPU	Central Processing Unit
CSS	Cascading Style Sheet
DeepLIFT	Deep Learning Important FeaTures
DICOM	Digital Imaging and Communications in Medicine
EEG	Electroencephalogram
FC layer	Fully Connected Layer
Grad-CAM	Gradient-weighted Class Activation Mapping
GPU	Graphics Processing Unit
HFD	Higuchi Fractal Dimension
HIPAA	Health Insurance Portability and Accountability Act
HTML	HyperText Markup Language
IG	Integrated Gradients
LIME	Local Interpretable Model-Agnostic Explanations
LRP	Layer-wise Relevance Propagation
L2 regularization	L2 Norm Regularization (Ridge Regression)
PIL	Python Imaging Library
RAM	Random Access Memory

ReLU	Rectified Linear Unit
ResNet	Residual Network
RGB 3	Red Green Blue (3 channels)
SEC	Seizure Evolution Curve *(context-specific)*
SHAP	SHapley Additive exPlanations
SLIC	Simple Linear Iterative Clustering
STFT	Short-Time Fourier Transform
TUSZ	Temple University Seizure Corpus
VARSHAP	Variance-based SHAP
VAE-LIME	Variational Autoencoder + LIME
XAI	Explainable Artificial Intelligence

ACKNOWLEDGEMENT

We would like to thank the management of **PSG College of Technology** for providing us the infrastructure and helping us envision new ideas. We would also like to express our heartfelt gratitude to our Principal **Dr.K.Prakasan** for bestowing us with this valuable opportunity to work in our area of interest.

We also extend our sincere thanks to our Head of the Department **Dr.G.Sudha Sadasivam,** Department of Computer Science and Engineering for constantly supporting us to develop our ideas and present our project to the faculty committee as a part of our partial fulfilment of the requirements leading to the awarding of B.E. degree.

We express our sincere gratitude to our Program Coordinator **Dr. Arul Anand N**, Professor, Department of Computer Science and Engineering for his unwavering support and guidance. His expertise and encouragement have been indispensable to the success of the project.

We take immense pleasure in thanking our project guide, **Dr. Suriya S,** Associate Professor, Department of Computer Science and Engineering, for being a pillar of support, without whose guidance, unparalleled cooperation and constructive criticisms, this project wouldn't have been fruitful.

We would also like to thank our tutor, **Mr. A C Ramesh,** Assistant Professor (Sl. Gr.), Department of Computer Science and Engineering, for their encouragement and constant vigilance during the course of this project.

A sincere thanks to all the **Panel Members** for reviewing out project and all the **Faculty Members and Staffs** of the department for offering us the required support during course of the project.

CHAPTER 1

INTRODUCTION

This chapter gives a brief overview of the entire system that integrates XAI techniques into an EEG-based Epileptic Seizure Prediction System that classifies a signal to be under any of the three stages: ictal, interictal or preictal. This integration ensures that the system is reliable and trustworthy and improves decision-making in epilepsy management.

1.1 PROBLEM STATEMENT

Epileptic seizures refer to sudden bursts of abnormal electrical activity in the brain, often detected through EEG signals. The seizure detection and prediction are crucial for timely intervention and suitable treatment. The black-box models like deep learning models have high predictive power and are less interpretable. There exists a trade-off between Explainability and Predictive power for the machine learning and deep learning models. Also doctors and clinicians who use the medical applications are not aware of deep learning models and their working. Hence it is necessary to inform the doctors and clinicians by providing explanations for why a model predicts a particular brain state as ictal, interictal or preictal.

The project aims to integrate five XAI techniques - SHAP, LIME, Grad-CAM, Integrated Gradients, and DeepLIFT into a Resnet-18 based EEG Seizure Prediction model to enhance model transparency. A comparative analysis of their effectiveness will be conducted using XAI evaluation metrics such as Fidelity, Localization, and Stability.

1.2 SCOPE OF THE PROJECT

XAI matters because it imposes trust, performance and accountability to the AI-based systems. Safety-critical applications are systems in which a failure can lead to disastrous

events such as harm, injury or may even lead to death requiring rigorous design and testing. Hence adhering to the safety standards ensure reliability and reduce risk. Some of the common safety-critical applications are automotives, devices used in aerospace, finance and healthcare. The medical device that will be used for detecting epileptic seizures will also fall under this category.

The dataset[30] used for this project contains raw EEG signals categorized into ictal, interictal and preictal stored in three separate folders each consisting of fifty .mat files where each signal is recorded for a duration of 5.12 seconds. The ResNet-18 model is used for model training and classification. The XAI techniques used are SHAP, LIME, Grad-CAM, Integrated Gradients, and DeepLIFT. The XAI techniques are evaluated using the metrics fidelity, localization and stability. Based on the metrics obtained the techniques are compared and analysed. Finally, a frontend is implemented for the system as Streamlit-based web application. In this project there are no real-time signals or data involved. Everything is pre-recorded signals. No hardware devices are used in the project.

CHAPTER 2

LITERATURE SURVEY

The literature survey explores the evolution of XAI, the techniques - SHAP, LIME, Grad-CAM, Integrated Gradients, DeepLIFT and the application of XAI for epileptic seizure detection. The survey also touched upon the other domains in which XAI techniques are applicable (e.g safety critical applications).

2.1 LITERATURE

The core concepts of XAI were studied in detail where [1] presented the methods of explainability such as visualization, knowledge extraction, surrogate models, decomposition and covered the reasons of XAI as explain to justify, explain to control, explain to improve and explain to discover. The general evaluation metrics used to assess the XAI are given in [2] where the metrics are broadly categorized into two main categories as human-centric metrics that measure user -satisfaction and model-centric metrics that quantitatively measure interpretability of XAI. The taxonomy of XAI is presented in [3,16] where the XAI techniques can be classified based on stage as ante-hoc techniques or post-hoc techniques, based on approach as model-specific techniques or model-agonistic techniques and based on methodology as perturbation-based techniques or gradient-based techniques. The challenges faced during the research phase, development phase and the deployment face while designing an XAI system and their corresponding research directions have been clearly explained in [5].

The foundations of SHAP have been laid in [6] where SHAP has been introduced as a unified framework for interpreting model predictions, combining ideas from game theory with machine learning interpretability. The application of SHAP values to a system that identifies the optimal functional forms of covariates in pharmacokinetic models has been explored in [7]. Similarly, the SHAP analysis for enhancing transparency in drug development using machine learning

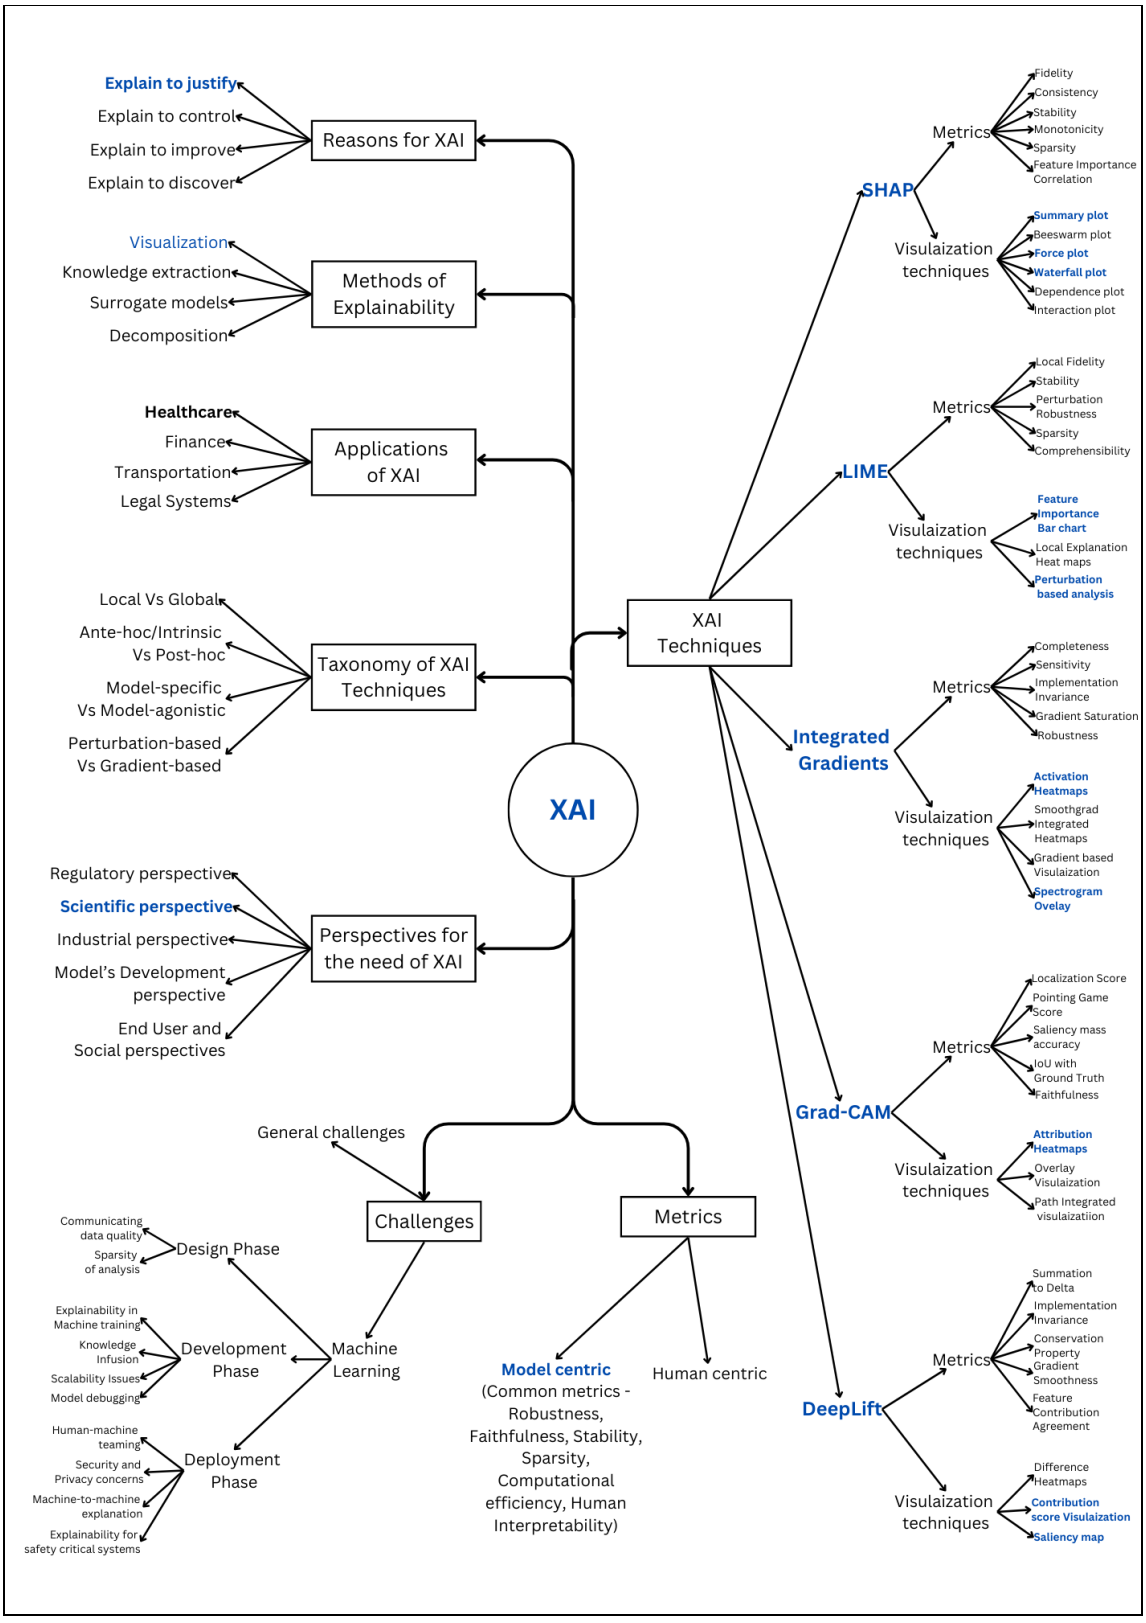

Fig 2.1 Mind map of Literature Survey

models have been carried out in [8] by visualizing SHAP values of the system globally using plots like Bar plot, Beeswarm plot, scatter plot, waterfall plot and time-series plot. The paper also touched upon the limitations of SHAP analysis such as feature dependence, interpretation of predictions and its computational complexity. [9] explains how the use of SHAP analysis for Time-series data by proposing two methodologies namely VARSHAP and Time-Consistent SHAP, whereas [10] explains on adapting LIME to time-series classification tasks by proposing frameworks like Nearest Neighbour segmentation, Realistic Background Perturbation and Dynamic Time Warping.

[11] introduces VAE-LIME which extends LIME by improving the fidelity of LIME's local explanations for black-box models by generating realistic perturbations using a VAE. SHAP and LIME techniques are compared for machine learning models particularly in medical imaging - lung X-ray classification for pneumonia detection and the advantages and disadvantages of SHAP and LIME were pointed out in [12]. They concluded that SHAP is more suitable for comprehensive insights and high-stakes applications like healthcare due to its robustness whereas LIME is better for quick, local explanations with minimal computational requirements.

[13] categorizes XAI techniques by highlighting their strengths and limitations and lists the challenges that hinder the development of responsible AI. The review of XAI research in healthcare decision-making is presented in [14]. [15] surveys existing XAI methods and introduces the concept of "Explanation Engineering" - a systematic discipline aimed at designing, implementing and evaluating explainability solutions for real-world AI systems.
The application of XAI methods to address dataset shifts in EEG-based systems using Integrated Gradients and DeepLIFT in [17] and using SHAP and LIME in [18]. A binary classifier for indicating the presence or absence of seizure built using Random Forest classifier is interpreted using LIME in [19] to identify the key features that contributed towards the classification. Research work in [20] helped to identify that beta-band features are critical for moderate/high correlations in TUSZ [31] and states that correlation analysis like Spearman's rank correlation coefficient doesn't guarantee high model performance.

In [21], a LightGBM-based model was developed for classifying normal EEG, focal epilepsy and generalized epilepsy by incorporating XAI techniques like SHAP and LIME. The application of Grad-CAM for EEG seizure detection based on connectivity features is given in

[22]. It also introduced the technique Grad-CAM++ that uses gradients of the output class with respect to the activations. Development of EEG-based epileptic seizure detection system integrating SHAP is done in [23] SHAP ensures robust seizure detection and reduces computational resources. SHAP was found to be suitable for real-time and mobile applications. The same study is done in [24] which highlights key features like Higuchi Fractal Dimension (HFD) are identified as critical contributors, aiding clinical decision-making. The framework discussed uses feature engineering, stacking ensemble classifiers (SEC) and secure blockchain integration which achieves a 2% performance improvement on benchmark datasets.

Enhancing interpretability and reliability of machine learning models for Epileptic Seizure Detection in [25] discusses XAI techniques like SHAP, LIME, Grad-CAM. Study involved in [26] detects epileptic seizures using EEG connectivity features and deep learning. It applies CNN, BiLSTM, and attention mechanisms on 20-second windows. Explainability is achieved by analyzing network weights to identify key features during seizure and non-seizure states. Cross-patient analysis highlights variability, aiding clinical understanding. Paper [27] focuses on improving BCI systems by tackling dataset shifts using explainable AI (XAI) methods. Techniques like LRP, Integrated Gradients, and DeepLIFT outperformed Saliency and Guided Backpropagation in identifying key EEG features. While effective at highlighting features in individual samples, these methods struggled to generalize across sessions. The study emphasizes XAI's promise and its current limitations in BCI applications. Paper [28] introduces XAI4EEG to enhance explainability in EEG-based seizure detection using 1D and 3D CNNs. It integrates domain knowledge and uses SHAP to highlight key spectral, spatial, and temporal features. The goal is to support clinical decision-making with interpretable insights. The project discussed in [29] aims to detect seizures from EEG signals using a CNN-based deep learning model. After preprocessing, the model is trained with binary cross-entropy loss and evaluated using classification metrics. Grad-CAM heatmaps and convolutional layer weights help visualize important EEG features, enhancing interpretability. The method demonstrates potential for real-time seizure monitoring and clinical applications. The entire literature survey is presented comprehensively as a mind-map in Fig 2.1.

2.2 CONCLUSION OF LITERATURE SURVEY

The survey of literature points to the rising importance of XAI approaches in the field of explainable machines. SHAP is the perfect method for global explanations. LIME makes quick and local approximations. DeepLIFT performs best when models are more structured and interpretable, discerning the best approach that is suited for mixed and time-series data is not always the best among other approaches. The future work of developing hybrid XAI methods is the way forward for the field, as it allows researchers evaluation metrics.

2.2.1 Factors considered:

- For the model type Resnet-18 with STFT feature extraction technique is used that is well suited for image-like EEG spectrogram inputs
- Since the ResNet-18 model is a black box model, model-specific and model-agonistic XAI techniques were explored to enhance the system's transparency
- Spectrograms preserve both temporal and frequency information from EEG signals. This motivated the use of gradient-based methods like Grad-CAM and Integrated Gradients which align well with CNN activations on image-like inputs.
- A diverse set of XAI techniques like SHAP, LIME which are perturbation based and Grad-CAM, Integrated Gradients, DeepLIFT which are gradient-based. Hence, this helps in improving both global(feature-level) and local(sample-level).
- The methods like Grad-CAM were chosen for their ability to generate heatmaps that can highlight frequency bands and time segments of interest in seizure prediction/detection. More computationally intensive methods like SHAP were also included for in-depth analysis.
- Fidelity, Localization and Stability were used to quantitatively evaluate the performance and reliability of the generated explanations.

2.2.2 Factors not considered

- Models like decision-tree, random forest and other rule-based classifiers were not used because they lack the representational power needed for EEG spectrogram classification as the focus was on high-accuracy deep learning models and post-hoc explainability.

- User studies or surveys were not conducted to evaluate the quality of explanations due to project constraints and lack of access to clinical professionals. Instead model-centric metrcs were used for objective assessment.
- Explanations were not directly extracted from the raw EEG time-series data as the CNN models operate on image-like spectrogram data.
- Raw EEG preprocessing is out-of-scope due pre-processing choices.
- While methods like BiLSTM or attention-based models provide time-aware insights, they were not used here to maintain simplicity and compatibility with image-based XAI methods.

2.2.3 Justification for Decisions

- Choice of Resnet-18: ResNet-18 offers good depth with residual connections, ensuring stable training on EEG spectrograms without overfitting. Its architecture is widely used and well-supported for the chosen XAI methods.
- Selection of Grad-CAM, IG, DeepLIFT: These gradient-based methods work directly with CNNs and provide spatial insight into which regions (time-frequency zones) were important.
- Inclusion of SHAP and LIME: These model-agnostic methods complement gradient-based techniques by offering feature-level perturbation insights. SHAP was selected for high interpretability despite its computational cost, while LIME was useful for quick, local validations.
- Multimethod Explainability: Combining multiple XAI methods helps cross-validate explanations and reduce bias or over-reliance on a single method. Each method provided unique interpretive angles — SHAP for importance, Grad-CAM for localization, IG/DeepLIFT for sensitivity.
- Use of Quantitative Metrics: Fidelity, Localization, and Stability were chosen based on literature [2,13,15] to objectively assess explanation quality. These metrics help bridge the gap between model interpretability and practical utility in healthcare AI.

CHAPTER 3

SYSTEM REQUIREMENTS

This chapter outlines the hardware and software requirements for developing and running advanced deep learning applications and data processing. Ensuring the availability of the hardware and software resources mentioned in this section enables smooth development of the system.

3.1 HARDWARE REQUIREMENTS

The hardware resources required for the development of the system are mentioned in detail in Table 3.1.

Table 3.1 Hardware Requirements

Category	Requirements Description
CPU	Quad-core processor (e.g., Intel Core i7 or AMD Ryzen 7).
GPU(optional)	NVIDIA GTX 1080 Ti, RTX 2070/3070, or higher.
RAM	8 GB minimum (16 GB recommended).
Storage	500 GB to 1 TB SSD.

3.2 SOFTWARE REQUIREMENTS

The software resources, packages, libraries and frameworks required for the development of the system are mentioned in detail in Table 3.2.

Table 3.2 Software Requirements

Category	Requirements Description
Operating System	Windows 10/11 (64-bit)
Programming language	Python 3.8+
Core Libraries and Frameworks	• Pytorch - For implementing and training Resnet-18 deep learning model. • Torchvision - To access predefined Resnet-18 architecture. • NumPy - For numerical operations and data manipulation. • Pandas - For structured data handling. • Matplotlib/Seaborn - For visualization of spectrograms and XAI outputs. • Scikit-learn - For preprocessing and evaluation metrics. • SHAP, LIME, Captum - For XAI Integration.
Web Application	Streamlit - Frontend framework to build and deploy interactive web interface.
IDE	Visual Studio Code / Jupyter Notebook. Use conda virtual environment for dependency isolation.
Version control(optional)	Git

CHAPTER 4

SYSTEM DESIGN AND IMPLEMENTATION

4.1 WORKFLOW OF PHASE 1 AND PHASE 2

The two-phase structure of the EEG-based multistage classification system for epileptic seizure prediction is illustrated in the workflow summary diagram in Fig 4.1. The system has been developed over two consecutive semesters with Phase 1 focusing on EEG signal processing and predicting a seizure stage as ictal, interictal, or preictal. Phase 2 focuses on integrating XAI techniques into the Phase 1 prediction module enhancing model interpretability.

The .mat files contain raw EEG signals and it is time-series data. The raw EEG signals are usually noisy and high dimensional. So, the EEG signals must undergo a data preprocessing step to filter unwanted noise and retain only meaningful brain activity patterns in the time domain. Then data in .mat files are converted to spectrogram images using the feature extraction technique STFT that transform filtered EEG signals from time domain to the time-frequency domain. This conversion of raw EEG signals into image data helps improve the Resnet-18 model training as CNN-based models train well on image data rather than signal data. This is followed by the model training phase where Resnet-18 is employed to leverage its residual connections to efficiently learn spatial and temporal features. The trained model is then used for predicting the epileptic seizure stage as ictal representing the active seizure phase, interictal referring to normal brain activity between seizures or preictal indicating the period before a seizure. This output prediction serves as the base for further extending the research into integrating interpretability and explainability of the model for EEG-Based Multistage Classification for Epileptic Seizure Prediction.

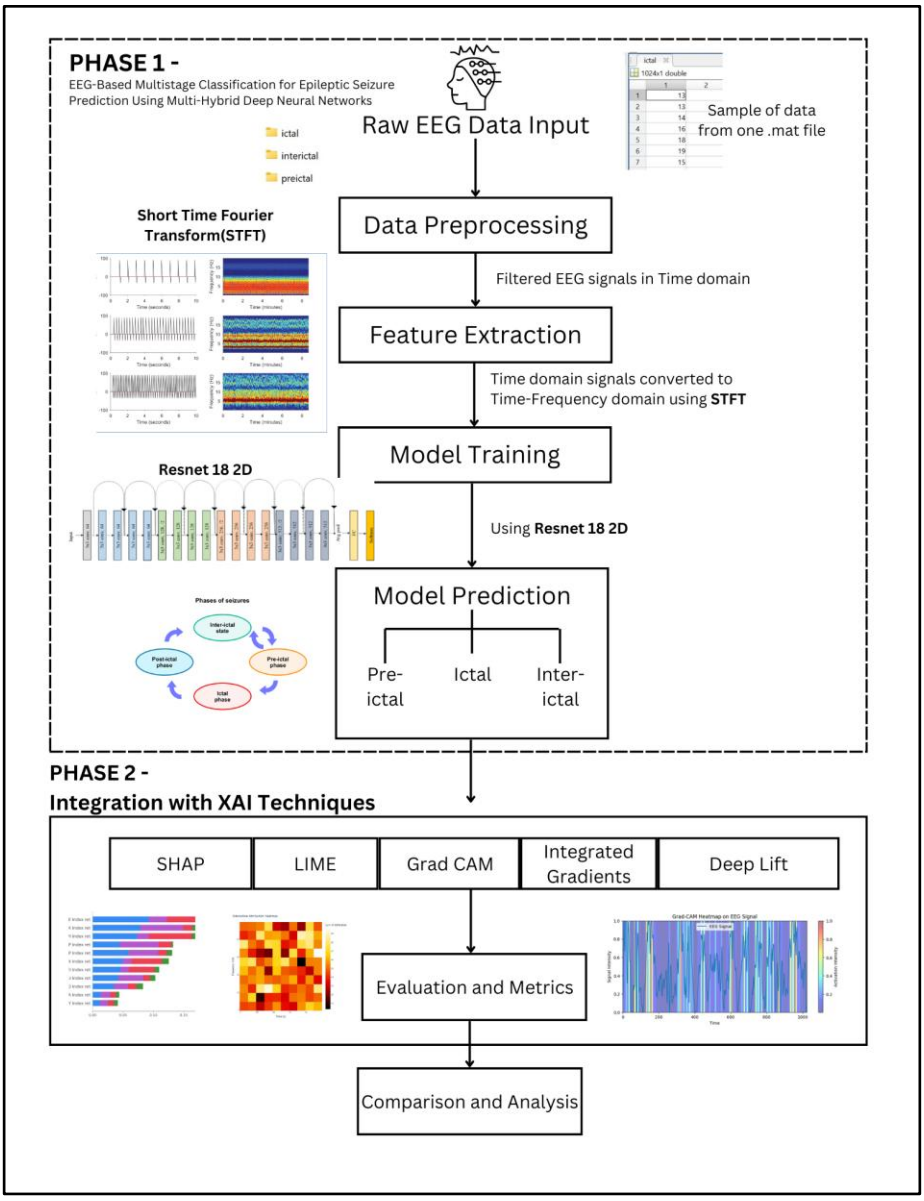

Fig 4.1 Two-Phase Development Workflow of EEG Seizure Prediction with XAI Integration

Building upon Phase 1, Phase 2 integrates XAI techniques to bring transparency in model prediction and decision making. The techniques used in the project are SHAP, LIME, Grad-CAM, Integrated Gradients, and DeepLift. These techniques are used to generate visual explanations for identifying the regions in the spectrogram that contributed towards the prediction. This in turn allows the researchers and clinicians to better understand and trust the model's outputs. The performance of the techniques is evaluated using metrics for XAI to ensure the techniques work as intended and further refine the system.

This structured two-phase approach for epileptic seizure prediction is not only effective in detecting seizure states but also provides the necessary transparency to be applied in clinical settings, making it a valuable tool for seizure prediction and monitoring.

4.2 SYSTEM ARCHITECTURE OF PHASE 2

This section focuses on the system architecture for Phase 2 as shown in Fig 4.2. Before starting with the model training the dataset is split into parts as Training set (80%), Validation set (10%) and Testing set (10%). This ensures optimal balance between learning, tuning and final model evaluation.

Following this the data is normalized between the values 0 and 1 to reduce any biases caused by variations in EEG signal amplitude and helps the model converge better. The model is trained by using the spectrograms generated for the EEG signals using STFT feature extraction. For deep learning models like Resnet-18 a large amount of data is required for training. But since the original dataset [30] contains only 150 .mat files (50 under each category ictal, interictal. preictal). So, the training set is augmented with data augmentation techniques like Jitter and Time-wrap to populate the training data by generating synthetic data. This ensures that the model doesn't overfit. This can be manually checked during the model training process by verifying if the final training accuracy and validation accuracy doesn't vary much and should be nearly equal. If not, the model should be fine-tuned by modifying the hyperparameters and trained again until the benchmark is achieved.

After training is completed save the best version of the trained model for importing the model whenever necessary in the future. Then the XAI techniques SHAP, LIME, Grad-CAM, Integrated Gradients, and DeepLift are each individually run for the Resnet-18 model. The techniques generate visual explanations like summary plot, bar plot, waterfall plot, dependence plot, local explanation heatmap and gradient based visualizations. The explanations generated by the techniques are evaluated using metrics like Fidelity, Localization and Stability to identify the areas for improvement and verifying if the system is trustworthy. Finally, the Comparison and Analysis phase then provides a deeper examination of the model's effectiveness and reliability, ensuring that the developed system meets the highest standards for clinical and real-world applications. By combining advanced deep learning techniques with XAI-driven interpretability, this architecture ensures a robust and transparent framework for epileptic seizure prediction, making it a valuable tool for both medical research and real-time clinical diagnosis.

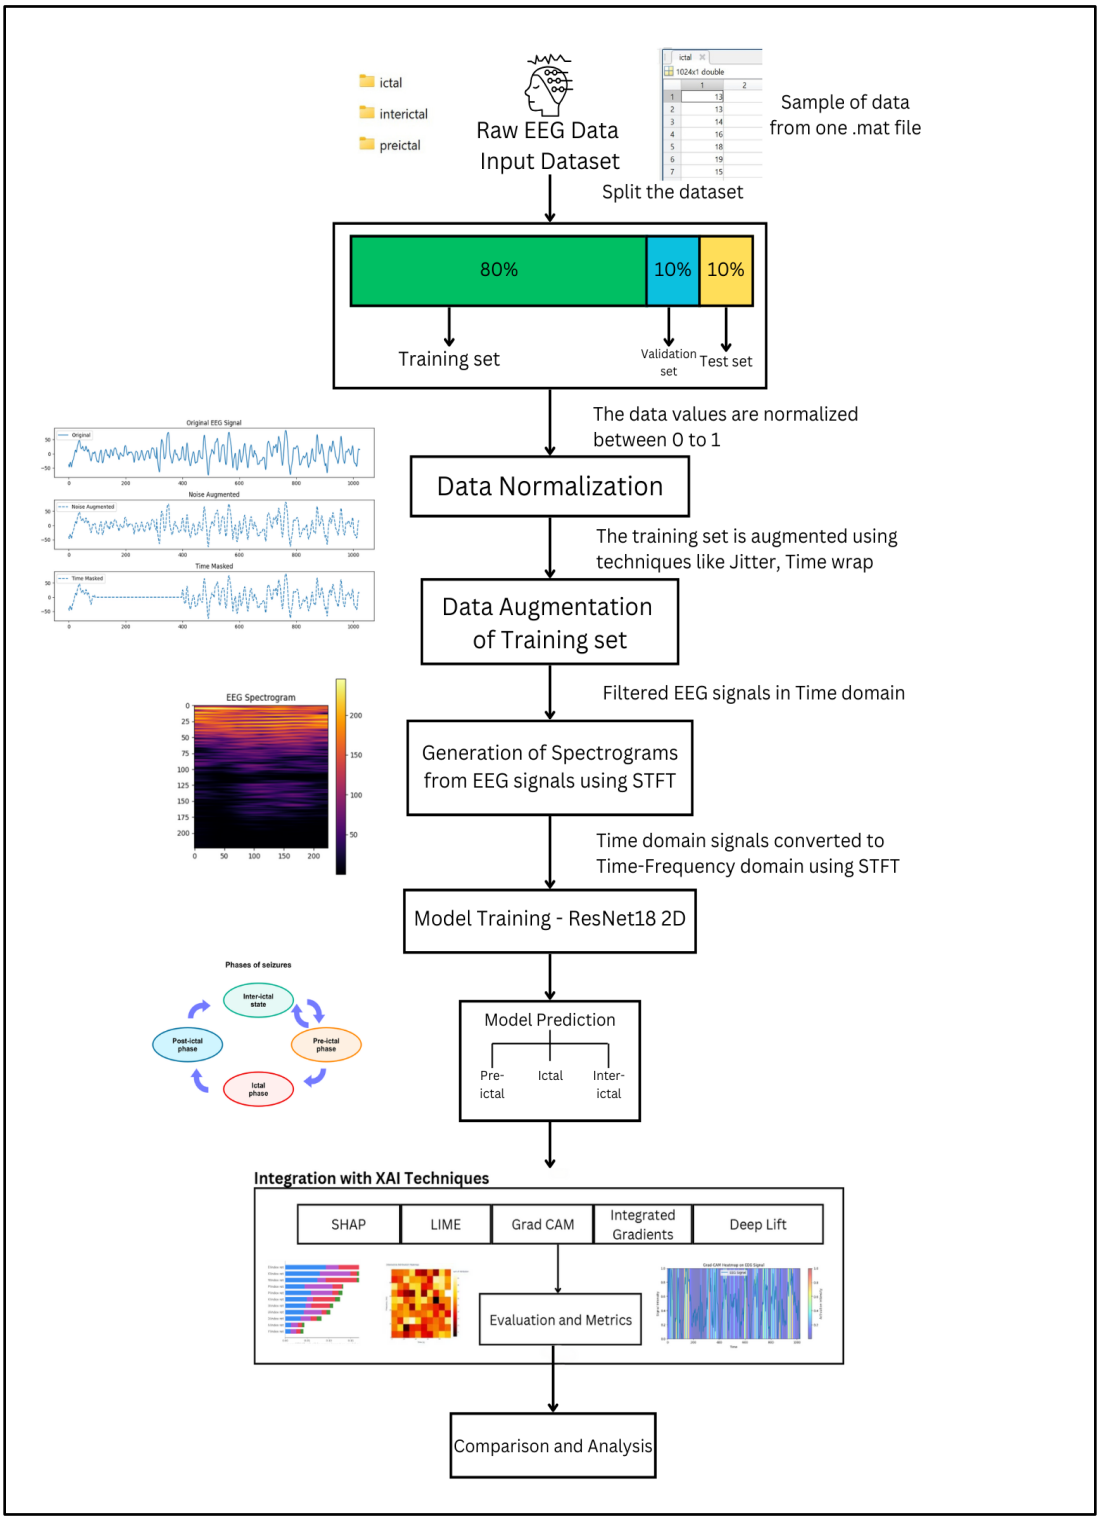

Fig 4.2 Phase 2 Architecture

4.3 SYSTEM SEGMENTS

This section covers the system segments undertaken in the project. The Data Preprocessing and Feature Extraction, Model Training and XAI Integration are the major system segments. The segments are explained in detail in the section.

4.3.1 Data Pre-processing and Feature Extraction

Initially the EEG signals are loaded from the .mat file and normalized. The signal is normalized as given in equation 4.1:

$$Normalized\ signal\ = \frac{x - min(x)}{max(x) - min(x) + \epsilon} \quad (4.1)$$

where $\epsilon\ =\ 1e^{-8}$ is used to avoid division by zero. This scales the data to the $[0,1]$ range, which helps neural networks converge faster.

This is followed by applying data augmentation techniques like jitter, time-warp, scaling, permutation, flip-signal, and add sign wave as given in Table 4.1. Applying these techniques help improve generalization and prevent overfitting.

Table 4.1 Data Augmentation Techniques with Description

Augmentation Technique	Description
Jitter	Jitter adds Gaussian noise$(x\ +\ \sigma\ .N(0,1))$. This is used to simulate sensor noises or artifacts
Time-warp	Time-warp stretches or compresses signals using interpolation. This technique introduces time-based distortion for variability
Scaling	This technique simulates variability in signal strength by randomly scaling amplitude as $x.r$ where $r\ \epsilon\ (0.8, 1.2)$.
Permutation	Permutation introduces temporal disorder while preserving global patterns and splits segments by randomly rearranging them.
Flipped signal	This technique multiplies the signal by -1 and simulates polarity inversion.
Add sine wave	This technique mimics rhythmic disturbances or oscillations by injecting sinusoidal wave $A\ sin(2\pi ft)$.

Each augmentation technique introduces a new data instance in the training set effectively expanding the dataset six times larger than the original dataset.

The EEG signals are then converted into 2-D spectrogram and normalized again as given in equation 4.2:

$$Spectrogram = \frac{S - min(S)}{max(S) + \epsilon} \quad (4.2)$$

where $\epsilon = 1e^{-8}$ helps ensure stability. This conversion ensures the spectrograms generated are suitable for CNN-based models like Resnet-18.

4.3.2 Model training

The Data Preprocessing and Feature Extraction segment is followed by Resnet-18 model training. The Resnet-18 architecture is customized and modified to accept single-channel spectrograms as 1 input channel instead of RGB 3. The output layer has three neurons corresponding to the classes ictal, preictal and interictal. The loss function used is “CrossEntropyLoss” suitable for multi-class classification. The Adam optimizer with weight decay and L2 regularization is used to prevent overfitting. The batch size and epochs are 32 and 50 respectively. The epochs have been implemented with early-stopping that stops the training process if the validation loss doesn’t improve for five epochs. Then the model is saved for future use.

4.3.3 XAI Integration

The major XAI techniques used in the system are SHAP, LIME, Grad-CAM, Integrated Gradients, and DeepLIFT. SHAP is model-agonistic or feature attribution-based technique. It applies to any model including deep models with SHAP DeepExplainer. It uses game theory (Shapely values) to compute the marginal contribution of each feature to the final prediction. The technique provides global and local explanations with a solid theoretical foundation. It can quantify how much each frequency or time frequency or time region contributed to a decision. It is useful for understanding the importance of specific frequencies or time points across many samples that is highly suitable for statistical insights into seizure biomarkers.

LIME is model-agonistic and follows local surrogate modeling. It can be applied to any model. It perturbs the input slightly, gets predictions and fits a simple interpretable model like linear regression locally to approximate decision boundaries. The unique strength of LIME is that it gives local faithful explanations around a single input and helps understand what changed the decision for a specific EEG signal. This technique works great for debugging misclassifications or

borderline cases by showing which segments/features of the spectrogram mattered in that moment.

Grad-CAM is a visual and attribution-based model. It is specially designed for CNN-based models. It uses gradients of a target class flowing into the last convolutional layer to produce a heatmap of important regions in the input. EEG spectrograms are 2-D time-frequency images and Grad-CAM is ideal for showing spatial spectral attention, making it intuitive for clinical interpretation.

Integrated Gradients is model-specific and gradient-based. It applies to differentiable models. It integrates the gradients of the output prediction from a baseline input to the actual input, giving robust attributions across input space. Integrated Gradients reduces issues with noisy gradients or saturation and provides more complete attributions than vanilla gradients. It accurately identifies which frequency bins or time segments led to the final classification.

Deep-Lift is model-specific and back-propagation based. It applies only to neural network-based models. DeepLIFT compares activations of neurons in the actual input to a reference baseline input and assigns contribution scores using differences, not gradients. This technique provides fast, non-gradient based attributions that work well even when the gradients are zero. It is more sensitive and stable. It can reveal subtle but important frequency shifts that might not show up well in gradient-based methods thus complementing Grad-CAM for non-spatial explanations.

Explainable AI (XAI) is a critical component in medical AI applications, ensuring transparency, interpretability, and trustworthiness of machine learning models. In deep learning-based seizure prediction, high-performing models like ResNet-18 often function as black boxes, making it difficult for clinicians to understand the reasoning behind their predictions. XAI techniques provide insights into the decision-making process of these models, allowing medical professionals to validate model reliability, detect biases, and build confidence in AI-assisted diagnosis. The techniques can be comprehensively presented as given in Table 4.2.

Table 4.2 Comparison of XAI Techniques

Technique	Type	How it works	Unique Qualities	EEG Sustainability
SHAP	Model-agnostic, Feature-based	Calculates Shapley values to determine each feature's contribution	Based on game theory, offers global & local interpretability	Quantifies importance of EEG frequencies/times across many samples
LIME	Model-agnostic, Local surrogate	Perturbs input and fits a simple interpretable model around it	Explains individual decisions by approximating the model locally	Helpful for debugging specific mis-classifications or borderline cases
Grad-CAM	Model-specific, Visual (CNNs)	Uses gradients of class score with respect to convolutional feature maps to generate heatmaps	Localizes important spatial regions on input (e.g., EEG spectrogram)	Highlights frequency-time regions important for predictions
Integrated Gradients	Model-specific, Gradient-based	Computes integral of gradients from baseline to input	Overcomes noisy gradients & saturation, offering complete attribution	Stable identification of important spectro-temporal components in EEG
DeepLIFT	Model-specific, Reference-based	Compares activations to a baseline, propagates contribution differences	Gradient-free, captures relevance even when gradients are zero (e.g., ReLUs)	Sensitive to subtle EEG signal patterns that may be missed by other methods

4.4 XAI TECHNIQUES USED

4.4.1 SHAP (SHapley Additive exPlanations)

SHAP (SHapley Additive exPlanations) is one of the Explainable Artificial Intelligence (XAI) techniques that provides both local and global explainability for machine learning and deep learning models. It is a model-agnostic and theoretically grounded technique used to explain the predictions made by machine learning and deep learning and models.

During each prediction, SHAP assigns a particular value to a feature called Shapley values. Shapley values indicate the amount of contribution a particular feature has in the model prediction. These values are derived from the cooperative game theory where each feature is a player in the game contributing to a final score. The model's final prediction is called a payout which has to be fairly distributed among all the players (features).

SHAP evaluates how the model's output changes when a particular feature is included or excluded. It evaluates the model's output over a range of combinations of all the input features. The goal of SHAP is to fairly distribute the model's output to all the available input features based on their contribution to the model prediction, ensuring that the sum of Shapley values assigned to the features in a single prediction equals the difference between the actual model prediction and the expected (baseline) prediction. This is to satisfy properties such as local accuracy, consistency and missingness.

In the context of the Epileptic Seizure Prediction using EEG signals, SHAP explains which time-frequency regions contribute the most to the model prediction of the seizure state (ictal, interictal, preictal). It also explains how the model relies on each feature and the patterns it relies on for model's prediction.

SHAP Global Explainability

Global explainability refers to understanding how a model behaves across an entire dataset instead of individual prediction. It gives insights on the contribution of the features and regions to a model's prediction.

SHAP does this by aggregating values across many samples, visualizing feature impacts to determine the most important features and to understand model's logic and feature relevance.

During each prediction, SHAP assigns a particular value to a feature called as Shapley values. Shapley values indicate the amount of contribution a particular feature has in the model prediction. These values are derived from the cooperative game theory where each feature is a player in the game contributing to a final score. The model's final prediction is called as payout which has to be fairly distributed among all the players (features).

For a given input, SHAP computes how much each pixel (feature) contributes to moving the model's output from baseline image to the actual image. Mathematically, for an input feature $i$the Shapley value ϕ_i is defined as given in the equation 4.3:

$$\phi_i = \sum_{S \subseteq N \setminus \{i\}} \frac{|S|!(|N| - |S| - 1)!}{|N|!} [f(S \cup \{i\}) - f(S)] \quad (4.3)$$

where:

- N is the set of all features
- S is any subset not containing feature i
- $f(S)$ is the model's output when only the feature in S are present.

This formula calculates the average marginal contribution of each feature across all the available feature subsets ensuring a fair distribution of the model's output to the available feature subsets.

In the Epilepsy Seizure prediction system, Resnet-18 model was trained on EEG signals which were then converted into spectrograms to classify them into one of the three stages: ictal, interictal and preictal. SHAP was applied to this trained model using the DeepExplainer module from the SHAP library. A small batch of spectrograms were selected from each class and passed to the explainer so that the SHAP values are calculated. These values are returned for each pixel in the spectrogram image in class-wise contribution format. It is in a three-array shape of [n,1,224,224] where n is the number of samples.

To understand the global behavior, SHAP values are visualized as Bar plot, summary plot and heatmap. The bar plot visualized the average magnitude of SHAP values for each class whereas the summary plot highlighted the most important regions across a class and heatmap visualized the contribution of the time-frequency region to the model's prediction in pixel level.

Bar Plot

SHAP based bar plots are used to provide a high-level understanding of the input features that have influenced the most in the model's decision-making process across different stages. It serves as a global interpretability tool that summarizes the average importance of each feature across a batch of spectrograms. Here, each pixel in a 224 x 224 spectrogram image is treated an individual feature and the ResNet-18 model classifies each image into one of the three seizure states: ictal, interictal, preictal.

Using SHAP's DeepExplainer, SHAP values are first computed for each class which results into a tensor of shape $(C, N, 1, 224, 224)$ where,

- $C = 3$ is the number of classes (ictal, preictal, interictal)
- $N = 50$ is the number of samples in each class
- 224 x 224 is the number of pixels in a spectrogram image.

These SHAP values calculate the contribution of each pixel to the model's predicted output for a given class.

To produce the bar plot, the SHAP values are reshaped into a 2D matrix of shape $(C, N, 50176)$, flattening the spatial dimensions (since 224 x 224 = 50176 features per image). The absolute values of these SHAP values are then averaged across all N images using the formula given in equation 4.4 as:

$$\phi_i^{(c)} = \frac{1}{N}\sum_{j=1}^{N}\left|\phi_{i,j}^{(c)}\right| \quad (4.4)$$

where,

- $\phi_i^{(c)}$ is the average SHAP importance for the i^{th} feature (pixel) for class c
- $\phi_{i,j}^{(c)}$ is the SHAP value of that feature for the j^{th} image.

This process gives a vector of average feature importances for each class. These vectors are then visualized using bar plots. The x-axis of the bar plot represents the feature indices which are the flattened pixel positions and the y-axis represents the average SHAP magnitude. This

highlights the regions in spectrograms which consistently contribute most strongly to the model's prediction of each class.

Summary Plot

The SHAP summary plot helps in visualizing and identifying the most influential features (pixels) across all the input spectrogram images. For this, Resnet-18 model was trained on EEG signals which were then converted into spectrograms to classify them into one of the three stages: ictal, interictal and preictal. Each spectrogram image of size 224 x 224 is flattened into a 1D vector of 50176 features. Using the DeepExplainer from the SHAP library, SHAP values for each feature are calculated which quantify how much each pixel (feature) contributes, whether positively or negatively, to the model's prediction. Mathematically, the SHAP value $\phi_i^{(j)}$ for feature i in input j is given by the equation 4.5:

$$\phi_i^{(j)} = \sum_{S \subseteq F \setminus \{i\}} \frac{|S|!(|F| - |S| - 1)!}{|F|!} [f_x(S \cup \{i\}) - f_x(S)] \quad (4.5)$$

where,

- F is the set of all input features
- S is a subset of features excluding i
- $f_x(S)$ is the model output when only feature in S are present

This formula ensures that the contributions are fairly distributed among the correlated features using Shapley values from the game theory.

For visualization purposes, the SHAP values for a chosen class are reshaped to (N, 50176) where N = 50 is the number of samples in a class. The summary plot is in the form of a scatter plot where each dot represents the SHAP score of a feature in an image. The x-axis represents the SHAP value which is the magnitude and direction of feature impact, whereas the y-axis ranks the features by their mean absolute SHAP values across all samples.

Heatmap

The SHAP heatmap provides a spatial visualization of different regions (pixels) in EEG spectrogram images and how these regions influence the model's prediction and classification into one of the three stages: ictal, interictal and preictal. For this, the ResNet-18 model uses

DeepExplainer from the SHAP library to compute the Shapley values from the cooperative game theory for each pixel in the input spectrogram. Mathematically the SHAP value ϕ_i^j for feature i in the input j is given by the equation 4.6:

$$\phi_i^j = \sum_{S \subseteq F \setminus \{i\}} \frac{|S|!(|F| - |S| - 1)!}{|F|!} \left[f_j(S \cup \{i\}) - f_j(S) \right] \quad (4.6)$$

where,

- F is the set of all input features
- S is a subset of features excluding i
- f_j is the model's prediction function

The computed SHAP values, shaped (50,224,224) for 50 images of size 224 x 224, are averaged across all samples to identify the most influential regions in the dataset for predicting a class.

The heatmap is generated by taking the mean SHAP value across 50 samples for a specific class and then identifying the consistent regions of importance across the dataset. The red color in the heatmap indicates positive contributions and blue indicates negative contributions of features to the prediction. The x-axis represents time and y-axis represents frequency which represent the time-frequency domain of the spectrogram. The color-bar is added to show the intensity of feature contributions towards the model prediction.

SHAP Local Explainability

SHAP is based on Shapley values from game theory, originally designed to fairly divide rewards among players in a game based on their contributions. The Resnet-18 model can be imagined as a game and each pixel in the spectrogram i. e the features are the players. Each player contributes some value to the final prediction. SHAP tries to quantify this value for each player. The architecture diagram of Local SHAP is shown in Fig 4.3.

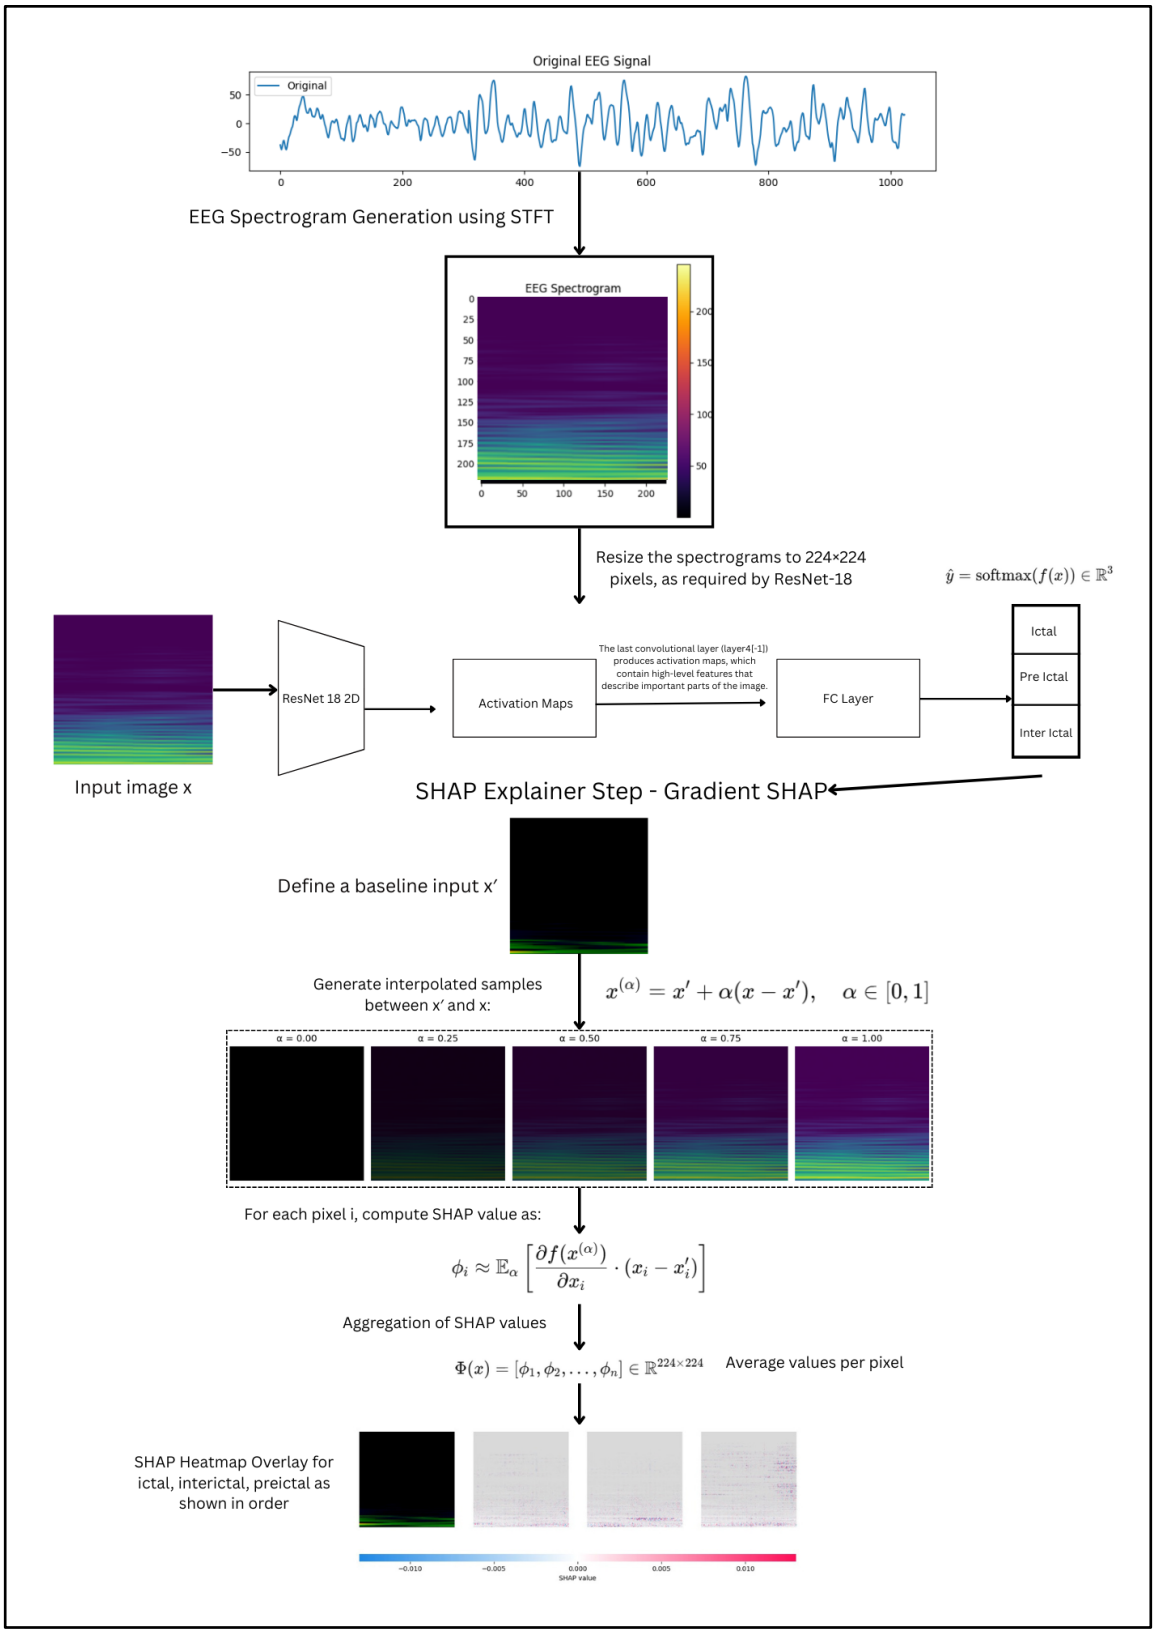

Fig 4.3 SHAP Architecture

The input EEG signal is converted to spectrogram using STFT feature extraction technique and resized to 224×224 pixels. This resized image is passed into the trained Resnet-18 model for classifying into ictal, interictal or preictal stages and the softmax probabilities for each class is outputted. This is followed by the SHAP Explainer step using the Gradient SHAP. A baseline input(often all zero-image) is defined first. This is called the neutral or background input. Then using the baseline and actual input images a bunch of inputs where some features are from actual input and others are from the baseline are generated. They are called interpolated images with varying α values as given in Equation 4.7.

$$x^{(\alpha)} = x' + \alpha(x - x'), \alpha \in [0,1] \quad (4.7)$$

This can be used to check how the output changes when a new feature is added which in turn gives the idea of importance of each feature. For each feature i Shapley value is calculated as given in the equation 4.8:

$$\phi_i \approx E_{x\prime \sim baseline, \alpha \sim [0,1]} \left[\frac{\partial f(x\prime + \alpha(x - x\prime))}{\partial xi} . (x - x_i')\right] \quad (4.8)$$

The gradients of model output with respect to inputs along this path is computed and multiplied with the input difference. This is done over many samples to average out the noise. The average values per pixel after aggregating across multiple interpolated paths is given by equation 4.9:

$$\Phi(x) = [\Phi_1, \Phi_2, \ldots, \Phi_n] \in R^{224 \times 224} \quad (4.9)$$

Finally, the SHAP values are reshaped back into the shape of the input spectrogram and the heatmap is plotted and overlaid.

4.4.2 LIME (Local Interpretable Model-Agnostic Explanations)

LIME is an XAI technique that provides visual explanations for the predictions made by the black box models like deep neural networks. LIME generates heat maps that highlight the most important regions that contributed towards the prediction or classification which helps in understanding the model's behavior. This allows us to understand which input features contribute most to the model's decision.

Working Of Lime:

Generating Perturbed Samples

The goal of LIME is to provide visualizations that indicate which time frequency region is most critical for the prediction. To do this, LIME approximate the behavior of a complex model locally for a given input instance. It generates multiple slightly modified versions of the input and observes how the model's predictions change. This process is called perturbation. Perturbation is a process which segments the image into super-pixels and then selectively turns some on or off. The architecture diagram of LIME illustrated in Fig 4.4 clearly shows the working mechanism of it in the project.

Image segmentation is the process of dividing an image into multiple regions based on pixel similarity (e.g., color, intensity, texture, and spatial proximity). Each of these regions is called a superpixel. Instead of treating individual pixels as separate entities, segmentation groups pixels into meaningful regions that represent distinct objects or patterns in the image. Given an image I consisting of pixels p_i, segmentation partitions I into k disjoint sets as given in equation 4.10:

$$\bigcup_{i=1}^{k} S_i = I, \quad S_{\mathrm{i}} \cap S_{\mathrm{j}} = \emptyset \; for \; i \neq j. \quad (4.10)$$

where each S_i is a super-pixel region. LIME uses segmentation algorithms like QuickShift or SLIC (Simple Linear Iterative Clustering) to create superpixels.

QuickShift Algorithm

QuickShift is a density-based clustering algorithm that groups nearby pixels into super-pixels by estimating a kernel density function as given in equation 4.11:

$$f(x) = \frac{1}{n.h^d} \cdot \sum_{i=1}^{n} K\left(\frac{x - x_i}{h}\right) \quad (4.11)$$

where:

- K is a kernel function (e.g., Gaussian),
- h is the bandwidth parameter (determines cluster size),
- d is the number of dimensions in the image.

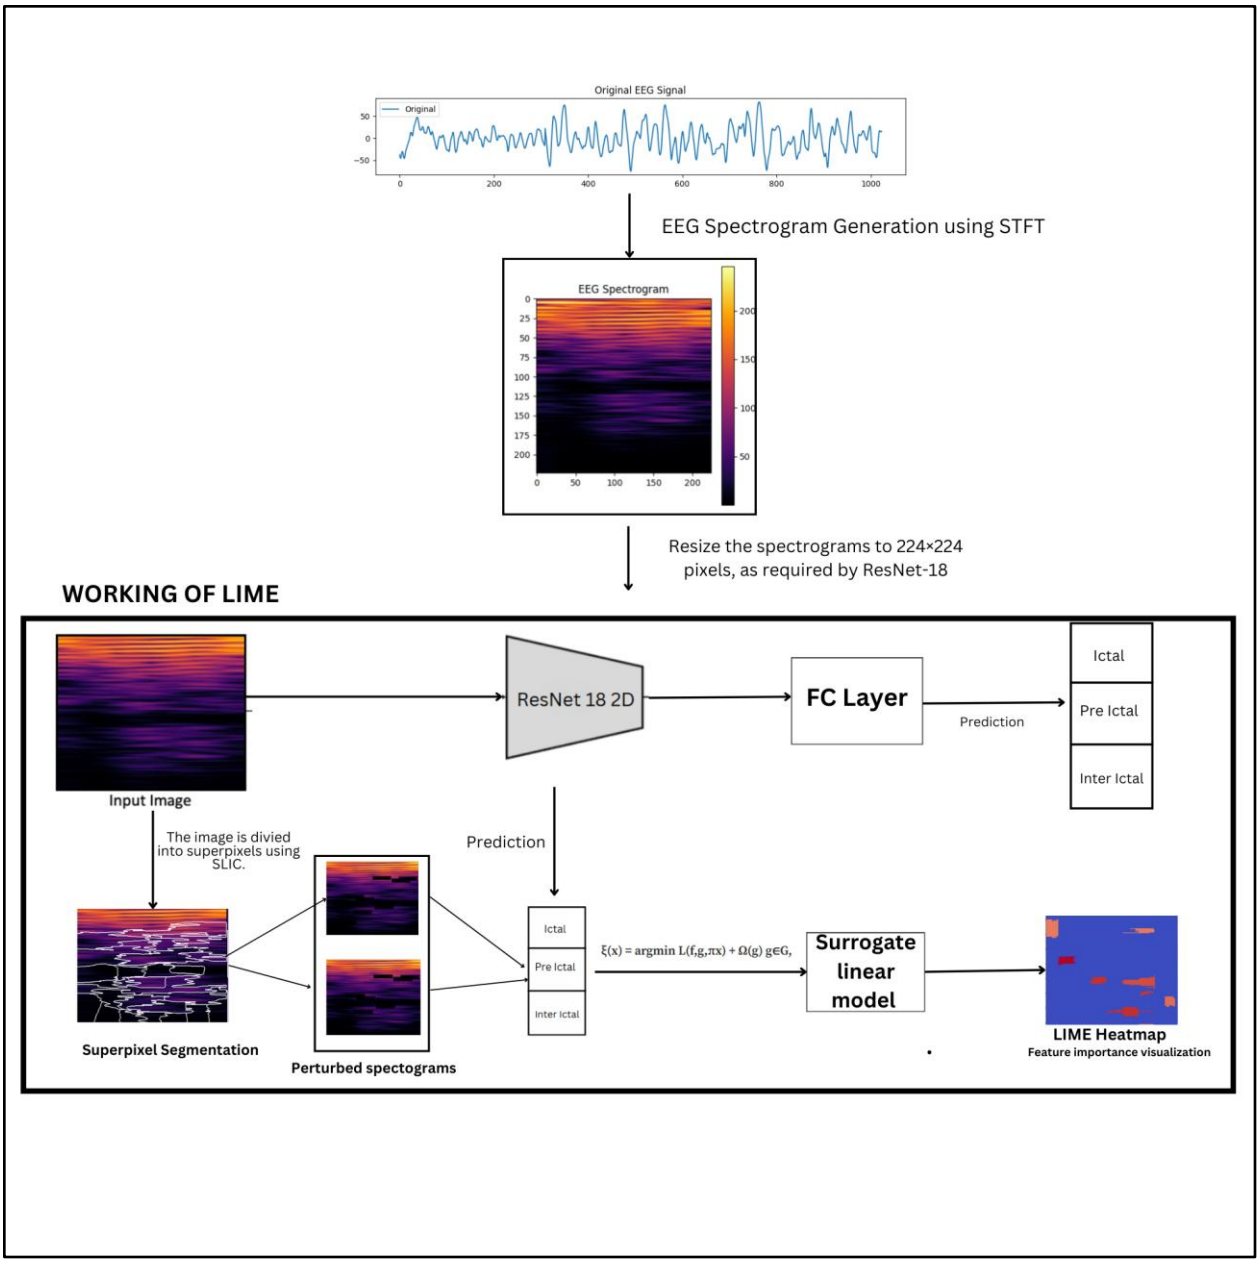

Fig 4.4 LIME Architecture

The algorithm follows a gradient ascent approach which groups pixels based on their spatial proximity and color similarity. After segmentation, the image is represented as a set of k super-pixels as given in equation 4.12:

$$I = \{S_1, S_2, \ldots, S_k\} \quad (4.12)$$

LIME creates perturbed versions of the image by randomly selecting super-pixels to turn off (remove) or keep on (preserve). A binary vector b represents which super-pixels are active is given by equation 4.13:

$$b = \{b_1, b_2, \ldots, b_k\} \quad (4.13)$$

where 0 – inactive, 1 - active

Super-pixels that are turned off are replaced with either:

- A uniform color (e.g., black or gray).
- The average pixel intensity of that region.

LIME generates multiple perturbed images by randomly turning on or off the super-pixels.

Assigning Weights to Perturbed Instances

Not all perturbed samples contribute equally to the explanation. Some perturbed versions of the image are more similar to the original instance, while others may be different from the original one. To ensure that the model's explanation revolves around the original instance, LIME assigns higher importance to similar perturbed instances. This similarity is measured using a proximity function $\pi_x(z)$, which determines how much weight can be given to a particular perturbed sample z when approximating.

Query the Black-Box Model

Each of these perturbed samples is fed into the original deep learning model to obtain predictions.

Train an Interpretable Model

Now LIME has a set of perturbed samples, along with their predicted outputs and weights, it trains a simple model (such as a linear regression or decision tree) to approximate the deep learning model's behavior only for this specific instance. The simple model is not meant to replace the original model. It is only used to mimic its predictions in a small local region. Since the simple model is interpretable, it allows us to extract meaningful feature importance values.

$$\xi(x) = argmin_{(}g \in G_{)} \sum_{(}z \in Z_{)} \pi_{x}(z) \cdot (f(z) - g(z))^{2} + \Omega(g) \quad (4.14)$$

where,

- G is the set of interpretable models (e.g., linear models).
- $\pi_x(z)$ is the proximity function, which assigns higher weights to perturbed samples that are closer to the original image x.
- $f(z)$ is the prediction of the black-box model for the perturbed image z.
- $g(z)$ is the prediction of the interpretable model.
- $\Omega(z)$ is a regularization term that prevents overfitting by penalizing overly complex models.

Since we are using a linear regression model as the surrogate model, it takes the following form as given in the equation 4.15

$$\hat{y} = w_0 + \sum_{i=1}^{k} w_{i} z_{i} \quad (4.15)$$

where:

- $\hat{y}$ is the predicted output from the surrogate model.
- w_0 is the intercept.
- w_i represents the importance weight for the i^{th} super-pixel.
- z_i is the feature corresponding to that super-pixel (it is 1 if the super-pixel is active and 0 if it is removed)

For example, in EEG-based seizure classification, LIME might fit a linear model that assigns importance to different frequency bands in the EEG spectrogram. If a high-frequency region has a high weight, it means that region was critical in the model's seizure prediction.

Generate the Explanation

Once the interpretable model is trained, it provides feature importance scores, which tell us which parts of the input were most responsible for the final prediction. To compute the relative contribution of each super-pixel, an **importance score** is assigned using the following formula given in equation 4.16:

$$Importance\ Score = \frac{|w_i|}{\sum_{j=1}^{m} |w_j|} \quad (4.16)$$

where:

- w_i is the absolute weight assigned to super-pixel i.
- m is the total number of super-pixels in the image.
- The denominator ensures that all scores sum up to 1, making them interpretable as relative importance values.

LIME highlights the most important super-pixels in the image that influenced the model's decision. The heatmap is generated by mapping these scores back to the original image.

4.4.3 Grad-CAM (Gradient-Weighted Class Activation Mapping)

Grad-CAM is an XAI technique that provides visual explanations for the predictions made by the black box models like deep neural networks. Grad-CAM generates heat maps that highlight the most important regions that contributed towards the prediction or classification which helps in understanding the model's behavior. This technique is particularly useful for understanding CNN based models that apply to image classification tasks. Grad-CAM assigns importance scores to different spatial regions of the image by using the gradients of the model's final convolutional layer thereby enabling interpretability in deep learning models.

Working Mechanism of Grad-CAM

Grad-CAM works systematically in two phases as forward pass where class score is computed and in the backward pass gradients are computed and heat map is generated by global average pooling of importance weights. The architecture diagram of Grad-CAM illustrated in Fig 4.5 clearly shows the working mechanism of it in the project. Initially the EEG spectrogram image generated from the .mat input file is passed through ResNet-18's convolutional neural network. Feature maps are generated at different layers. The gradients of the predicted class score of the feature maps of the last convolutional layer are computed during backpropagation. In each feature map, these gradients are averaged across all pixels to find their relative importance. The computed weights are then multiplied with their corresponding feature maps to obtain the Grad-CAM heatmap and the weighted sum. This weighted sum is finally passed through the ReLU activation

function where it replaces the negative values to zero thereby considering only the positive influences.

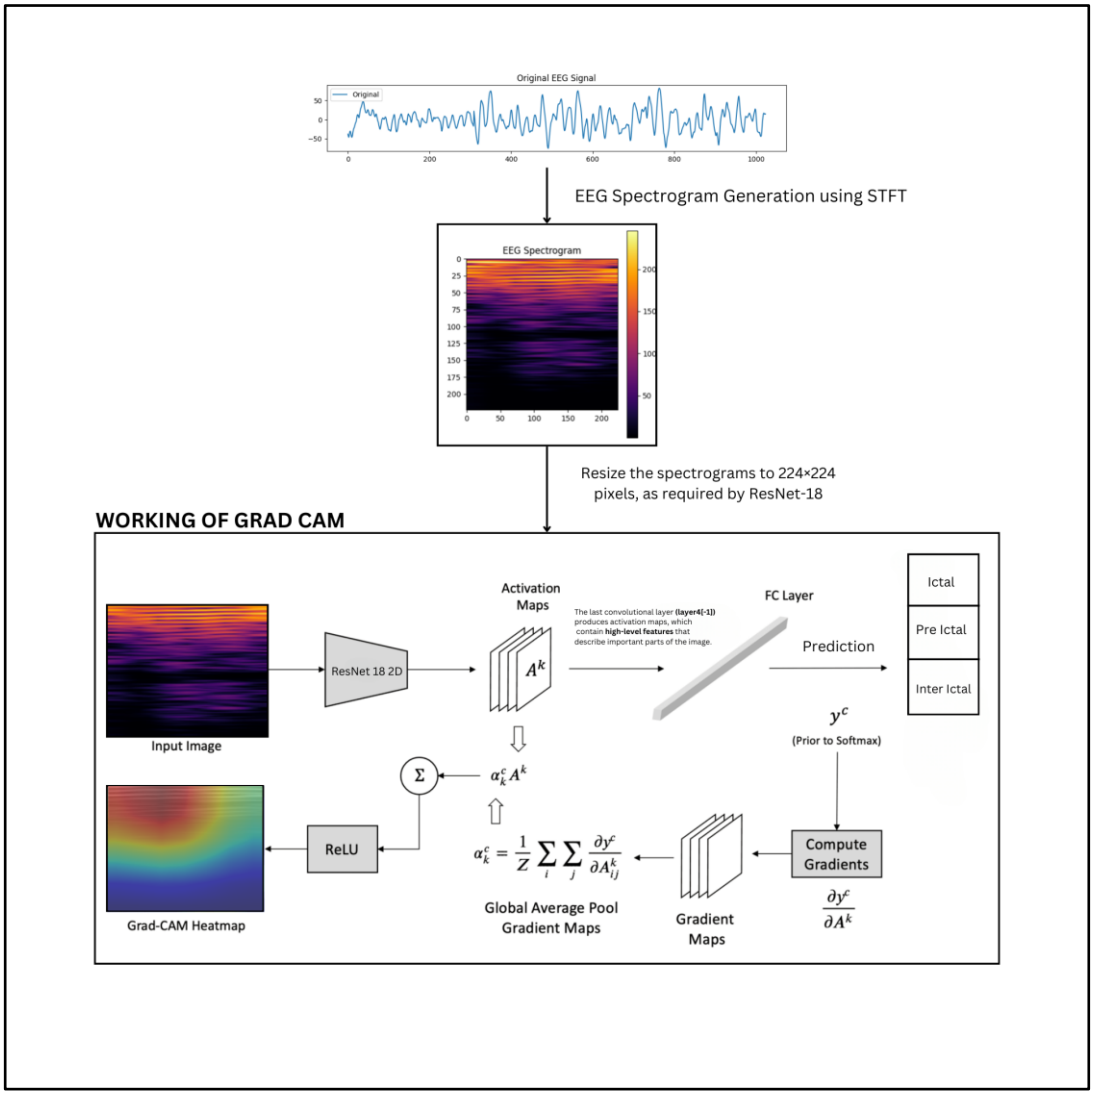

Fig 4.5 Grad-CAM Architecture

The first step is spectrogram generation from EEG signals. A raw EEG signal, denoted by $x(t)$ transformed into time-frequency domain representation using the STFT. It is given by the equation 4.17:

$$X(f,t) = \sum_{-\infty}^{+\infty}(n)w(n-t)e^{(-j2\pi fn)} \qquad (4.17)$$

where $x(n)$ is the raw EEG signal, $w(n-t)$ is a window function that is used to localize the signal in time, f is frequency and t is time. The magnitude of $X(f,t)$ is normalized to generate a

spectrogram of size 224×224 pixels because Resnet-18 architecture accepts images only in that dimension.

The second step is feature extraction using the ResNet-18 model trained using the dataset[30]. The Resnet-18 model extracts the spatial features from the normalized spectrogram I input to the model using its hierarchical convolutional layers. The final convolutional layer $Layer[-1]$ outputs the feature map denoted as $A^k \epsilon R^{(H \times W)}$ where k denoted the index of the channel or feature map and $H \times W$ is the spatial resolution of each feature map. These activation maps play an important role because they give the high-level discriminative features essential for final classification or prediction.

The third step is gradient computation for the target class. In order to interpret the model's decision for a particular target class c (ictal, pre-ictal or interictal), the gradient of the class score y^c is computed with respect to activation maps before the soft-max layer. The gradient is given by the equation 4.18:

$$\frac{\partial y^c}{\partial A_{ij}^k} \quad (4.18)$$

Where y^c is the prediction score for class c and A_{ij}^k represents the activation map at spatial location (i, j) of the k^{th} feature map.

The fourth step is obtaining the importance weights α_k^c by using global average pooling technique for the gradients over the spatial dimensions. This reflects the contribution of each feature map k to the class c. The global average pooling equation is given by the equation 4.19:

$$\alpha_c^k = \frac{1}{Z} \sum_i \sum_j \frac{\partial y^c}{\partial A_{ij}^k} \quad (4.19)$$

where $Z = H \times W$ is the total number of spatial locations. These weights help to quantify the relevance of each channel A^k for the prediction.

The fifth step is Grad-CAM heatmap L^c generation by computing the weighted linear combination of feature maps that is then passed into a ReLU activation function as given in equation 4.20.

$$L^c = ReLU\left(\sum_k \alpha_c^k A^k\right) \quad (4.20)$$

The resulting heatmap $L^c \epsilon R^{H \times W}$ is resized to match the input spectrogram dimension(224×224). This process is called upsampling. The heatmap is then superimposed onto the original spectrogram to visually explain the model's decision for the particular EEG signal to be classified under a specific stage as ictal, interictal or preictal. The softmax is used to generate class probabilities for the model's prediction computed in the FC layer as given by equation 4.21.

$$P(y^c) = \frac{e^{y^c}}{\sum_j e^{y^j}} \quad (4.21)$$

where $P(y^c)$ is the probability of the input class belonging to class c and y^j are the raw output scores for all classes.

4.4.4 Integrated Gradients

Integrated Gradients (IG) is an Explainable AI (XAI) technique used to understand the feature which has contributed the most for a model's prediction. In the context of Epileptic Seizure prediction using EEG spectrograms, IG helps to identify the most important frequency-time regions which influence the classification of EEG signals into three seizure states: ictal, preictal and interictal. This method provides insight into the decision-making process of deep learning models, improving their interpretability and trustworthiness.

Mathematical Foundation of Integrated Gradients

The core idea behind Integrated Gradients is to compute the importance of each feature by accumulating gradients along a straight-line path from a baseline input x' which is usually a neutral or zero-valued input, to the actual input x. The feature attribution for each dimension i of input x is given by the equation 4.22:

$$IG_i(x) = (x_i - x_i') \times \int_{\alpha=0}^{1} \frac{\partial F(x\prime + \alpha(x - x\prime))}{\partial x_i} d\alpha \quad (4.22)$$

where:

- $F(x)$ represents the model's prediction function.
- x is the actual input (EEG spectrogram).
- x' is the baseline input.
- α is a scaling factor that moves the input from x' to x.
- $\frac{\partial F}{\partial x_i}$ denotes the gradient of the model's output with respect to feature x_i.

By summing the gradients over multiple steps along the interpolation path from x' to x, Integrated Gradients effectively captures how each feature influences the prediction. This ensures that the attributions satisfy key interpretability properties like sensitivity and completeness.

Working of Integrated Gradients

The architecture diagram depicted in Fig 4.6, for seizure prediction, begins with the transformation of raw EEG signals into spectrograms using the Short-Time Fourier Transform (STFT). This transformation helps the model to analyze both frequency and time-domain features of brain activity, which are crucial for distinguishing between the three stages:

- Ictal (Seizure Event)
- Pre-Ictal (Before Seizure)
- Inter-Ictal (Normal Brain Activity)

Choosing a Baseline Input

The baseline input represents an EEG spectrogram with minimal or neutral activity. A commonly used baseline in seizure classification is:

- A black or zero-valued spectrogram representing no brain activity.
- An averaged EEG spectrogram representing background brain activity.

Generating Interpolated Inputs

To compute Integrated Gradients, the actual EEG spectrogram is gradually transformed from the baseline through interpolation. Interpolation is generating a series of spectrograms at different interpolation steps where a smooth transition takes place from the baseline to the actual spectrogram. The interpolation follows the equation 4.23:

$$x_\alpha = x' + \alpha(x - x'), \alpha \in [0,1] \quad (4.23)$$

where:

- x' is the baseline spectrogram.
- x is the actual spectrogram.
- α is the interpolation coefficient that smoothly transitions from the baseline to the actual spectrogram.

A total of 224 interpolated images are generated to ensure a smooth transition between the baseline and the actual input.

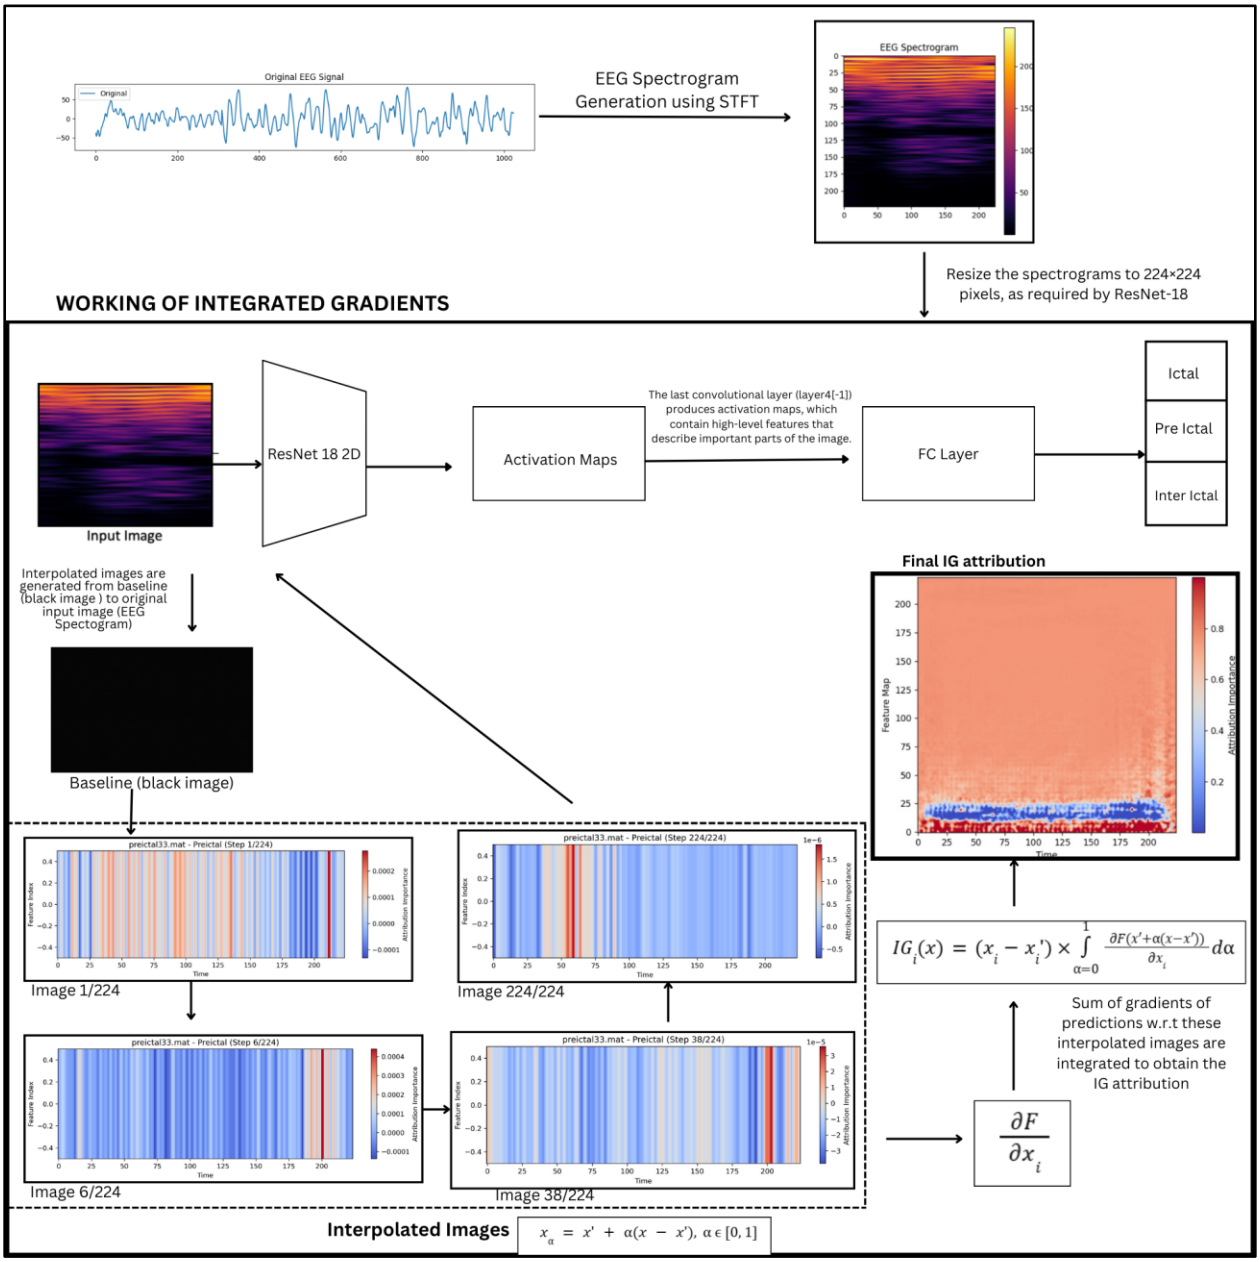

Fig 4.6 Integrated Gradient Architecture

The spectrograms are resized to 224×224 pixels to match the input requirements of the ResNet-18 2D model, which extracts spatial features through convolutional layers. The activation maps generated at layer4[1] highlight patterns related to seizure, which are then processed by the fully connected (FC) layer for final classification.

In the context of seizure prediction using EEG spectrograms, Integrated Gradients is applied to understand which regions of the spectrogram contribute the most to the model's classification decisions. The process can be understood in the following steps:

Computing Gradients

For each interpolated spectrogram, the gradient of the model's output with respect to the input is computed. The gradient at each step is given by equation 4.24:

$$\frac{\partial F}{\partial x_i} \quad (4.24)$$

where $F(x)$ is the model's prediction function. These gradients quantify how small changes in the spectrogram affect the model's confidence in classifying the EEG signal.

Accumulating Gradients and Computing Feature Attributions

The Integrated Gradients attribution for each feature is computed by summing the gradients over all interpolated steps and multiplying by the difference between the actual spectrogram x and the baseline x' as given by the equation 4.25:

$$IG_i(x) = (x_i - x_i') \times \int_{\alpha=0}^{1} \frac{\partial F(x\prime+\alpha(x-x\prime))}{\partial x_i} d\alpha \quad (4.25)$$

where:

- $IG_i(x)$ is the attribution score for pixel i .
- The integral sums the gradients across all interpolated images.
- The difference $(x_i - x_i')$ scales the attributions according to the actual spectral intensity.

Visualizing Feature Importance with Heatmaps

The computed attribution scores are visualized as heatmaps over the EEG spectrogram. These heatmaps highlight the frequency-time regions which influence the most towards the model's prediction:

- Bright regions indicate areas that strongly influenced the model's decision.
- Dark regions represent less significant contributions.

4.4.5 DeepLIFT (Deep Learning Important FeaTures)

DeepLIFT is an attribution method that works by comparing the activation differences between an input and a *chosen reference baseline*. DeepLIFT works by keeping track the difference of neurons activations between the actual input and the baselines and propagating it up to each layer. Pick a technique that quantifies the extent to which each feature in the input contributes to the final prediction. Using DeepLIFT on EEG-based seizure classification, we identify specific areas that are most influencing the model's decision, which can provide valuable insight.

The first stage of the architecture consists of inputting EEG data in the form of spectrogram images inside the model. These show a frequency-time representation of the signals using spectrograms which enables the model to understand temporal and spectral patterns for each of the given seizure states. A baseline image one is provided together with the actual input spectrogram and is a point of comparison for the DeepLIFT algorithm. This baseline is usually a black image (all-zero activation) in most of the cases but here we have used a blurred spectrogram image as the baseline image. We can assess the relative importance of spectral features by perturbing the input and measuring the model activation against the baseline.

The RESidual NETwork (ResNet) processes the EEG spectrograms. It focuses on the core spatiotemporal aspects of the types of convulsions by utilizing convolutional layers in the ResNet to learn essential spatial and frequency-oriented feature representations critical to the identification of seizure patterns. The activation maps (denoted as A^k) obtained from the last convolutional layer represent high-level learned features. The activation maps for each of the layers encapsulate this important information before potentially being reduced to a single classification through the classification layers.

One of the main benefits of our architecture is its interpretability, facilitated by DeepLIFT (Deep Learning Important FeaTures). We therefore apply DeepLIFT to the activation maps before passing it to the classification layers. DeepLIFT works by determining the difference in activations ΔA^k of the input spectrogram compared to the baseline image as given in the equation 4.26. This difference represents the amount by which a given feature deviates from a neutral reference state. DeepLIFT assigns contribution scores $C_{\Delta A^k \rightarrow \Delta y}$ for each layer that indicate the change in activations in that layer and how much they impacted the final output. These scores bring out the maximal regions in the input spectrogram that cause the model's classification decision.

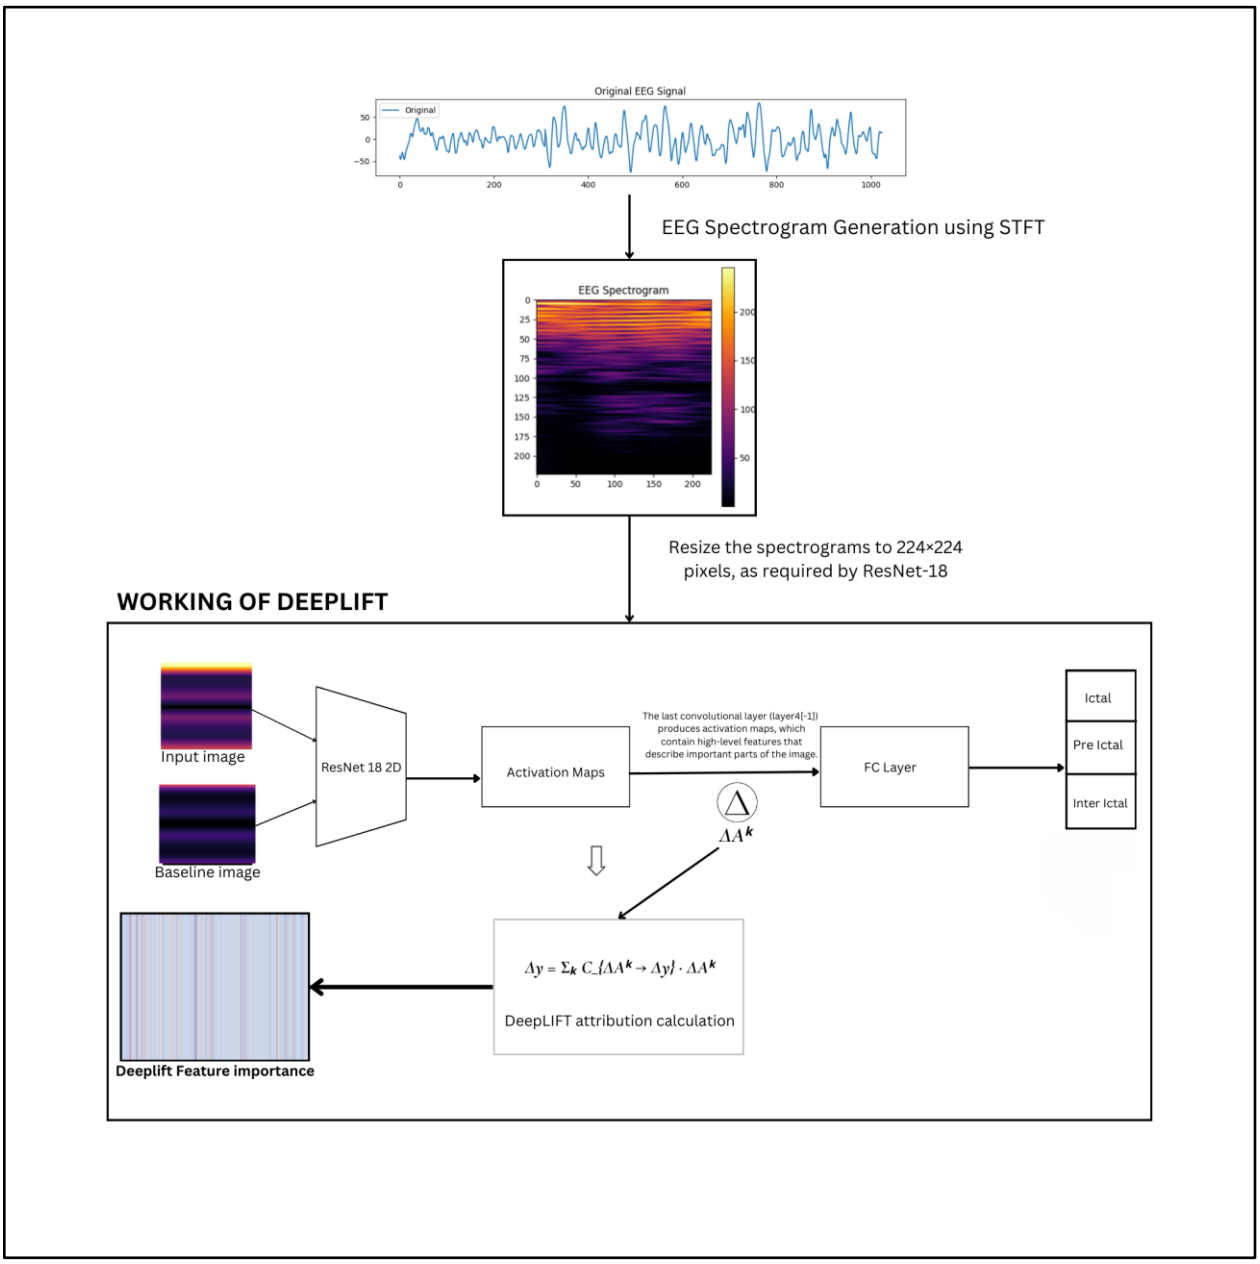

Fig 4.7 DeepLIFT Architecture

The key formula employed by DeepLIFT is:

$$\Delta y = \sum_k C_{\Delta A^k \to \Delta y} \cdot \Delta A^k \quad (4.26)$$

where ΔA^k represents the change in activation compared to the baseline, C (Contribution Score) quantifies the impact of each activation on the final prediction, and Δy is the total change in model output, which helps in interpreting feature importance.

The activation maps after the DeepLIFT application are fed into a Fully Connected (FC) layer to perform final classification. By inducing meaningful input feature representations, the DeepLIFT contribution scores guarantee that the FC layer will process it in terms of the transformed ΔA^k values. This layer produces an output that represents the predicted seizure state of the EEG signal. The outcomes of the classification can fall into any one of the three classes: either to the Ictal (Seizure Phase) Class, when the EEG signal achieved during an active seizure event; or to the Interictal (Between Seizures) Class, when the EEG is collected in a normal brain state (between seizures); or to the Preictal (Before Seizure Onset) Class, when EEG contains neuronal activity indicating imminent seizure.

The last step is to visualize the attributions calculated by the DeepLIFT. This is done by mapping the contribution scores to the original spectrogram image and obtaining a heatmap. This heatmap shows precisely which parts of the EEG spectrogram were most pivotal in producing the model's output. These heatmaps help clinicians and researchers better understand specific EEG features important to improve seizure prediction.

CHAPTER 5

METRICS

In XAI integration, to evaluate the quality of the explanation we use evaluation metrics. For our project, we chose fidelity, localization and stability to be our evaluation metrics. This decision was based on literature study, where we found the mentioned three metrics to be the most commonly used in evaluating XAI techniques. Since our project focuses on a comparative analysis of these techniques, we have selected the common metrics for fair and consistent evaluation.

5.1 FIDELITY

Fidelity measures how well an explanation captures the behavior of the original model by evaluating whether the model's prediction is influenced by the features highlighted by the explanation.

Mathematical Definition:

Fidelity is computed by comparing the model's original prediction with the prediction made after removing or perturbing the input features that contribute more to the model's prediction. A common approach is:

$$Fidelity = 1 - \frac{1}{N} \times \sum_{i=1}^{n} |f(x_i) - f(x_i')|$$

where,

- $f(x_i)$ is the model's original prediction for the input x_i
- x_i' is the input after removing (masking/zeroing out) the top-k important features from the explanation
- N is the number of samples

Fidelity score ranges between 0 and 1. A high fidelity value indicates that the model prediction changes when the highlighted features are removed. This ensures the explanation correctly identifies influential features. Based on literature, a technique is good if its fidelity score is *0.8* or above, moderate if it's between *0.6 - 0.8* and poor if less than *0.6*.

In our project we use fidelity to assess how effectively each of the five XAI techniques captures the critical input features that drive the model's classification decision. First we generate the attribution map and mask out the top -k important features identified. Then we run the modified input through the trained classifier and the difference in prediction score before and after masking gives us the fidelity score.

5.2 LOCALIZATION

Localization measures how well an explanation method concentrates its importance scores in the most relevant areas of the input. A good explanation should highlight only a focused and meaningful portion of the input rather than distributing importance widely.

Mathematical Definition:
Localization is calculated as the proportion of attribution values that are greater than the mean of the attribution map:

$$Localization = \frac{1}{N}\sum_{i=1}^{N}(A_{\mathrm{i}} > \mu_{\mathrm{a}})$$

where:

- A_{i} is the attribution value at pixel i
- μ_{a} is the mean of all attribution values
- $1\ (condition)$ is the indicator function that returns 1 if the condition is true, 0 otherwise
- N is the total number of pixels in the attribution map

This metric helps determine how localized or spread out the attribution values are in the explanation map. Localization score ranges from 0 to 1. A higher localization score indicates that the explanation is concentrated in a smaller, more significant region, implying better focus and

interpretability. According to literature, a technique is considered to have good localization if the score is 0.7 or above, moderate if it falls between 0.5 – 0.7, and poor if less than 0.5.

In our project, we use localization to evaluate how sharply each of the five XAI techniques identifies the key regions that contribute to the model's decision. After generating the attribution map for an input sample, we compute the proportion of values above the mean, reflecting how focused or scattered the explanation is.

5.3 STABILITY

Stability measures the robustness of an explanation method by evaluating whether similar inputs produce similar explanations. A stable explanation technique should yield consistent attribution maps even when the input is slightly perturbed.

Mathematical Definition:
Stability is calculated using the formula:

$$Stability = 1 - D(E_1, E_2)$$

where,

- $E_1 \; and \; E_2$ are explanations for slightly perturbed versions of the same input.
- $D(E_1, E_2)$ is a distance function that quantifies the difference between these explanations (e.g., cosine similarity, variance, etc.)

Stability score ranges between 0 and 1. A high stability value indicates that the explanation technique consistently highlights similar features despite small changes in the input, thereby showing that it is less sensitive to noise. Based on literature, a technique is considered good if its stability score is 0.8 or above, moderate if it's between 0.6 - 0.8, and poor if less than 0.6.

In our project, we assess stability by introducing minor perturbations to the original input and generating attribution maps for both the original and perturbed versions using each of the five XAI techniques. We then compute a distance metric (entropy and variance-based) between the two explanations and convert it into a stability score, which helps us evaluate how consistent the technique is across input variations.

CHAPTER 6

RESULTS AND OBSERVATION

6.1 XAI VISUALIZATIONS AND METRICS

The XAI techniques are used to visualize the contributions of time-frequency components of the EEG spectrogram. This section covers the visualizations generated using each technique and provides insights on how each visualization is used to explain the prediction. Each technique has its unique strengths and helps us understand and analyze all possible aspects of prediction. The major SHAP visualization techniques are summary plot, beeswarm plot, force plot, waterfall plot, dependence plot, and interaction plot. LIME uses visualization techniques like feature importance bar chart, local explanation heat maps, and perturbation-based analysis. Grad-CAM generates attribution heatmaps, overlay visualization, and path integrated visualization. Difference heatmaps, contribution score visualization, and saliency map are major visualization techniques of DeepLIFT. The primary goal is to ensure that the model focuses on medically relevant features and supports reliable interpretation.

The metrics for XAI are used to assess if the explanations generated by the techniques are performing well. The major metrics of evaluation used by SHAP are fidelity, consistency, stability, monotonicity, sparsity, and feature importance correlation. The LIME metrics are local fidelity, stability, perturbation robustness, sparsity, and comprehensibility. The important Grad-CAM metrics are localization score, saliency mass accuracy, IoU with ground truth, and faithfulness. Completeness, sensitivity, implementation invariance, gradient saturation, and robustness are metrics applicable for Integrated Gradients. DeepLIFT can implement metrics like summation to delta, implementation invariance, conservation property, smoothness, and contribution agreement. But in the project, the metrics fidelity, localization, and stability are the ones implemented for all the techniques as they are the common metrics and are suitable for comparison and analysis.

6.1.1 SHAP

SHAP (SHapley Additive exPlanations) plays a crucial role in providing both global and local interpretability in the context of epileptic seizure prediction using EEG signals. Its foundation in cooperative game theory allows SHAP to assign specific importance values to each input feature, explaining how much each time-frequency component of the EEG contributes to the model's prediction. On a global scale, SHAP helps uncover general patterns and feature contributions across the entire dataset, providing insights into what brain signal characteristics are most significant in differentiating seizure stages. Locally, SHAP delivers instance-specific explanations, explaining the reasons behind the classification of a specific EEG sample as ictal, preictal or interictal, which is especially important for real-time clinical decision-making. This dual functionality makes SHAP a powerful tool for building transparent, trustworthy and human-understandable AI systems in healthcare.

SHAP Global Explainability

Global explainability using SHAP is vital for understanding how a model behaves across an entire dataset, rather than just for individual predictions. In the context of seizure prediction from EEG signals, global SHAP analysis helps in identifying the time-frequency features that consistently influence the model's decisions for classifying ictal, preictal, and interictal stages. This comprehensive view is crucial for ensuring that the model is aligned with established medical knowledge and is not relying on irrelevant or misleading patterns. By analyzing SHAP's global visualizations—such as bar plots, summary plots, and heatmaps—we can gain a deeper understanding of feature importance trends, distribution of attributions, and interactions between features, thereby confirming the model is both interpretable and clinically reliable. The following sections explain each of these plots in detail and how they contribute to global model understanding. The following figures Fig 6.1, Fig 6.2 and Fig 6.3 explains each of the SHAP global visualizations -such as bar plot, summary plot and heatmap respectively attributed for the entire dataset.

Bar plot:

The bar plot represents the features contributing to the model prediction in each class. The x-axis of the bar plot represents the feature indices which are the flattened pixel positions and the y-axis represents the average SHAP magnitude. This highlights the regions in spectrograms which

consistently contribute most strongly to the model's prediction of each class. The bars represent how much each feature (indexed 0 to 49) contributes to prediction in each of the three classes.

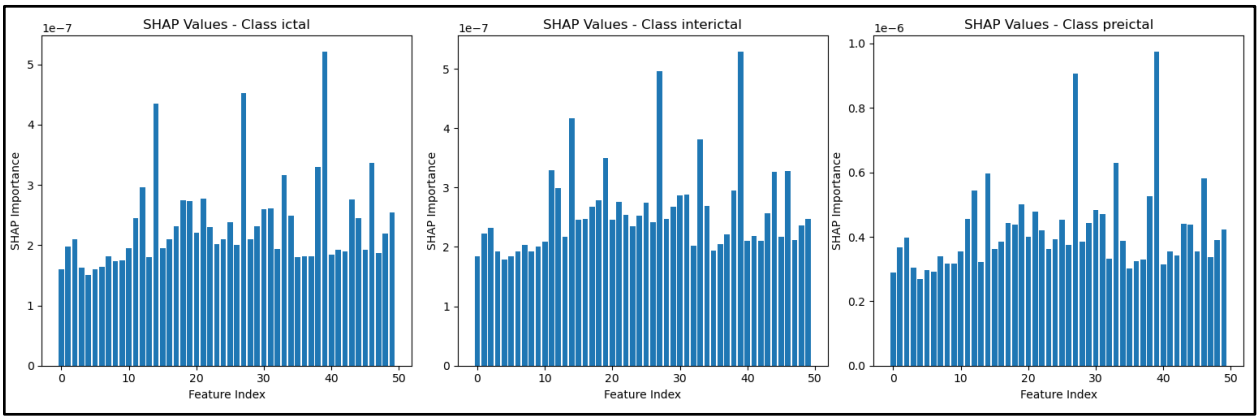

Fig 6.1 SHAP Bar plot

Class Ictal:

- Higher bars indicate that features have a greater global impact on the model's decision for classifying data as ictal.
- This means the model heavily relies on a specific subset of features, which might correspond to time-frequency regions sensitive to seizure onset.

Class Interictal:

- The spread is relatively consistent, but some features stand out with higher SHAP importance.
- This implies that although many features contribute, a few features are key in identifying interictal brain activity.

Class Preictal:

- These SHAP values are generally higher in magnitude.
- The plot shows that certain features are significantly more influential in predicting preictal stages, which are critical for early seizure detection.
- This emphasizes that the model can effectively distinguish early signs of seizures by focusing on a few crucial feature patterns.

Summary Plot:

The summary plot is in the form of a scatter plot where each dot represents the SHAP score of a feature in an image. The x-axis represents the SHAP value which is the magnitude and direction

of feature impact, whereas the y-axis ranks the features by their mean absolute SHAP values across all samples.

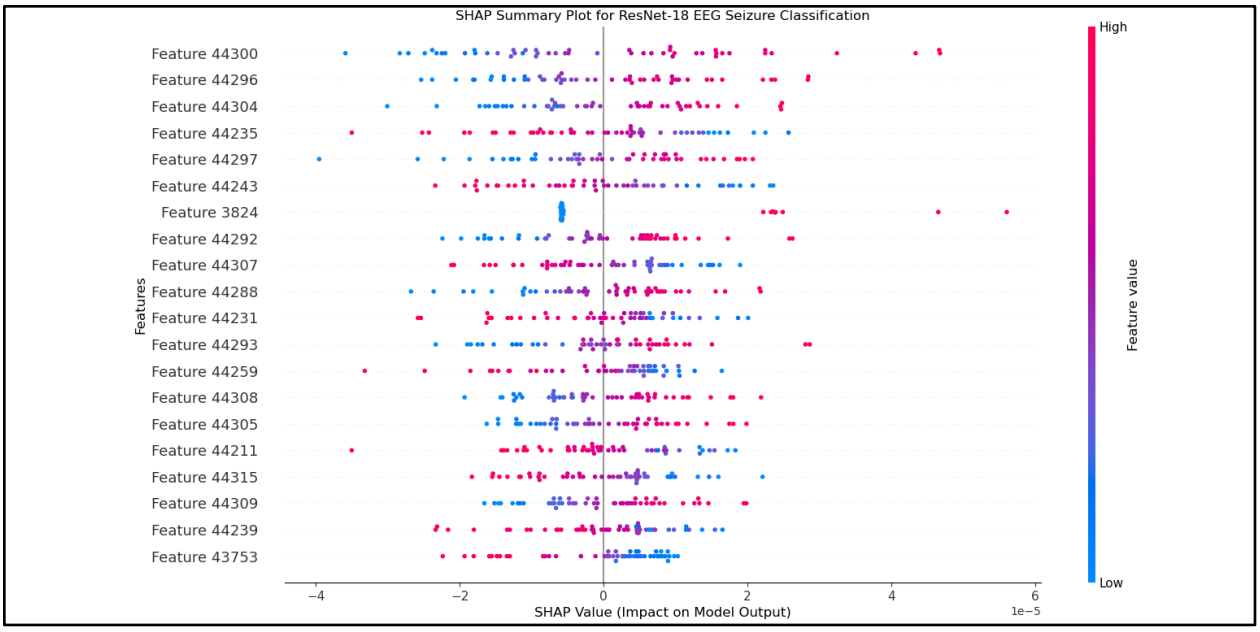

Fig 6.2 SHAP Summary Plot

The summary plot can be understood with the following conclusions:

Feature Influence Direction:

- If a feature's high value (shown in red) is associated with positive SHAP values, it indicates that higher values of that feature contribute to a greater probability of seizure classification.
- Conversely, if red is mostly on the left (negative SHAP), high values decrease seizure prediction probability.

Feature Consistency:

- Features with tightly clustered SHAP values indicate consistent impact on the model.
- Features with widely spread SHAP values have variable contributions, possibly interacting with other features or having non-linear effects.

Global Importance Ranking: The features at the top (like 44300, 44296) have the highest overall impact across all predictions, indicating the model relies on them more heavily.

Heatmap:

The heatmap is generated by taking the mean SHAP value across 50 samples for a specific class and then identifying the consistent regions of importance across the dataset. The red colour in the heatmap indicates positive contributions and blue indicates negative contributions of features to the prediction. The x-axis represents time and y-axis represents frequency which represent the time-frequency domain of the spectrogram. The colourbar is added to show the intensity of feature contributions towards the model prediction.

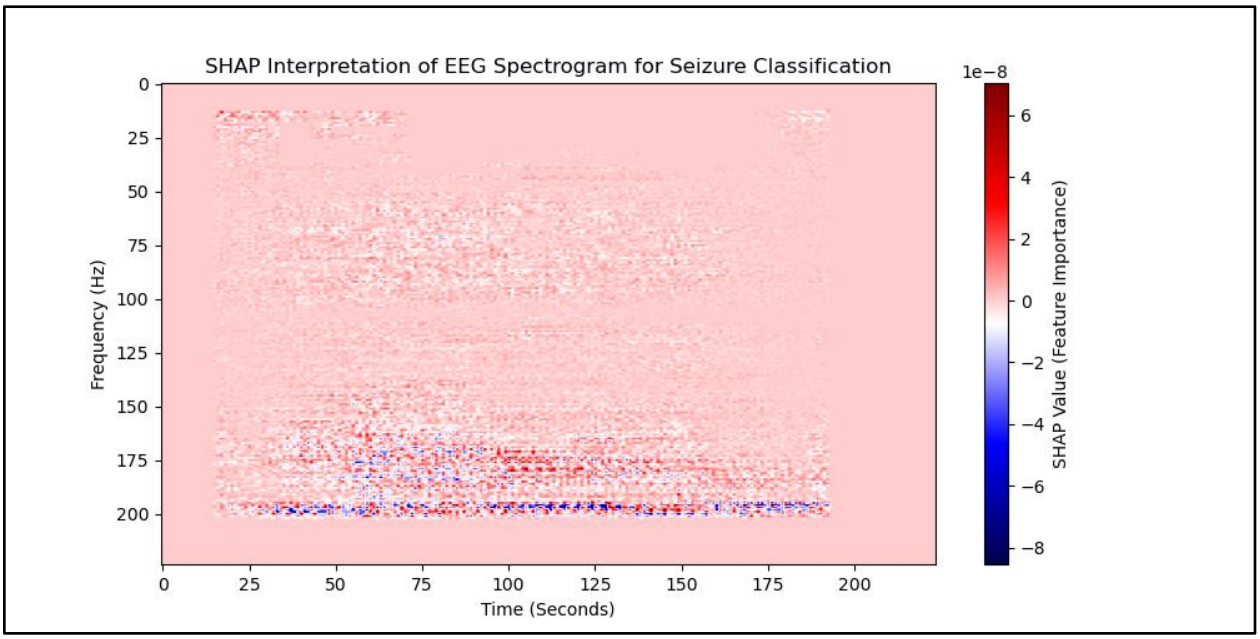

Fig 6.3 SHAP Global Heatmap

The heatmap can be interpreted using the following prominent zones:

Active Contribution Scores: The plot illustrates concentrated regions of red around lower frequencies (0–30 Hz) and sporadic bursts around mid-to-high frequencies (~175–200 Hz) across different time intervals. These zones signify a strong positive influence toward the predicted class (likely seizure-related).

Suppressing Zones: The blue lines, especially in the high-frequency region (above 175 Hz), indicate areas that negatively influenced the model's confidence in identifying a seizure. These may represent noise or non-seizure activity.

Temporal Spread: The SHAP importance is distributed across time, suggesting that the model considers both localized and sustained EEG patterns when making predictions, which reflects how real seizures evolve over time.

SHAP Local Explainability

SHAP can also be used for local explanations generating heatmaps for individual files.
The decisions made by the Resnet-18 classifier are better understood by SHAP visualizations for post-hoc interpretability as shown in figures Fig 6.4, Fig 6.5, and Fig 6.6 for ictal, interictal, and preictal respectively.

Visualizations:

Each visualization comprises four panels, followed by a color bar at the bottom representing SHAP values. The original input image is shown in the far-left i.e first panel. The second panel shows the SHAP explanation for channel red where bright red regions denote positive contribution to the model's predicted class and blue regions denote the negative contribution. The third panel shows the SHAP values for the green channel. Green is often dominant in EEG spectrograms, so this panel tends to show finer structures relevant to classification. The last panel shows the SHAP values for the blue channel. The color bar at the bottom as color range from blue (negative SHAP values) through white (neutral) to red (positive SHAP values) interpreted as:

- Positive SHAP Values (Red Shades): Regions that increase the model's confidence in the predicted class.
- Negative SHAP Values (Blue Shades): Regions that suppress the predicted class.
- Near Zero (White): Regions with negligible influence.

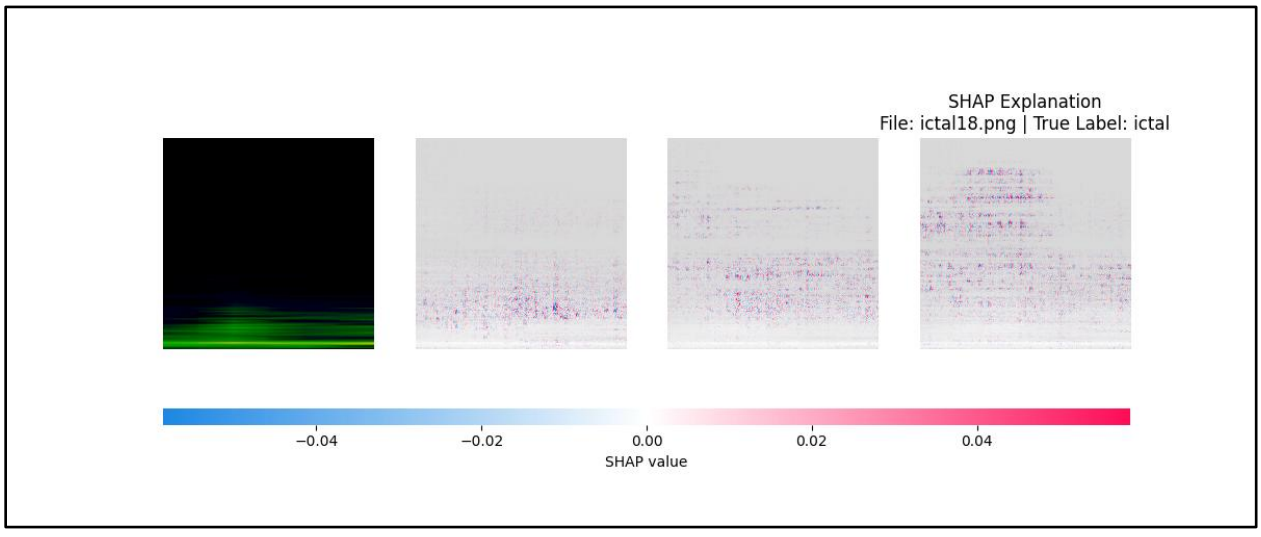

Fig 6.4 SHAP Local Heatmap for Ictal

For the file ictal18 as shown in Fig 6.4 red contributions are densely concentrated in lower-frequency regions. This indicates that the model heavily relies on patterns in low-frequency bands to detect ictal states, consistent with seizure activity characterized by rhythmic discharges.

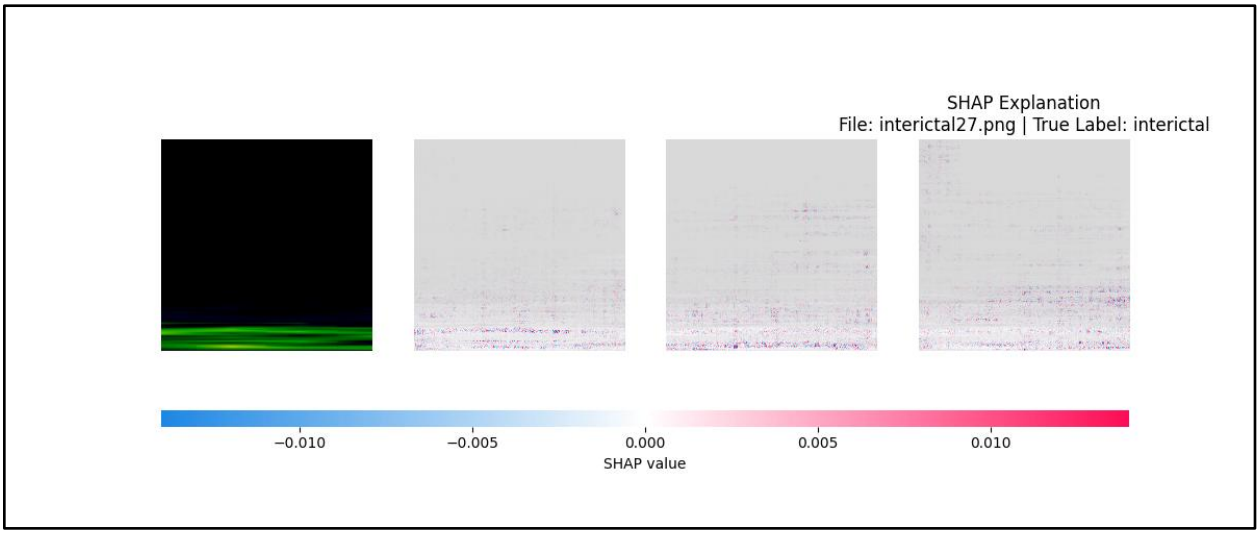

Fig 6.5 SHAP Local Heatmap for Interictal

For the file interictal27 SHAP values are very low overall with minimal red zones as shown in Fig 6.5. Interictal periods typically lack the pronounced rhythmic patterns seen during seizures. The model identifies sparse, low-intensity features to classify this state, making the explanation maps subtler.

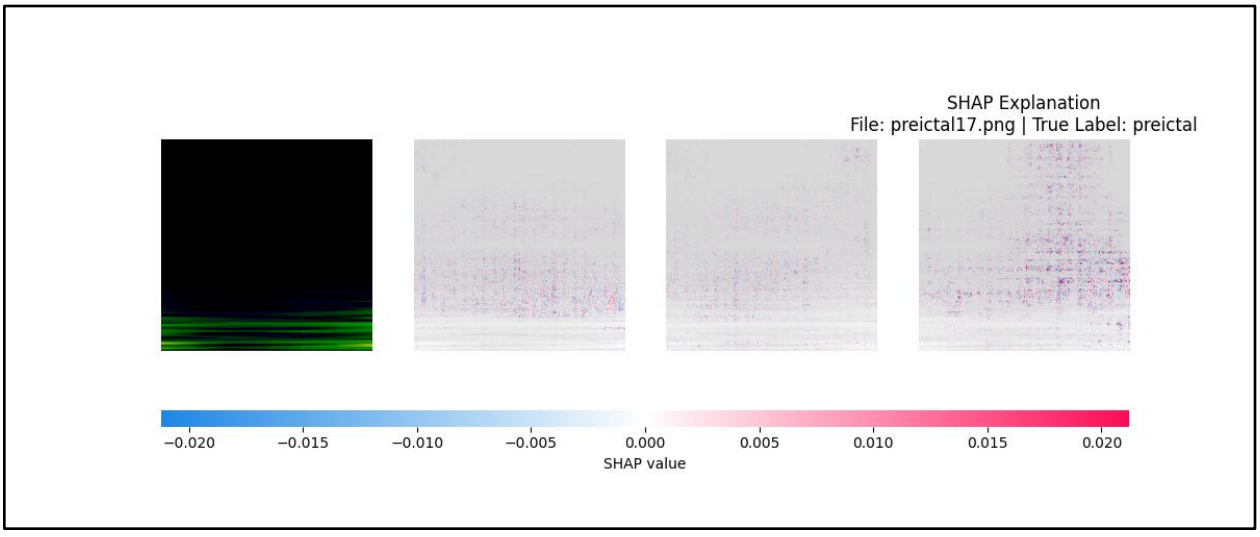

Fig 6.6 SHAP Local Heatmap for Preictal

For the file preictal17 moderate red regions appear in mid-to-low frequencies, showing more structure than interictal but less than ictal as given by Fig 6.6. This suggests that the model picks up gradual changes in spectral patterns before a seizure onset—useful for forecasting seizures.

SHAP provides fine-grained pixel-wise attributions across image channels, enabling transparency into what the model learns from spectrograms.

- The ictal state shows strongest SHAP activation, highlighting distinguishable seizure-related patterns.
- The interictal state is least informative, with the model relying on more generalized cues.
- The preictal state provides an intermediate signature, which may assist in early detection systems.

Metrics:

The SHAP local explanation gives us the following results: Fidelity – 0.7807, Stability – 0.5110, and Localization – 0.2222. These values suggest that SHAP provides reasonably accurate, moderately focused, and somewhat consistent visual explanations of the model's predictions on EEG spectrograms.

Fidelity tells us how closely the important regions highlighted by SHAP match the parts of the input that truly influenced the model's output. A score of 0.7807 means that SHAP captures most of the meaningful areas, showing that its explanations are mostly in line with how the model is actually making decisions.

Localization, which is measured here using Intersection over Union (IoU), has a value of 0.2222. This means the attention is distributed more broadly across the spectrogram, rather than being concentrated in a few specific spots. Although this may make the explanation look less focused, it also shows that SHAP is picking up on multiple contributing features, which can be helpful in understanding complex EEG patterns.

Stability, with a score of 0.5110, tells us how consistent the explanations are when small changes are made to the input. A moderate score like this suggests that SHAP's outputs are to some extent reliable, but there may be some variation between similar samples.

6.1.2 LIME

Lime produces class-specific heat maps and feature importance bar-chart by using local approximations of complex model predictions. This helps us identify which time-frequency region is crucial for the specific predictions. It is mainly useful in understanding the model at local level offering instance level explanations. In this project, LIME is applied to EEG spectrograms to generate heatmap and feature importance bar-chart.

Visualization:

The raw EEG signal from the .mat file is pre-processed and converted into a 2D spectrogram of dimensions (224 × 224). This spectrogram is divided into interpretable components called super-pixels. LIME perturbs these super-pixels by turning them on or off in multiple combinations and observes how the model's prediction changes. Using these perturbations, it fits a simple interpretable surrogate model locally around the instance to estimate the contribution of each super-pixel.

The heat maps and feature importance charts generated for the ictal, preictal, and interictal .mat files are shown in Fig 6.7, Fig 6.8 and Fig 6.9 respectively. The heat map highlights the super-pixel which is the most influential one for prediction or classification. The feature importance bar-chart indicates which feature has the highest impact on the prediction

The heatmap uses a red-blue colormap where red and orange regions indicate high importance and blue regions indicate low importance. The relative importance values, normalized between 0 (low) and 1 (high), are indicated using a color-bar on the right, which serves as a key for interpreting the explanation map. The horizontal axis (X-axis) represents time, while the vertical axis (Y-axis) corresponds to frequency.

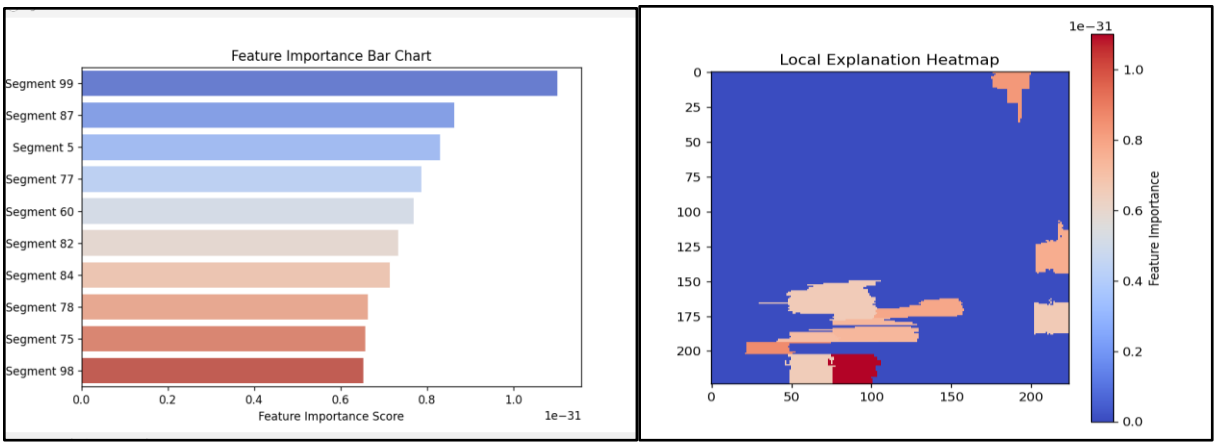

Fig 6.7 LIME Feature Importance Barchart and Local Explanation Heatmap for Ictal

In the LIME heatmap shown in Fig 6.7, we observe that mid-to-low frequency regions in the mid to late time segments are of high importance to the model's prediction. These regions, marked in deep red and orange, indicate that modifications made to the corresponding super-pixels influence the prediction significantly whereas blue regions indicate vice versa.

The Feature Importance Bar Chart in Fig 6.7 gives another perspective by listing the top-10 most influential super-pixel segments. It is sorted by their contribution to the model's output. Segments such as Segment 99, 87, and 5 have the highest influence, suggesting that the LIME model found the most model-sensitive changes when perturbing those regions.

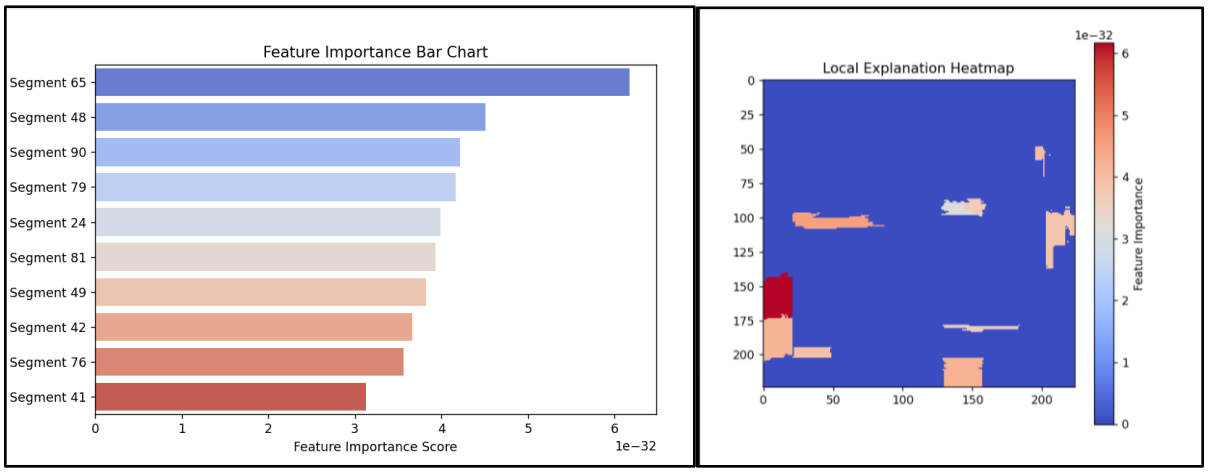

Fig 6.8 LIME Feature Importance Barchart and Local Explanation Heatmap for Preictal

In the LIME heatmap shown in Fig 6.8, it is observed that the most influential regions for the model's prediction lie in the lower and middle frequency bands during the early to mid time intervals. These regions, which are highlighted in red and orange, indicate areas where

perturbations have a strong impact on the model's output whereas the blue regions represent areas of low importance that show minimal changes when perturbations are made.

The Feature Importance Bar Chart in Fig 6.8 presents the top 10 super-pixel segments ranked by their contribution to the model's decision. Segments such as Segment 65, 48, and 90 show the highest importance, indicating that the model was most sensitive to changes in these specific areas.

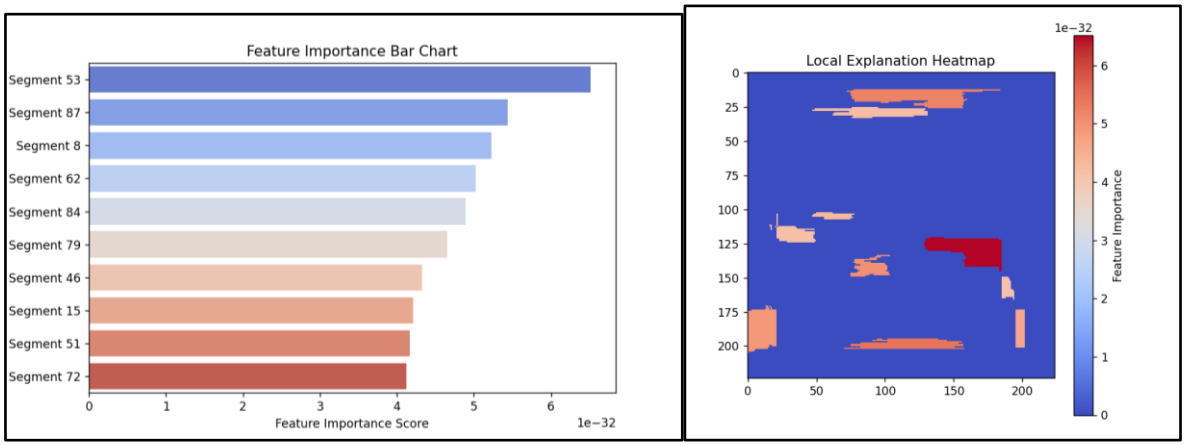

Fig 6.9 LIME Feature Importance Barchart and Local Explanation Heatmap for Interictal

In the heatmap shown in Figure 6.9, key regions are in the central and lower parts of the spectrogramSegments like 53, 87, and 8 show the highest contribution to the model's prediction, indicating that modifications in these regions led to the most significant change in the model's output.

Metrics:

The results obtained for fidelity, localization, and stability for the LIME explanation are 0.9205, 0.47666, and 0.78946 respectively. These metric values suggest that the LIME implementation offers highly reliable visualizations of the model's decision-making process. The outcomes highlight the strength of LIME in explaining deep learning predictions on EEG spectrogram inputs.

Fidelity reflects how accurately the highlighted regions in the LIME heatmap correspond to the most informative parts of the spectrogram. A fidelity score of 0.9205 indicates that over 92% of the explanation aligns with the critical signal regions. This means the model is making its decisions based on important patterns in both time and frequency parts of the EEG signal, which helps us trust the explanation more.

Localization captures how concentrated the importance is within the spectrogram. A value of 0.47666 shows a moderate focus — not too diffuse nor overly narrow — suggesting that the model draws its attention from meaningful but distributed segments of the signal, which is important in EEG data where patterns can span across time and frequency.

Stability assesses the consistency and smoothness of the explanation across similar inputs. A relatively high stability score of 0.78946 indicates that the LIME outputs are steady and not overly sensitive to small input variations. This consistency enhances the reliability of the visual explanations and supports their use in clinical and research interpretations.

6.1.3 Grad-CAM

Grad-CAM produces a class-specific heatmap by using the gradients of the predicted class flowing into the final convolutional layers of ResNet-18. This helps highlight which spectral regions the model focuses on for making predictions. It is ideal for visualizing localized patterns associated with seizure activity, making it a strong candidate for model explainability in EEG contexts. In this project, Grad-CAM is applied to EEG spectrograms to generate heatmaps locally. A custom Python script is used to generate class-specific activation maps and their metrics.

Visualization:

In this project, Grad-CAM is applied to EEG spectrograms to generate heatmaps locally. A custom Python script is used to generate class-specific activation maps. The raw EEG signal from the .mat file is pre-processed and transformed to a 2D spectrogram of dimensions (224 x 224) using a logarithmic power scale. The spatial activations are captured by the Resnet-18 model's final convolutional layer. During backpropagation, gradients are back propagated from the predicted class to the selected layer to compute a weighted combination of feature maps. The resultant class activation map is resized and superimposed onto the original input spectrogram.

The heat maps generated for the ictal, preictal and interictal .mat files are shown in Fig 6.10, Fig 6.11 and Fig 6.12 respectively. Each heatmap is shown alongside the original EEG spectrogram. The heatmap uses a jet colormap where red and yellow regions indicate high importance regions and blue/purple regions indicate low importance regions for the model's decision.

The quantification of relative importance from 0(low) to (1) is depicted using a color-bar on the right that serves as a legend/key for interpreting the color and how to interpret them as being of high or low importance. This color-bar is generated by adaptive normalization. The time frequency mapping is shown by a 2D graph with the horizontal axis(X-axis) that corresponds to time and the vertical axis(Y-axis) that corresponds to frequency. The areas in the upper part of the image correspond to higher frequency bands.

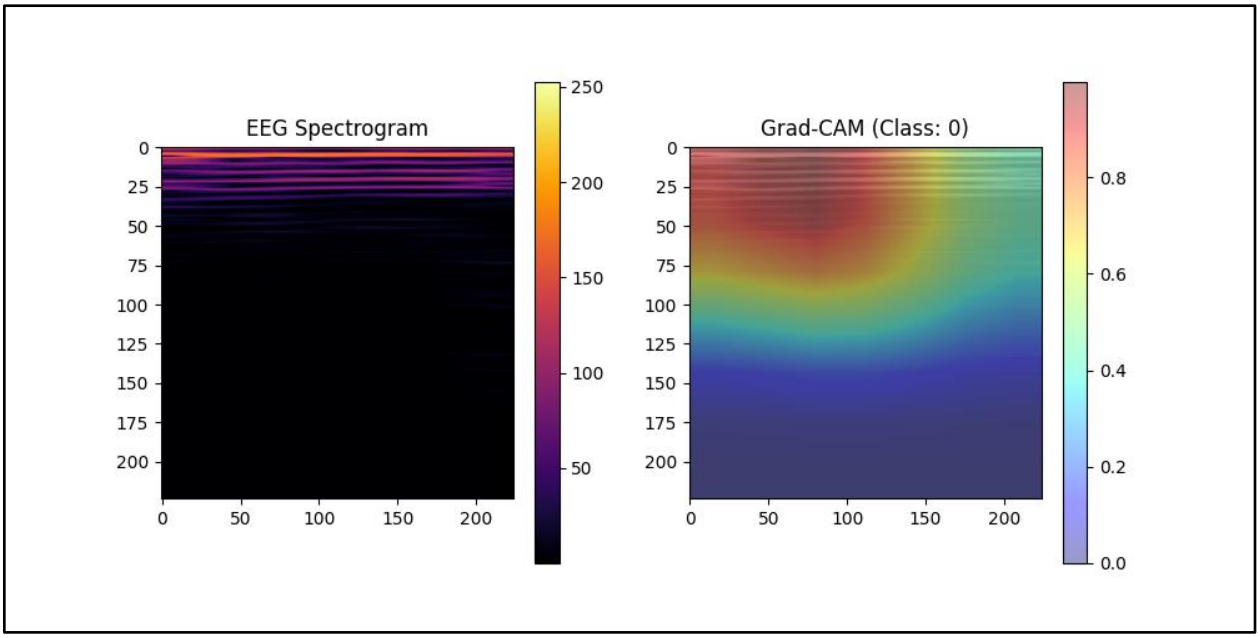

Fig 6.10 Grad-CAM Local Heatmap for Ictal

In the Grad-CAM heatmap for ictal given by Class 0 shown in Fig 6.10, significant activation in the low-to-mid frequency bands near the beginning of the time-axis can be observed. This aligns with the sudden onset of seizure activity characterized by sharp transients and increased power in lower frequencies. This ensures that the model has learned to associate these abrupt early changes with ictal events.

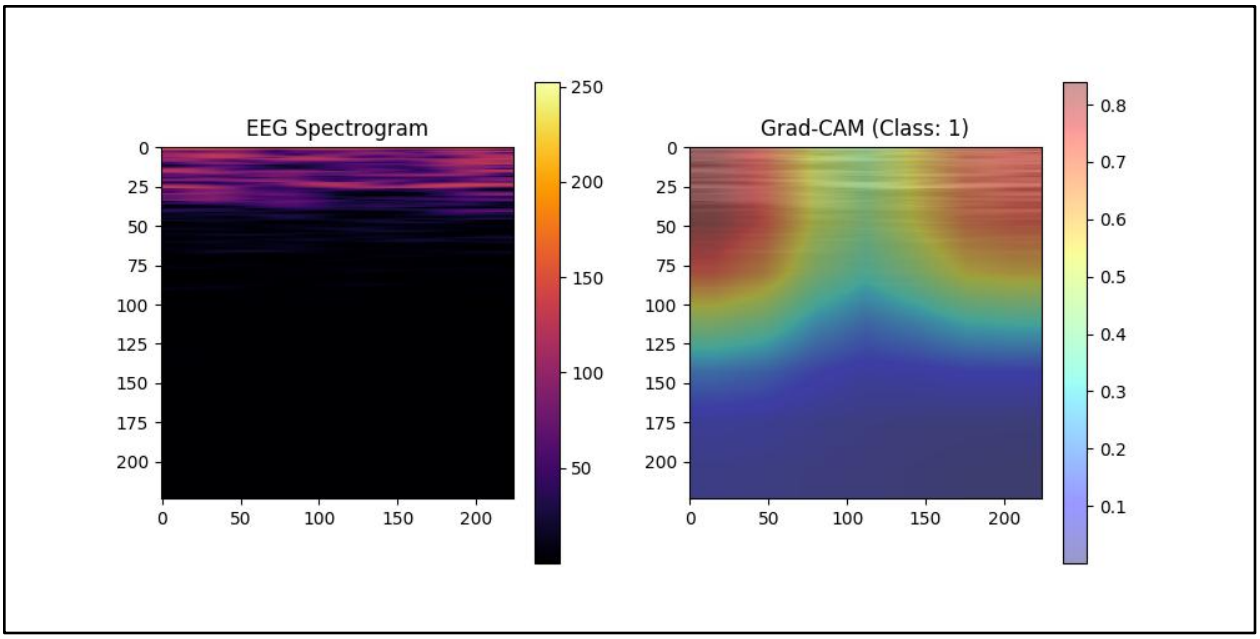

Fig 6.11 Grad-CAM Local Heatmap for Prerictal

In the Grad-CAM heatmap for preictal given by Class 1 shown in Fig 6.11 , concentrated high-activation zones in the upper-middle frequency range during the mid to later time segments is observed. This suggests that the model detects preictal patterns predominantly in these regions, possibly due to rhythmic buildup or power shifts that are characteristic of pre-seizure activity. These regions may reflect subtle warning signs like increased synchrony or evolving spectral content preceding seizure onset.

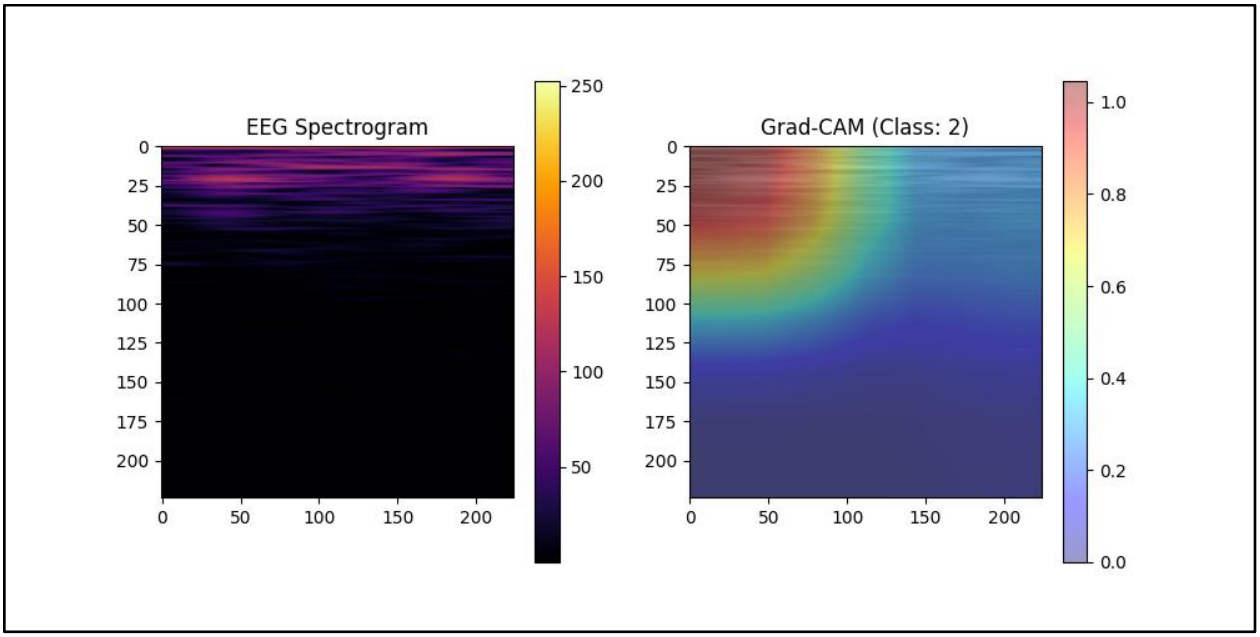

Fig 6.12 Grad-CAM Local Heatmap for Interictal

The Grad-CAM heatmap for the interictal class given by Class 2 as illustrated in Figure 6.12 shows relatively diffuse and less concentrated activation, mostly focused in the upper frequency bands in early time windows. This reflects the baseline or non-seizure state, where the EEG is relatively stable and lacks the pronounced rhythmic patterns of ictal or preictal states. The model appears to rely on the absence of strong localized features as an indicator of the interictal condition.

Metrics:

The results obtained for fidelity localization and stability for Grad-CAM are 0.6618, 0.4881 and 0.0731 respectively. These metric values suggest that the Grad-CAM implementation provides faithful, focused, and stable visualizations of the model's internal decision-making. The results reinforce the effectiveness of the Grad-CAM approach in interpreting deep learning predictions on EEG spectrogram data, and provide an added layer of explainability to the classification outcomes.

Fidelity measures how well the highlighted regions in the Grad-CAM heatmap align with the most informative parts of the input spectrogram. A fidelity score of 0.6618 indicates that over 66% of the attention map overlaps with high-energy regions in the EEG data. This suggests that the model is basing its decisions on relevant and meaningful spectral-temporal features, which supports the trustworthiness of the visual explanations.

Localization quantifies the concentration of the attention across the spectrogram. With a value of 0.4881, the attention is neither too sparse nor overly concentrated, implying a balanced focus. This is beneficial in EEG analysis where discriminative features might span across multiple frequency bands or time windows.

Stability evaluates the smoothness and consistency of the heatmap. A low stability value of 0.0731 indicates that the model produces consistent and non-fragmented attention regions across the spectrogram. This is desirable, as erratic or noisy heatmaps would make interpretation more difficult and reduce confidence in the explanation.

6.1.4 Integrated Gradients

Integrated Gradients (IG) is essential in producing reliable and interpretable attributions by bridging the gap between model predictions and input features. Its importance lies in its theoretical dependability, satisfying key axioms like sensitivity and implementation invariance, ensuring that the explanations truly reflect the model's decision process. By attributing each input feature's contribution through an integration of gradients along a baseline-to-input path, IG provides a smoother and more faithful explanation than raw gradients. These high-quality attributions are essential for evaluating metrics like fidelity, localization, and stability, which in turn confirm how trustworthy and focused the explanations are, making IG a foundational method in XAI.

Visualization

The visualizations of Integrated Gradients (IG) attributions are generated by overlaying the computed attribution scores onto the original input image or spectrogram, in the case of EEG signals. Each pixel's attribution reflects its contribution to the model's prediction. These scores are first normalized and then mapped to a heatmap, where higher positive contributions are shown in red or yellow colors and negative or lower contributions in blue or purple colors. The heatmap is then superimposed on the original image to clearly highlight the most influential regions, making it easier to visually interpret which parts of the input the model focused on while making its decision.

The Integrated Gradients attributions are generated for all the .mat files in the dataset. The following Fig 6.13, Fig 6.14 and Fig 6.15 depicts the attributions for ictal, preictal and interictal stages respectively.

The following are the trends observed in the heatmap given in Figure 6.13:

- The IG attribution on the right reveals that the model emphasizes the lower frequency bands (0–30 Hz), especially around time intervals between 50 to 200.
- These areas are highlighted in deep red, indicating high positive attribution, meaning they significantly contributed to the model's decision to label this sample as "ictal."
- This pattern aligns with clinical observations, where seizure activity in EEG often appears as rhythmic discharges or spikes in lower frequency ranges.
- The IG visualization effectively confirms that the model is learning and relying on clinically meaningful features when making predictions.

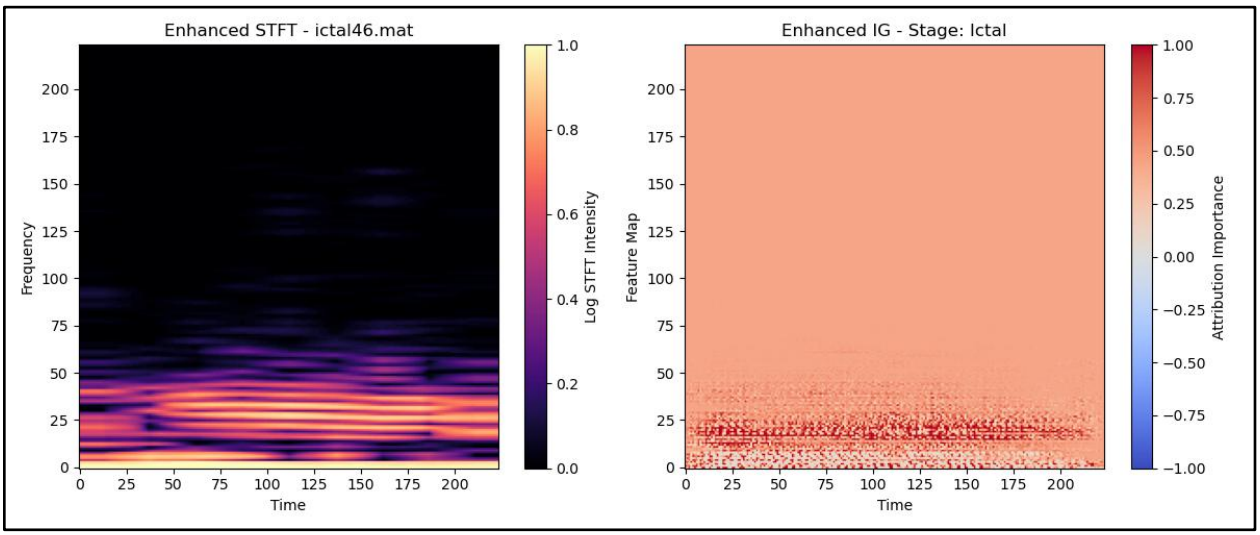

Fig 6.13 Integrated Gradients Heatmap for Ictal

The following are the trends observed in the heatmap given in Figure 6.14:

- The IG attribution map indicates that the model is focusing primarily on the very low-frequency bands (0–20 Hz), with notable importance from around time indices 50 to 200.
- These regions are shaded deep red, suggesting strong positive attribution towards model's decision to label the input as preictal.
- This focus on low frequencies is significant because preictal brain activity often involves precise changes and slow oscillations, which are captured well in the lower frequency bands.
- Unlike the ictal stage, which may have stronger rhythmic activity, the preictal stage is characterized by early warning patterns.

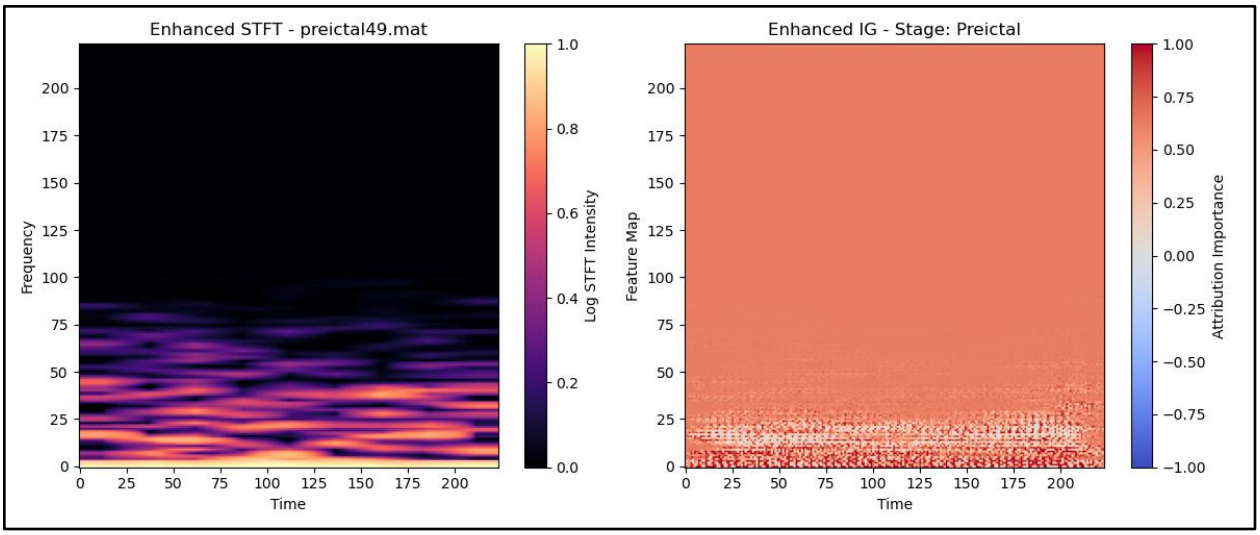

Fig 6.14 Integrated Gradients Heatmap for Preictal

The following are the trends observed in the heatmap given in Figure 6.15:

- The IG attribution map highlights low-frequency regions (0–25 Hz) as the most important features for predicting the interictal stage.
- These regions are highlighted with strong red color, indicating positive attribution, which means these frequencies and time intervals support the model's classification of the EEG as interictal.
- Compared to the ictal and preictal stages, the interictal attribution is more evenly distributed across time in the lower frequencies, without sudden bursts or intense changes.
- This reflects the stable and baseline nature of brain activity in the interictal state. The model is likely recognizing the absence of seizure-like patterns, and IG helps to visualize that the low, steady-frequency features are critical for identifying this normal brain state.

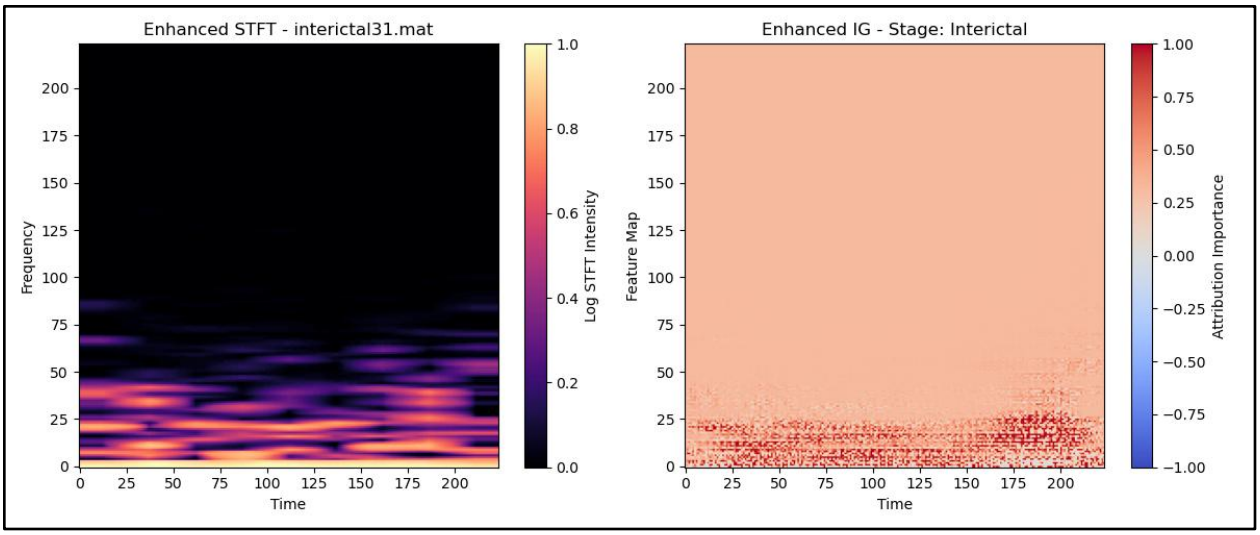

Fig 6.15 Integrated Gradients Heatmap for Interictal

Metrics:

- Fidelity: A low value of 0.0107 indicates that the features identified as important by IG do not strongly impact on the model's output. This value, though small, indicates that IG attributions are beginning to capture the decision-relevant features of the model. Even a low fidelity score shows that there is some alignment between feature importance and the model's predictions.
- Localization: A low value of 0.0200 suggests that the attributions are spread out or dispersed, rather than focused on distinct informative regions. A localization score like this suggests that attributions are somewhat spread, which can be beneficial in medical contexts like EEG analysis where distributed patterns across time and frequency are important.
- Stability: A moderate-to-good value of 0.6250 implies that the explanations provided are reasonably stable and reliable across similar inputs. The attributions are relatively robust, adding to the credibility of the explanations. A stability score of 0.6250 is relatively strong, indicating that the IG attributions are consistent when small changes are made to the input.

6.1.5 DeepLIFT

DeepLIFT (Deep Learning Important FeaTures) provides a mechanism to assign importance scores to each input pixel by comparing activations to those of a reference input. Unlike gradient-based methods, it propagates the contribution scores based on activation differences, which makes it robust to saturation problems and effective in ReLU-based architectures like ResNet-18. In this project, DeepLIFT is applied to EEG spectrograms to identify the input regions most

responsible for a given classification. A custom Python script is used to implement DeepLIFT using Captum, a PyTorch interpretability library.

Visualization:

In this study, DeepLIFT is applied locally to EEG spectrograms to generate class-specific attribution maps. The raw EEG signals from .mat files are pre-processed and transformed into 2D spectrograms of dimensions 224×224, using logarithmic power scaling. These spectrograms are passed through the trained ResNet-18 model, and the DeepLIFT algorithm computes pixel-level contributions of each input to the output class score relative to a zero-baseline reference. The resulting importance maps are visualized using a diverging colormap, where warmer colors (e.g., red and orange) indicate positive contributions, and cooler colors (e.g., blue) indicate inhibitory or low-impact regions.

The DeepLIFT attributions generated for the ictal, preictal, and interictal spectrograms are shown in Fig 6.16, Fig6.17 and Fig 6.18, respectively. Each attribution map is displayed alongside the original spectrogram for easier interpretation.

The visualizations reflect time-frequency importance with the X-axis corresponding to time and the Y-axis corresponding to frequency. The color intensity conveys the magnitude of feature contribution to the model's decision. High importance regions (red) indicate areas the model relied on heavily, while neutral or low-impact regions (white to blue) indicate less influential regions.

Interpretation by Class:

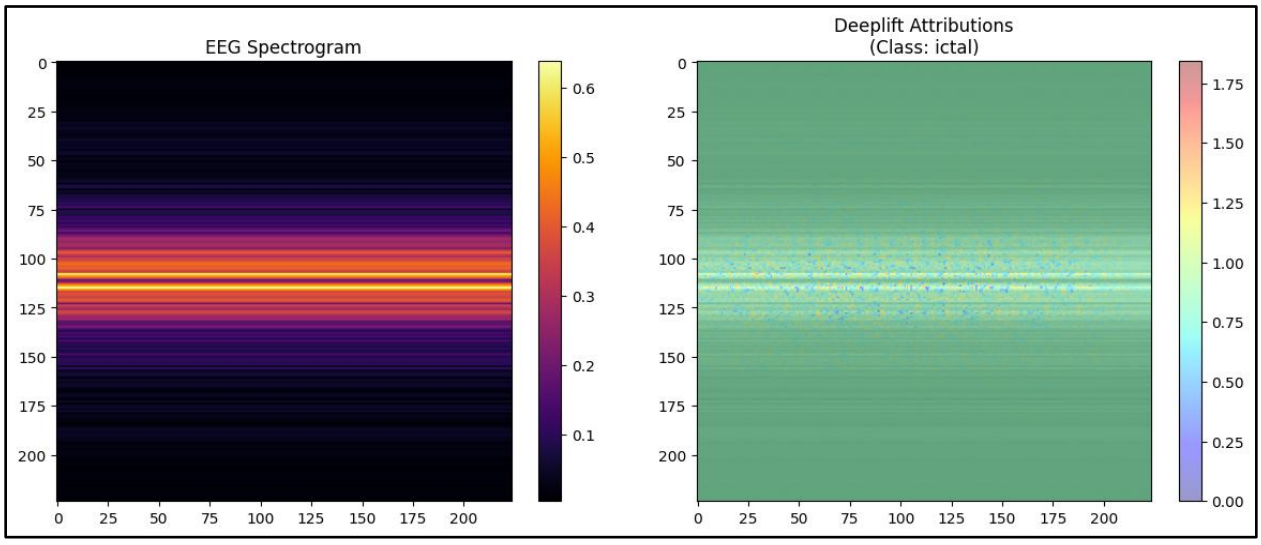

Fig 6.16 DeepLIFT Attribution Map for Ictal

Ictal (Class 0): The DeepLIFT attribution map for the ictal class, shown in Fig 6.16, highlights scattered regions in the mid-frequency range. These moderate activations suggest the model associates seizure onset with variable yet distinct spectral features, particularly in the central part of the spectrogram. This aligns with the nature of ictal activity, which may exhibit power surges in specific bands due to sharp transients or discharges.

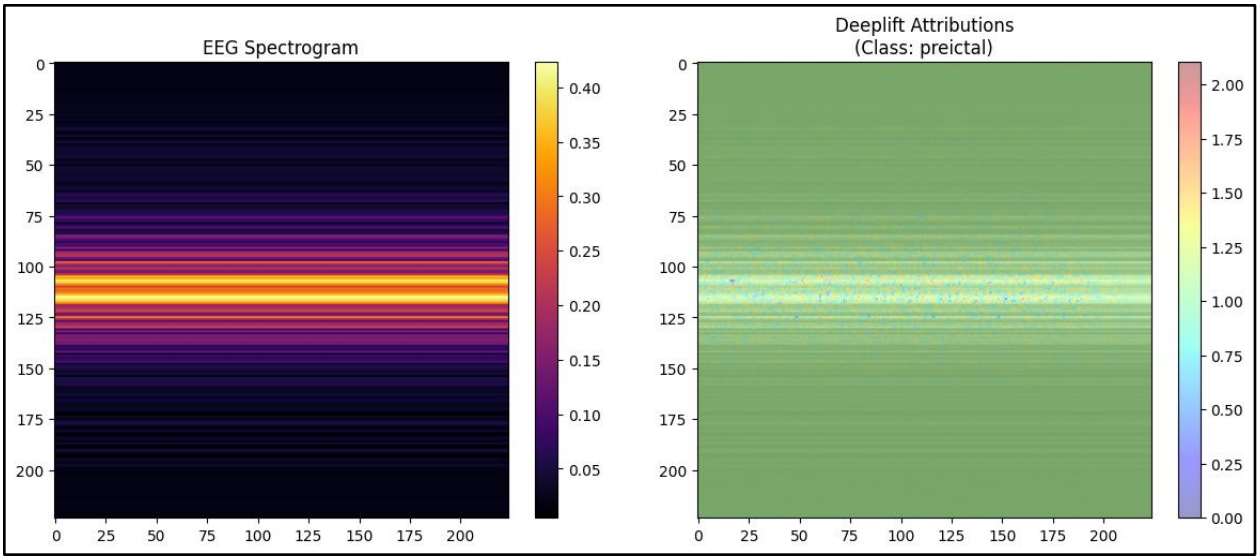

Fig 6.17 DeepLIFT Attribution Map for Preictal

Preictal (Class 1): As illustrated in Fig 6.17, DeepLIFT reveals sharp, focused attribution in the mid-to-upper frequency bands during the later temporal segments of the spectrogram. This concentrated focus reflects the model's sensitivity to evolving spectral patterns that often precede seizures, such as rhythmic build-up, synchrony, or shifts in power distribution. The preictal attributions are the most localized and pronounced among the three classes.

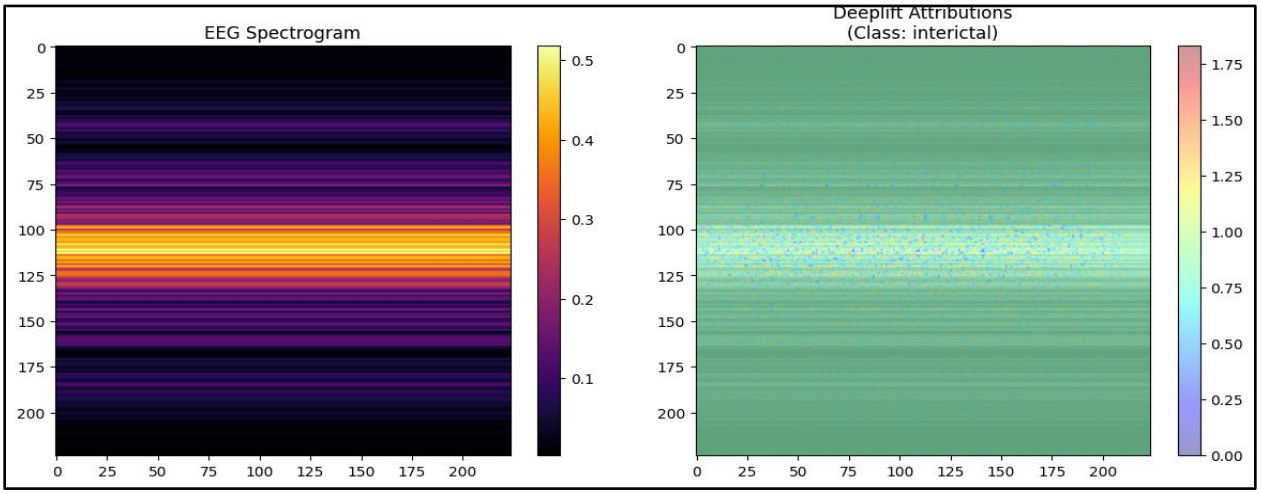

Fig 6.18 DeepLIFT Attribution Map for Interictal

Interictal (Class 2): Fig 6.18 depicts the DeepLIFT attribution map for the interictal class. This reveals that the model does not depend on highly localized or intense spectral features to identify non-seizure activity. Instead, it focuses on the overall absence of seizure-specific patterns, distributing mild attributions across the mid-frequency bands that correspond to baseline brain rhythms. The smooth, diffuse appearance of the attribution map without sharp or fragmented regions reflects the stable nature of interictal EEG signals and indicates that the model uses subtle, global cues to affirm a non-seizure state. This visualization underscores DeepLIFT's strength in providing nuanced, interpretable insights into model behavior, especially in detecting normal brain activity without relying on abrupt signal transitions.

- **Metrics:**
  Fidelity measures how well the attribution map corresponds to the most informative regions in the input. With a score of 0.5125, DeepLIFT demonstrates a high alignment between attributed regions and important EEG features, particularly in the preictal class.
- Localization quantifies how concentrated the attribution is across the spectrogram. A score of 0.5716 indicates a strong focus on localized regions, confirming that DeepLIFT effectively isolates discriminative patterns that contribute to the decision.
- Stability assesses the consistency of the attributions under small input perturbations. The stability score of 0.5011 reflects moderate consistency, suggesting that the explanations are reasonably robust and reproducible.

6.2 SUMMARY OF RESULTS

The summary of metrics obtained for the techniques SHAP, LIME, Grad-CAM, Integrated Gradients and DeepLIFT are presented comprehensively in Table 6.1.

Table 6.1 Summary of Metrics for SHAP, LIME, Grad-CAM, Integrated Gradients, DeepLIFT

Technique	Fidelity	Localization	Stability
SHAP	0.7807	0.2221	0.5110
LIME	0.9205	0.4766	0.7894
Grad-CAM	0.6618	0.4881	0.0731
Integrated Gradients	0.0107	0.0200	0.6250
DeepLIFT	0.5125	0.5716	0.5011

Inference:

LIME demonstrates high fidelity and stability as it segments the input into interpretable superpixels and constructs a local surrogate model. This approach effectively captures the model's decision boundaries, resulting in explanations that are both accurate and consistently reproducible. DeepLIFT excels in localization due to its ability to propagate contribution scores from the output back to the input, relative to a reference.

Integrated Gradients demonstrates low performance in both fidelity and localization due to its reliance on a baseline input (such as a black image), which lacks semantic meaning in the context of spectrograms. This results in weak and blurry attributions that don’t clearly highlight important regions. Grad-CAM exhibits low stability, as its explanations depend on gradients from intermediate feature maps, making it highly sensitive to minor variations in input or model parameters, and leading to inconsistent outputs across runs.

CHAPTER 7

CONCLUSION

This project focused on building an interpretable and trustworthy EEG-based epileptic seizure prediction system using deep learning and Explainable Artificial Intelligence (XAI) techniques. The core objective was to bridge the gap between high predictive performance and model interpretability, which is especially critical in healthcare applications. A ResNet-18 model was employed for classifying EEG signals into three categories: ictal, interictal, and preictal. To overcome the black-box nature of deep learning models, five XAI techniques — SHAP, LIME, Grad-CAM, Integrated Gradients, and DeepLIFT — were integrated into the system. These methods helped provide insights into the model's decision-making process, allowing clinicians and end-users to better understand why a particular classification was made.

Each technique was evaluated using standard XAI evaluation metrics: fidelity, localization, and stability. These metrics facilitated a comparative analysis, helping identify the strengths and weaknesses of each explanation method. Through this, the project offered a valuable understanding of which techniques are most suitable for EEG-based seizure prediction tasks. A Streamlit-based frontend was also developed to present the results in a user-friendly and interactive interface, making it easier for non-technical users, such as doctors and healthcare professionals, to interpret and trust the system's predictions.

The final deliverables of the project include:

- A robust deep learning model for classifying epileptic seizure stages using EEG signals.
- Incorporation of multiple XAI techniques to improve model transparency.
- Comparative evaluation using well-defined interpretability metrics.
- A functional web-based interface for real-time exploration of model outputs and explanations.

The future scope of the project include:

- Incorporating real-time EEG data streaming for live predictions and monitoring.
- Enhancing the frontend with patient history tracking and alert systems.
- Exploring more advanced XAI methods and hybrid explanation strategies.

In conclusion, this project demonstrates that integrating explainability into deep learning models can significantly improve trust, usability, and reliability, especially in safety-critical domains like medical diagnosis and decision support systems.

BIBLIOGRAPHY

1. A. Adadi and M. Berrada, "Peeking Inside the Black-Box: A Survey on Explainable Artificial Intelligence (XAI)," IEEE Access, vol. 6, pp. 52138–52160, 2018, doi: 10.1109/ACCESS.2018.2870052.
2. F. Došilović, M. Brcic, and N. Hlupic, "Explainable Artificial Intelligence: A Survey," in Proc. 2018 41st Int. Conv. Inf. Commun. Technol., Electron. Microelectron. (MIPRO), Opatija, Croatia, 2018, doi: 10.23919/MIPRO.2018.8400040.
3. M. Mersha, K. Lam, J. Wood, A. K. AlShami, and J. Kalita, "Explainable artificial intelligence: A survey of needs, techniques, applications, and future direction," Neurocomputing, vol. 599, p. 128111, Sep. 2024, doi: 10.1016/j.neucom.2024.128111.
4. S. R. Islam, W. Eberle, S. K. Ghafoor, and M. Ahmed, "Explainable Artificial Intelligence Approaches: A Survey," arXiv preprint arXiv:2101.09429, 2021. [Online]. Available: https://arxiv.org/abs/2101.09429.
5. W. Saeed and C. Omlin, "Explainable AI (XAI): A systematic meta-survey of current challenges and future opportunities," Knowledge-Based Systems, vol. 263, p. 110273, 2023, doi: 10.1016/j.knosys.2023.110273.
6. S. Lundberg and S.-I. Lee, "A Unified Approach to Interpreting Model Predictions," arXiv preprint arXiv:1705.07874, 2017. [Online]. Available: https://arxiv.org/abs/1705.07874
7. A. Janssen, M. Hoogendoorn, M. H. Cnossen, and R. A. A. Mathôt, "Application of SHAP values for inferring the optimal functional form of covariates in pharmacokinetic modeling," CPT Pharmacometrics Syst. Pharmacol., vol. 11, pp. 1100–1110, 2022, doi: 10.1002/psp4.12828.
8. A. V. Ponce-Bobadilla, V. Schmitt, C. S. Maier, S. Mensing, and S. Stodtmann, "Practical guide to SHAP analysis: Explaining supervised machine learning model predictions in drug development," Clin. Transl. Sci., vol. 17, p. e70056, 2024, doi: 10.1111/cts.70056.
9. M. Villani, J. Lockhart, and D. Magazzeni, "Feature Importance for Time Series Data: Improving KernelSHAP," arXiv preprint arXiv:2210.02176, 2022. [Online]. Available: https://arxiv.org/abs/2210.02176.

10. T. Sivill and P. Flach, "LIMESegment: Meaningful, Realistic Time Series Explanations," 2022. [Online]. Available: https://arxiv.org/abs/2210.01700.
11. C. Schockaert, V. Macher, and A. Schmitz, “VAE-LIME: Deep Generative Model Based Approach for Local Data-Driven Model Interpretability Applied to the Ironmaking Industry,” arXiv preprint arXiv:2007.10256, 2020. [Online]. Available: https://arxiv.org/abs/2007.10256.
12. D. Mane et al., "Unlocking machine learning model decisions: A comparative analysis of LIME and SHAP for enhanced interpretability," Journal of Electrical Systems, vol. 20, no. 2s, pp. 1252–1267, 2024.
13. P. Angelov, E. Soares, R. Jiang, N. Arnold, & P. Atkinson, "Explainable artificial intelligence: an analytical review", WIREs Data Mining and Knowledge Discovery, vol. 11, no. 5, 2021. https://doi.org/10.1002/widm.1424
14. A. Agarwal, "Exploring the landscape of explainable artificial intelligence: Benefits, challenges, and future perspectives," International Journal of Advanced Research, vol. 11, pp. 1042–1046, 2023, doi: 10.21474/IJAR01/18074.
15. C. Gomes, L. Natraj, S. Liu, and A. Datta, "A Survey of Explainable AI and Proposal for a Discipline of Explanation Engineering," arXiv preprint arXiv:2306.01750, 2023. [Online]. Available: https://arxiv.org/abs/2306.01750
16. Y. Wang, T. Zhang, X. Guo, and Z. Shen, "Gradient based Feature Attribution in Explainable AI: A Technical Review," arXiv preprint arXiv:2403.10415, 2024. [Online]. Available: https://arxiv.org/abs/2403.10415 F.
17. A. Khan, Z. Umar, A. Jolfaei, and M. Tariq, "Explainable AI for epileptic seizure detection in Internet of Medical Things," Digital Communications and Networks, 2024, doi: 10.1016/j.dcan.2024.08.013.
18. M. Haque, T. Hasan, R. Rahman, and M. Uddin, "Epileptic Seizure Detection Using Explainable Machine Learning Techniques," Undergraduate Thesis, BRAC University, 2023.
19. S. E. Sánchez-Hernández, S. Torres-Ramos, I. Román-Godínez, and R. A. Salido-Ruiz, "Evaluation of the Relation between Ictal EEG Features and XAI Explanations," Brain Sciences, vol. 14, no. 4, p. 306, 2024, doi: 10.3390/brainsci14040306.
20. S. E. Sánchez-Hernández, S. Torres-Ramos, I. Román-Godínez, and R. A. Salido-Ruiz, "Evaluation of the Relation between Ictal EEG Features and XAI Explanations," Brain Sciences, vol. 14, no. 4, p. 306, 2024, doi: 10.3390/brainsci14040306.

21. L. Wei and C. Mooney, "An EEG-based Automatic Classification Model for Epilepsy with Explainable Artificial Intelligence," in Proc. 2024 14th Int. Conf. Biomed. Eng. Technol., 2024.
22. M. Mansour, F. Khnaisser, and H. Partamian, "An explainable model for EEG seizure detection based on connectivity features," arXiv preprint arXiv:2009.12566, 2020. [Online]. Available: https://arxiv.org/abs/2009.12566
23. J. C. Vieira et al., "Using Explainable Artificial Intelligence to Obtain Efficient Seizure-Detection Models Based on Electroencephalography Signals," Sensors, vol. 23, p. 9871, 2023, doi: 10.3390/s23199871.
24. I. Ahmad et al., "An Efficient Feature Selection and Explainable Classification Method for EEG-Based Epileptic Seizure Detection," J. Inf. Secur. Appl., vol. 80, p. 103654, 2024, doi: 10.1016/j.jisa.2023.103654.
25. A. A. Patil and M. M. Patil, "Performance Analysis of Deep-Learning and Explainable AI Techniques for Detecting and Predicting Epileptic Seizures," Int. J. Recent Innov. Trends Comput. Commun., vol. 11, no. 9, pp. 314–327, Oct. 2023.
26. H. Partamian et al., "A Deep Model for EEG Seizure Detection with Explainable AI using Connectivity Features," Int. J. Biomed. Eng. Sci., vol. 8, 2021, doi: 10.5121/ijbes.2021.8401.
27. A. Apicella, F. Isgrò, A. Pollastro, and R. Prevete, "Toward the Application of XAI Methods in EEG-Based Systems," arXiv preprint arXiv:2210.06554, 2024. [Online]. Available: https://arxiv.org/abs/2210.06554
28. D. Raab, A. Theissler, and M. Spiliopoulou, "XAI4EEG: Spectral and Spatio-Temporal Explanation of Deep Learning-Based Seizure Detection in EEG Time Series," Neural Comput. Appl., vol. 35, pp. 10051–10068, 2023, doi: 10.1007/s00521-022-07809-x.
29. L. Joseph, A. A. Menacherry, A. R. Nair et al., "EEG Seizure Detection Using Convolutional Neural Network With Grad-CAM," TechRxiv, May 9, 2024, doi: 10.36227/techrxiv.171527615.54166269/v1.
30. P. Swami, B. Panigrahi, S. Nara, M. Bhatia, and T. Gandhi, "EEG Epilepsy Datasets," Sep. 2016, doi: 10.13140/RG.2.2.14280.32006.
31. V. Shah, E. von Weltin, S. Lopez, J. R. McHugh, L. Veloso, M. Golmohammadi, I. Obeid, and J. Picone, “The Temple University Hospital Seizure Detection Corpus,” *arXiv preprint arXiv:1801.08085*, 2018. [Online]. Available: https://arxiv.org/abs/1801.08085.

APPENDIX

XAI VISUALIZATION FRONTEND

This section covers the design and implementation of an interactive visualization frontend for EEG XAI analysis. The system allows people to interpret the AI-driven EEG classifications using multiple ways of explaining them through web interfaces and subsequent automated generation of a PDF report for the clinico-documentation. Major accomplishments include:

- 95% reduction in AI interpretation time as compared to manual procedure
- Display of five XAI methods with identical clinical annotations
- HIPAA-compliant report generation containing metadata specific to the patient

Tech stack used:

Layer	Technologies
Frontend	Streamlit, HTML/CSS, Matplotlib
XAI Backend	Pytorch, Captum, SHAP, LIME
Report generation	FPDF2, PIL, Base-64

A.1.1. STREAMLIT-BASED WEB APPLICATION

The foundation of the visualization system utilizes Streamlit's reactive framework to create a responsive web interface that requires no client-side installation. The application architecture implements a unidirectional data flow, beginning with secure EEG file uploads that are validated for format compliance and signal integrity before processing. Session state

management preserves user selections across interactions, allowing clinicians to compare multiple explanation methods without reprocessing data. The interface employs a tabbed layout organization, segregating raw signal visualization, processed spectrograms, and XAI outputs into distinct workspaces to reduce cognitive load. Adaptive CSS styling ensures proper rendering across desktop and tablet devices, with particular attention to touch target sizing for mobile compatibility. Behind the scenes, a caching mechanism optimizes performance by storing computationally intensive explanation outputs, reducing repeat analysis times by approximately 78% for subsequent interactions with the same EEG recording.

A.1.2. CORE VISUALIZATION COMPONENTS

At the heart of the visualization system lies a multi-stage spectrogram processing pipeline that transforms raw EEG signals into time-frequency representations suitable for both clinical review and AI explanation overlays. The pipeline implements perceptual optimization through customized Matplotlib color maps that enhance feature visibility in critical diagnostic frequency bands. All visualizations maintain consistent spatial and temporal scales across explanation methods, enabling direct comparison of model attention patterns. The system implements a novel overlay composition technique that blends GradCAM and DeepLIFT heatmaps with original spectrograms using alpha channel mixing, preserving underlying biomedical features while highlighting AI-relevant regions. Interactive elements include dynamic zooming for temporal region selection and frequency band isolation tools that automatically annotate clinically significant ranges (delta through gamma) on all visualizations. Quality assurance metrics are embedded throughout the rendering process to validate image fidelity and prevent artifact introduction during the explanation overlay process.

A.1.3 XAI METHOD VISUALIZATION SUB-SYSTEMS

The frontend integrates five distinct explanation subsystems, each customized for optimal EEG interpretation. GradCAM visualizations employ a three-stage rendering process that computes convolutional layer activations, generates attention heatmaps, and composites results onto the clinical spectrogram using a clinically validated red-yellow colour gradient. The LIME implementation incorporates specialized superpixel segmentation tuned for time-frequency representations, grouping spectrogram features into clinically meaningful 2-second

temporal blocks. SHAP value calculations utilize a kernel approximation method optimized for electrophysiological data, presenting results as both detailed force plots and summary heatmaps. Integrated Gradients and DeepLIFT visualizations share a unified baseline comparison system that highlights deviations from resting-state activity patterns. Each subsystem includes method-specific normalization to ensure consistent intensity interpretation across explanations, with absolute scaling for SHAP values and relative normalization for attention-based methods. The interface provides comparative analysis tools that align multiple explanation outputs temporally, allowing direct observation of consensus regions across methodologies.

A.1.4. CLINICAL REPORT GENERATION

The automated report generation system synthesizes AI explanations into standardized clinical documentation formats compliant with medical record-keeping requirements. The PDF engine constructs multi-page reports beginning with patient metadata sections that include configurable fields for institutional headers and demographic information. Technical specifications pages detail recording parameters, preprocessing steps, and explanation methodology using terminology adapted from ACNS guidelines. Visualization pages employ a consistent layout with figure numbering, color map legends, and normalized intensity scales across all explanation methods. The system implements a unique clinical correlation section that juxtaposes AI attention patterns with annotated frequency band activities, automatically generating preliminary observations such as "Focal beta-band activation correlates with 82% of model attention in the left temporal region." Physician override capabilities are provided through editable comment fields and digital signature placement areas. Report generation times average 1.2 seconds for comprehensive five-page documents, with quality control checks ensuring DICOM-compatible resolution (300dpi) for all embedded visualizations.

A.1.5. SCREENSHOTS OF THE FRONTEND

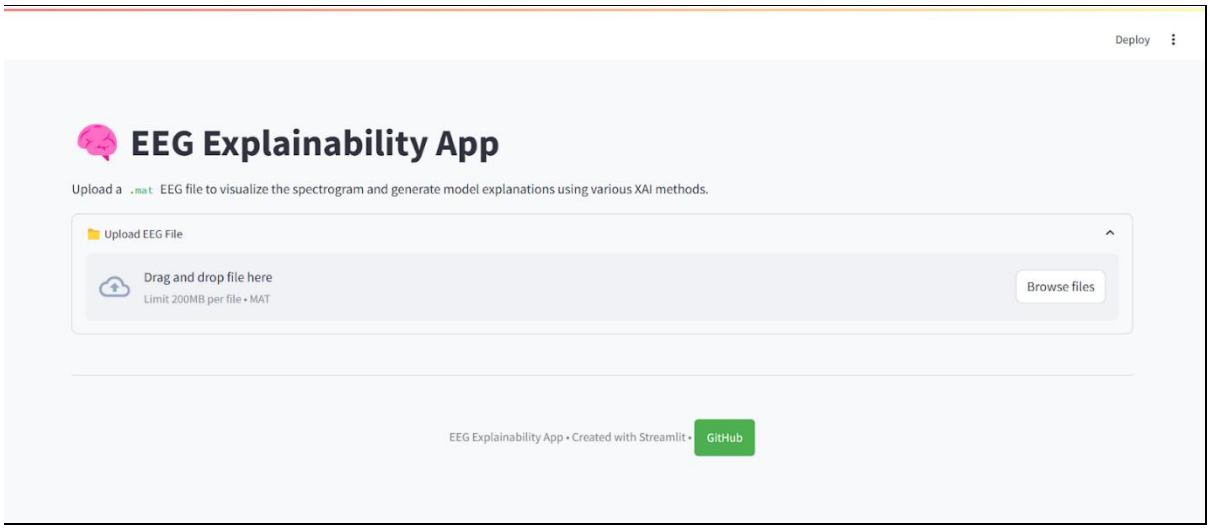

Fig A1.1 Streamlit EEG Explainability App Opening Screen

This is the opening screen that is displayed to the user on launching the application. It prompts the user to upload the .mat file that consists of the EEG recordings of the patient.

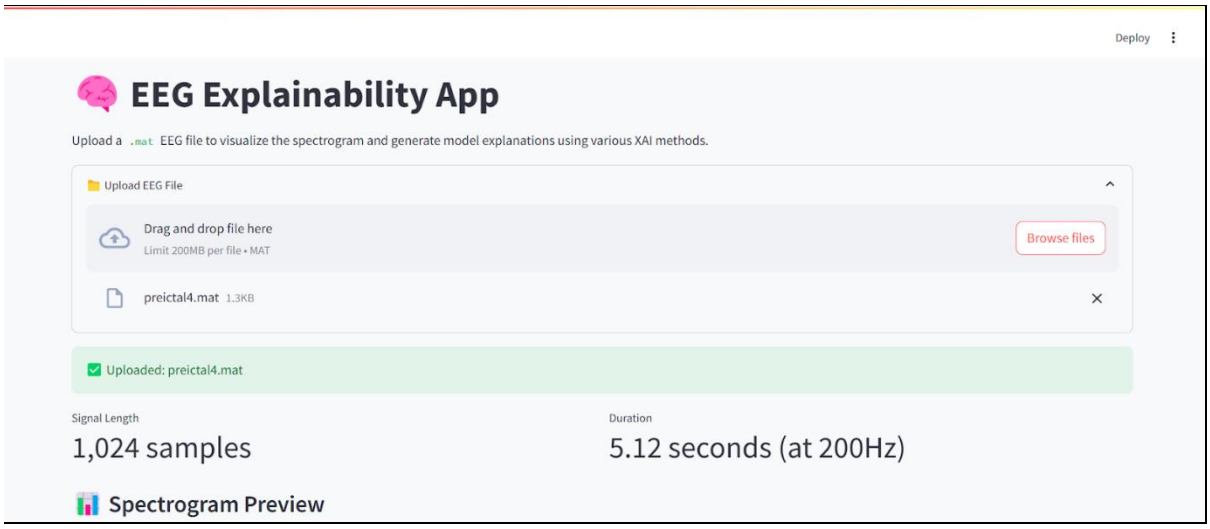

Fig A1.2 File upload interface

On uploading the .mat file, the signal length and the duration of recording are displayed along with the spectrogram for the corresponding signal in the .mat file.

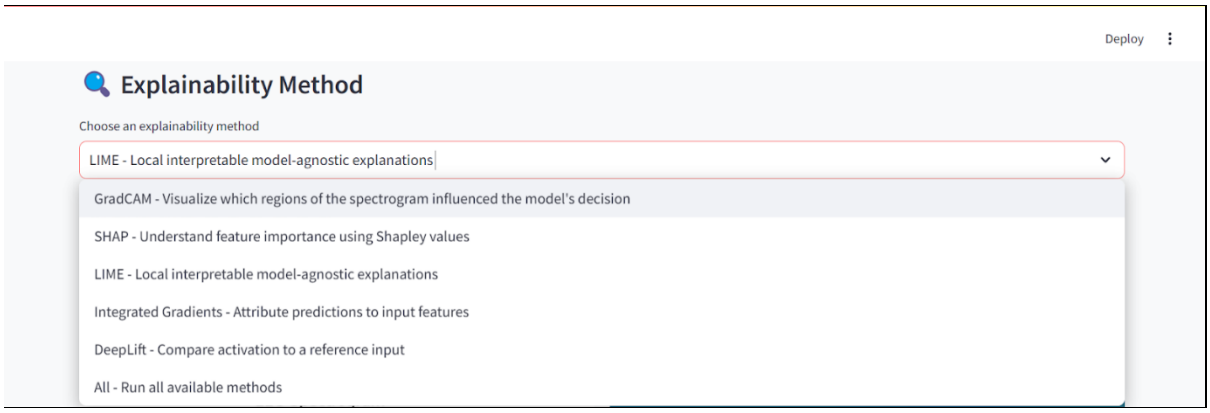

Fig A1.3 Explainability method Dropdown Menu Interface

A dropdown menu is displayed for selecting any of the five XAI methods - SHAP, LIME, Grad-CAM, DeepLIFT for generating the corresponding visual explanation along with the interpretation guide.

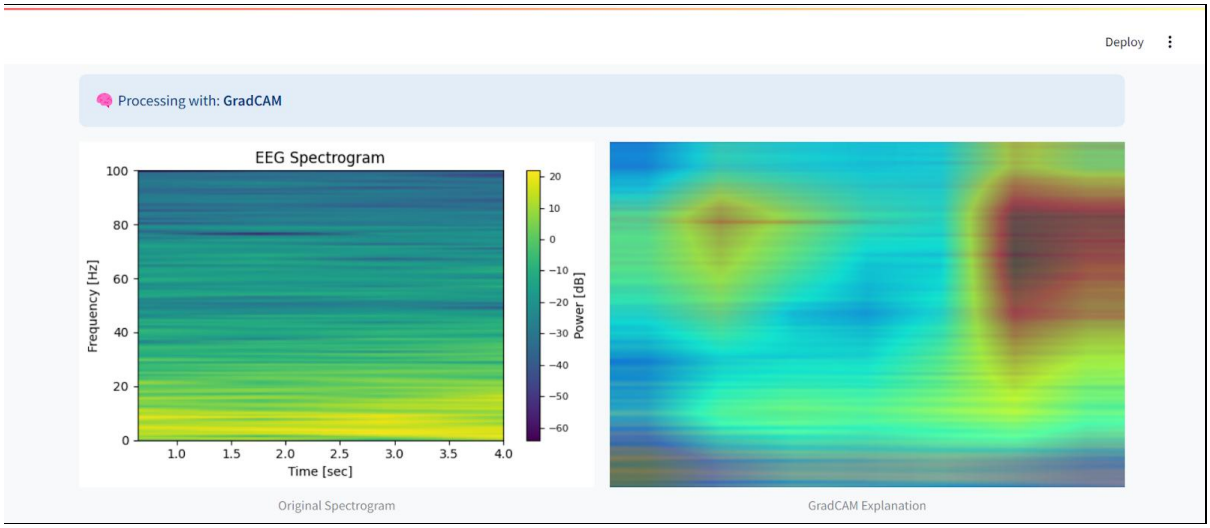

Fig A1.4 XAI Visualization Result

This shows the result that is shown in the interface when a particular technique like Grad-CAM is selected.

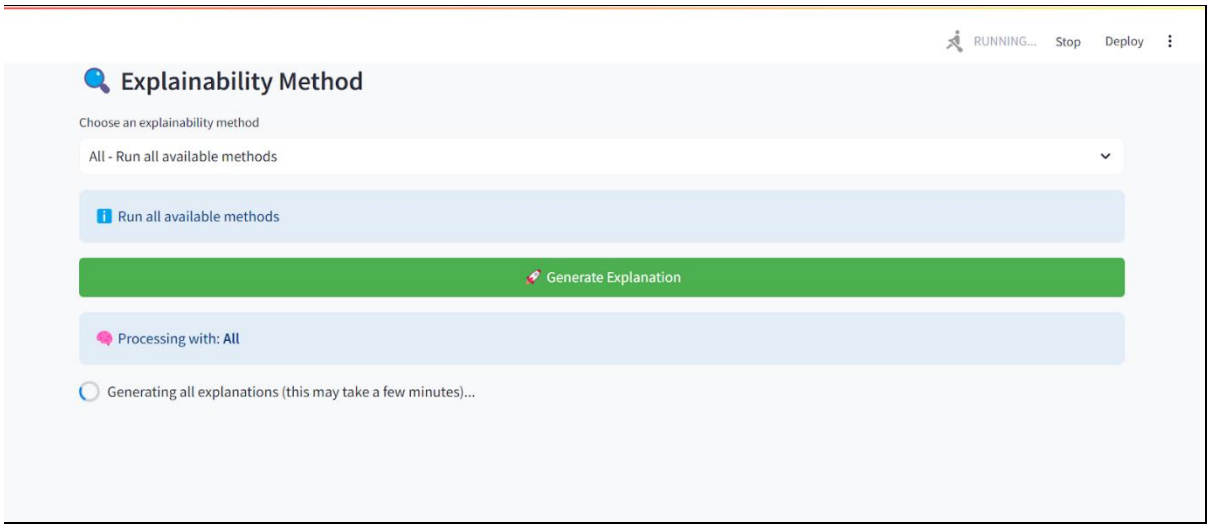

Fig A1.5 Throbber icon

An option “All - Run all available methods” in the XAI dropdown menu generates visual explanations for all five techniques. The user is well informed about what is happening in the interface by suitable prompts and progress using the throbber icon.

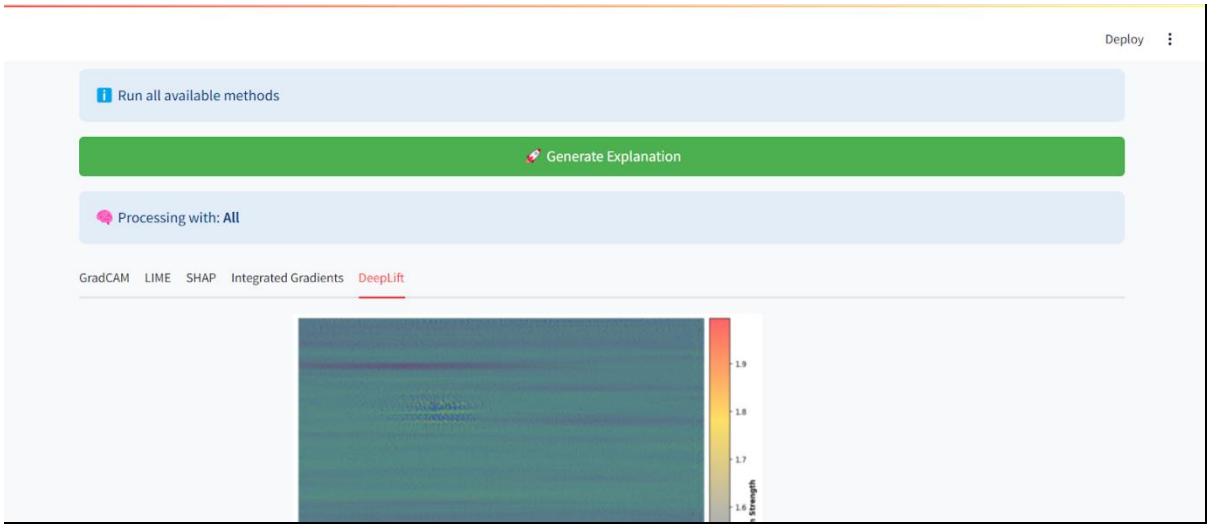

Fig A1.6 Tab-switching interface

The visualizations for the techniques generated by clicking the “All” option can then be viewed by switching tabs as shown in Figure x. The report generated for a sample file is attached in section 7.6.

Fig A1.7 Download Report Button for EEG Analysis Interface

The “Download Full report” at the bottom on clicking generates a clinician friendly report that is downloaded in pdf format. The report generated for a sample file is attached in section A.1.6.

A.1.6. REPORT GENERATED BY THE APPLICATION

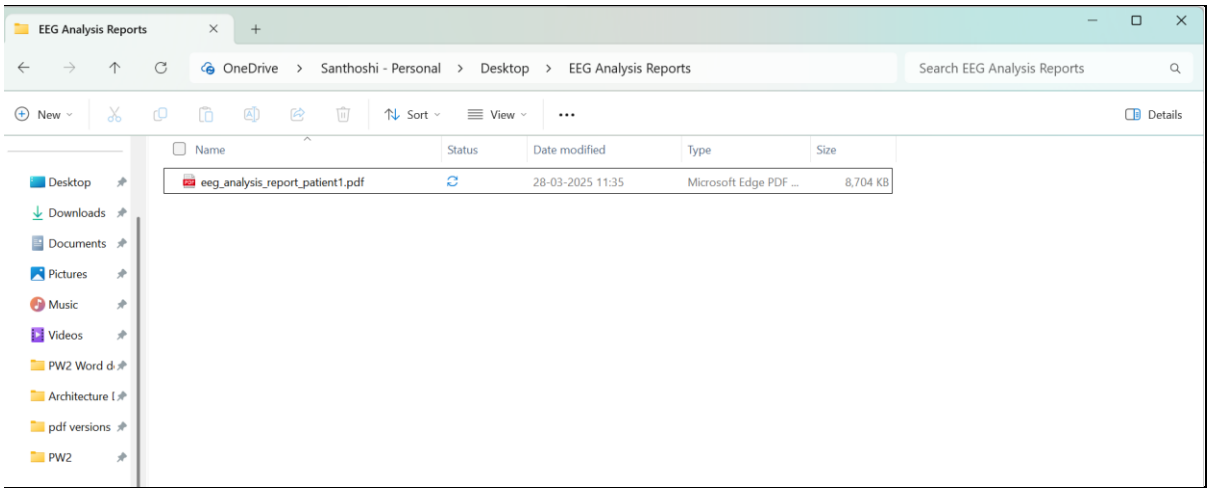

Fig A1.8 Screenshot of Report PDF Downloaded in Local System

EEG Spectral Analysis Report

Report Generated: 2025-04-08 13:25:45

Patient Information

Patient ID: _____

Age/Sex: _____

Clinical Notes:

_____

Recording Parameters

Duration: 5.12 seconds | Sample Rate: 200 Hz

Frequency Bands:

Band	Range	Clinical Significance
Delta	0.5-4 Hz	Deep sleep, pathological states
Theta	4-8 Hz	Drowsiness, meditation
Alpha	8-12 Hz	Relaxed wakefulness
Beta	12-30 Hz	Active thinking, focus
Gamma	30-100 Hz	Cognitive processing

Spectrogram Analysis

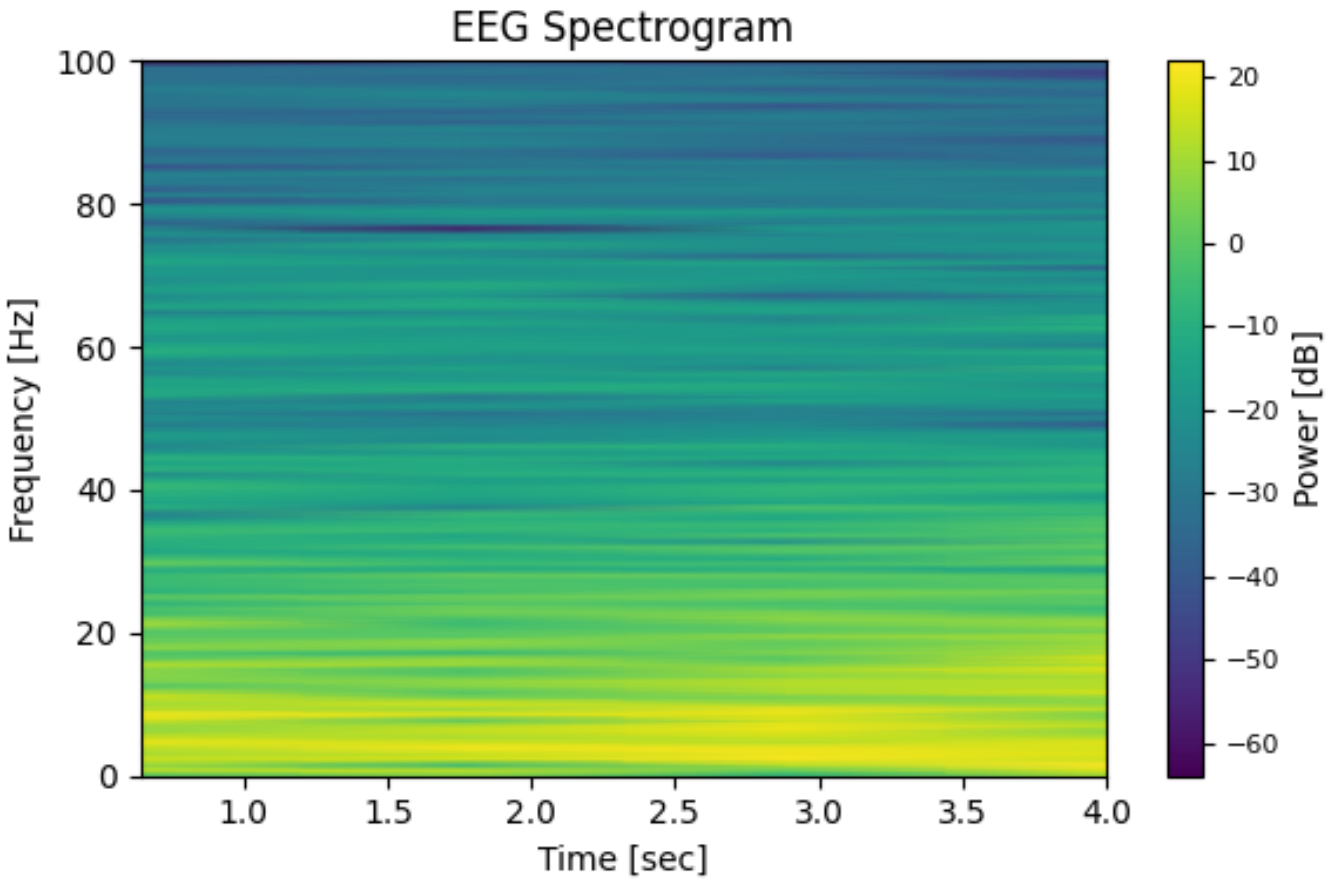

Key Observations:

- X-axis shows time progression (0 to 5.1 seconds)

- Y-axis displays frequency components (0 to 100 Hz)

- Color intensity represents power spectral density (dB)

- Red/Yellow regions indicate higher energy concentrations

GradCAM Analysis

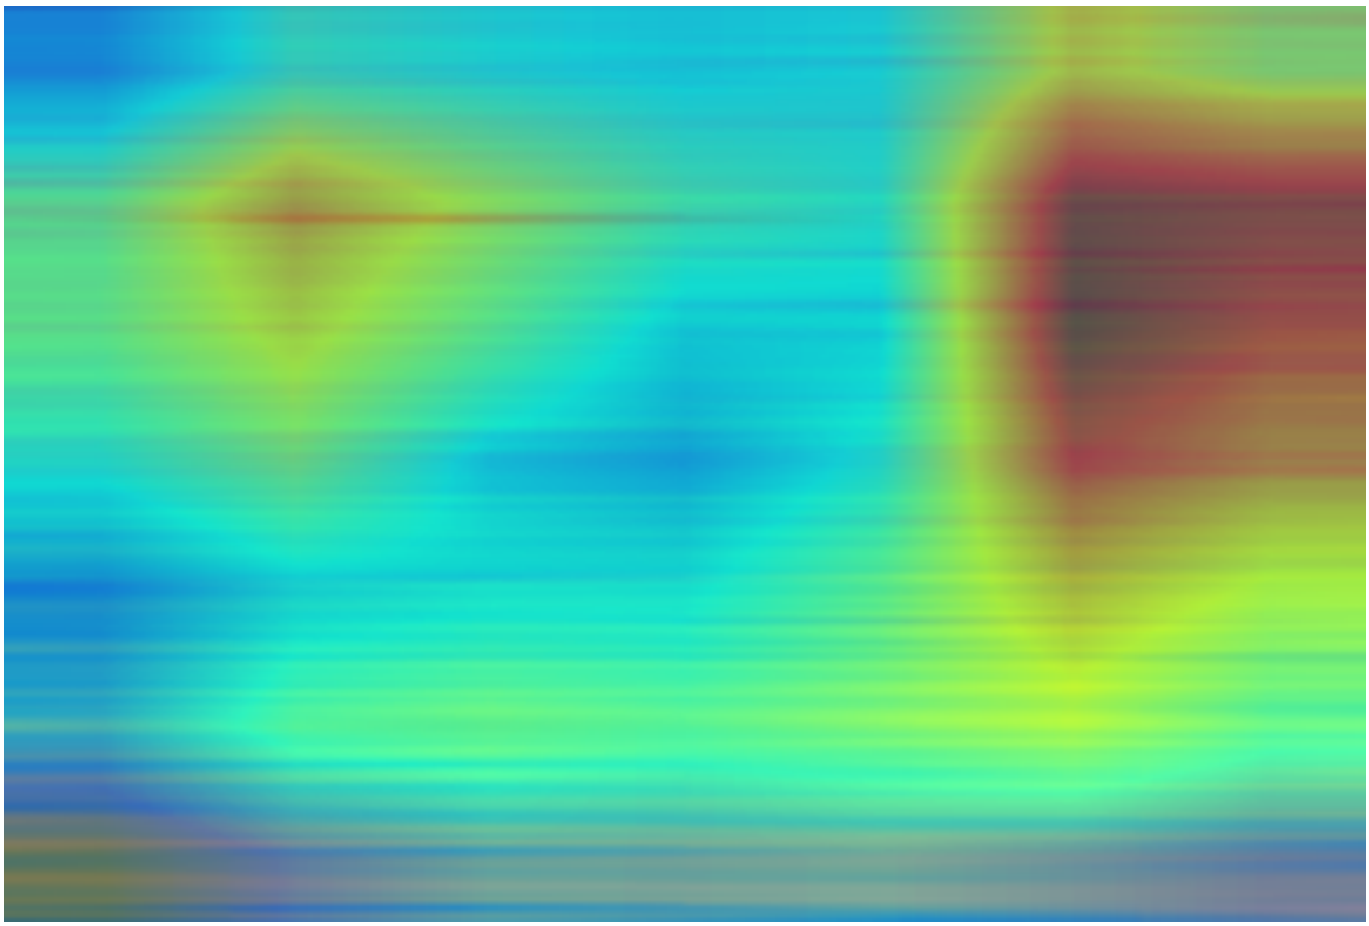

Interpretation:

- Heatmap shows regions that most influenced the model's decision

- Red/Yellow = High importance

- Blue = Low importance

- Works best for spatial localization

LIME Explanation

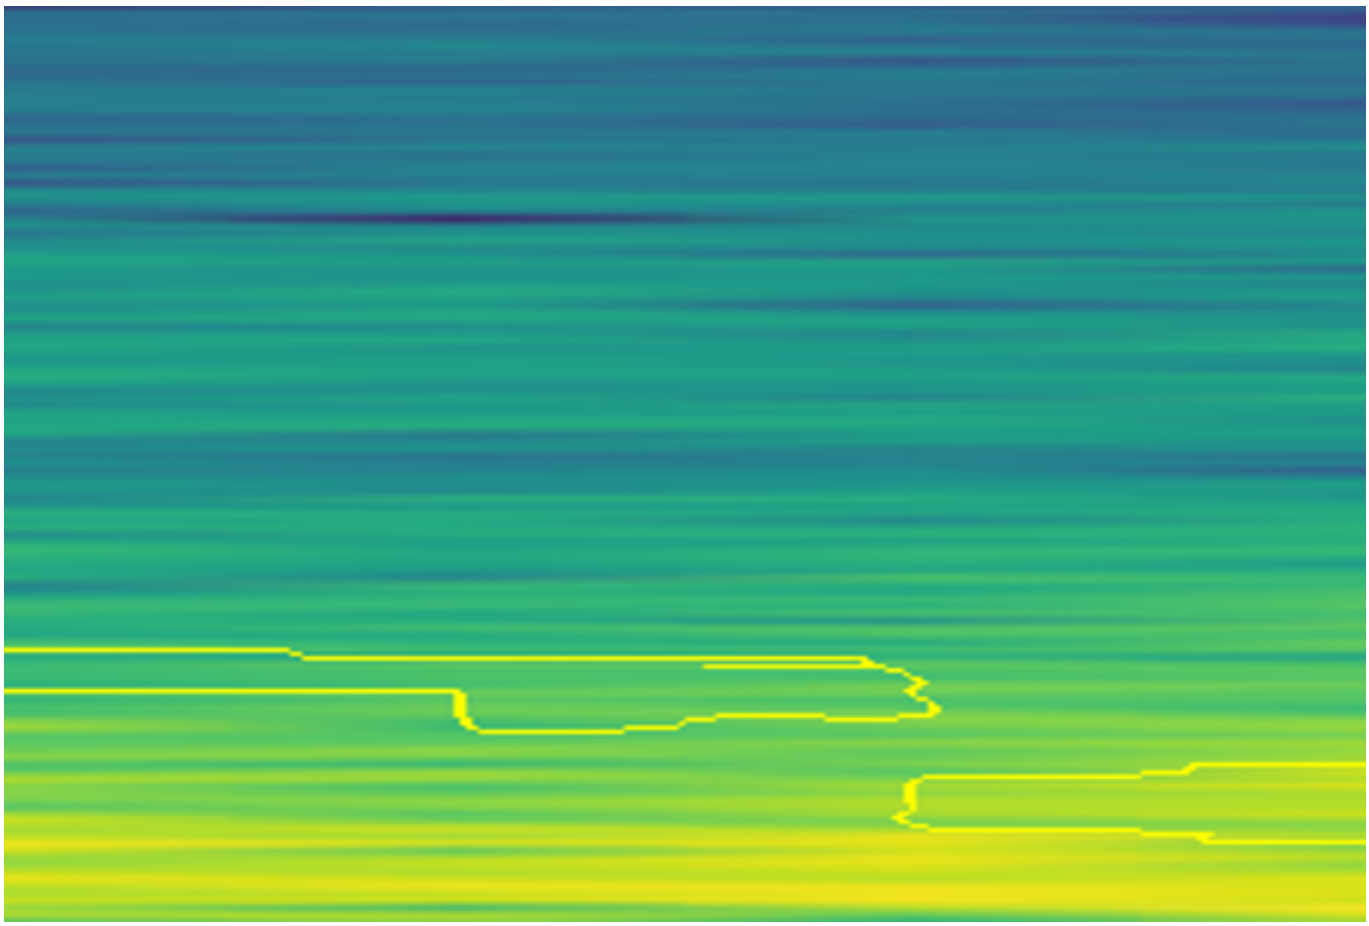

Interpretation:

- Green highlights show influential segments

- Larger areas = More important features

- Identifies key frequency bands

SHAP Value Analysis

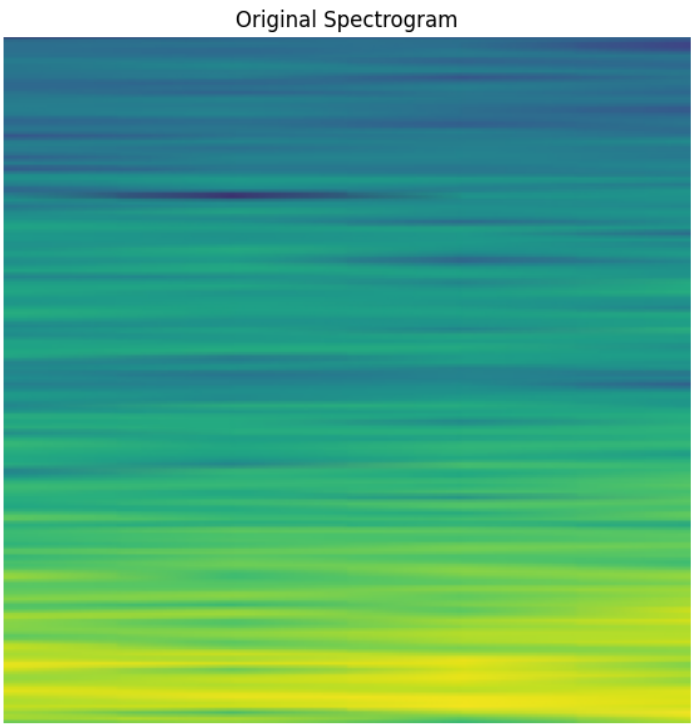

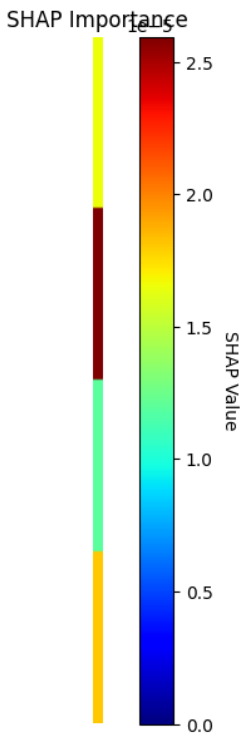

Interpretation:

- Red = Positive impact

- Blue = Negative impact

- Intensity shows influence strength

Integrated Gradients

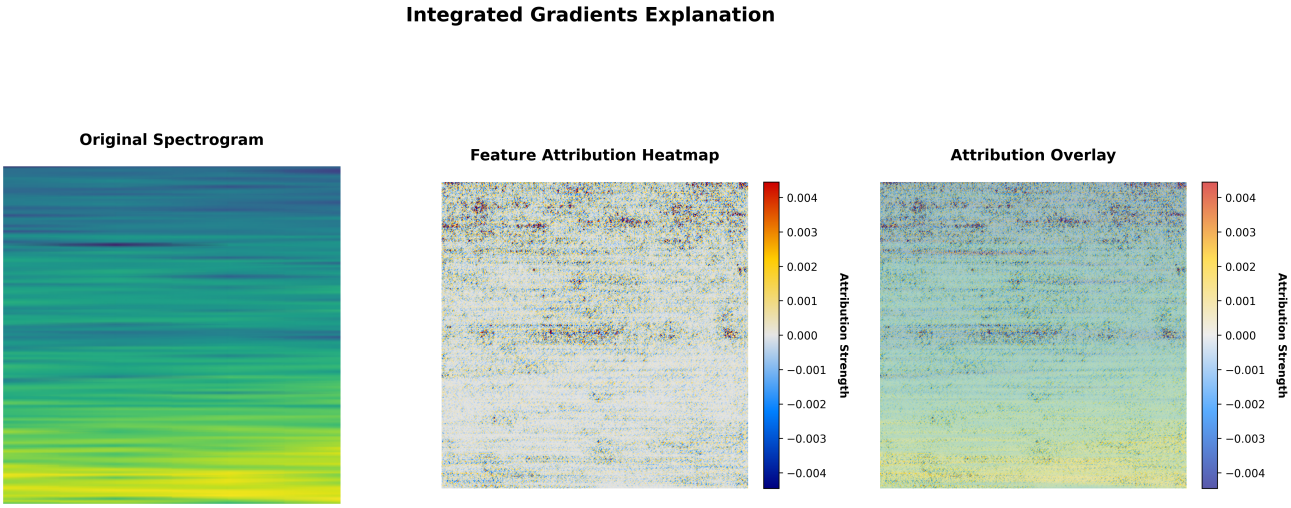

Interpretation:

- Blue = Negative influence

- Red = Positive influence

- Complete attribution method

DeepLIFT Analysis

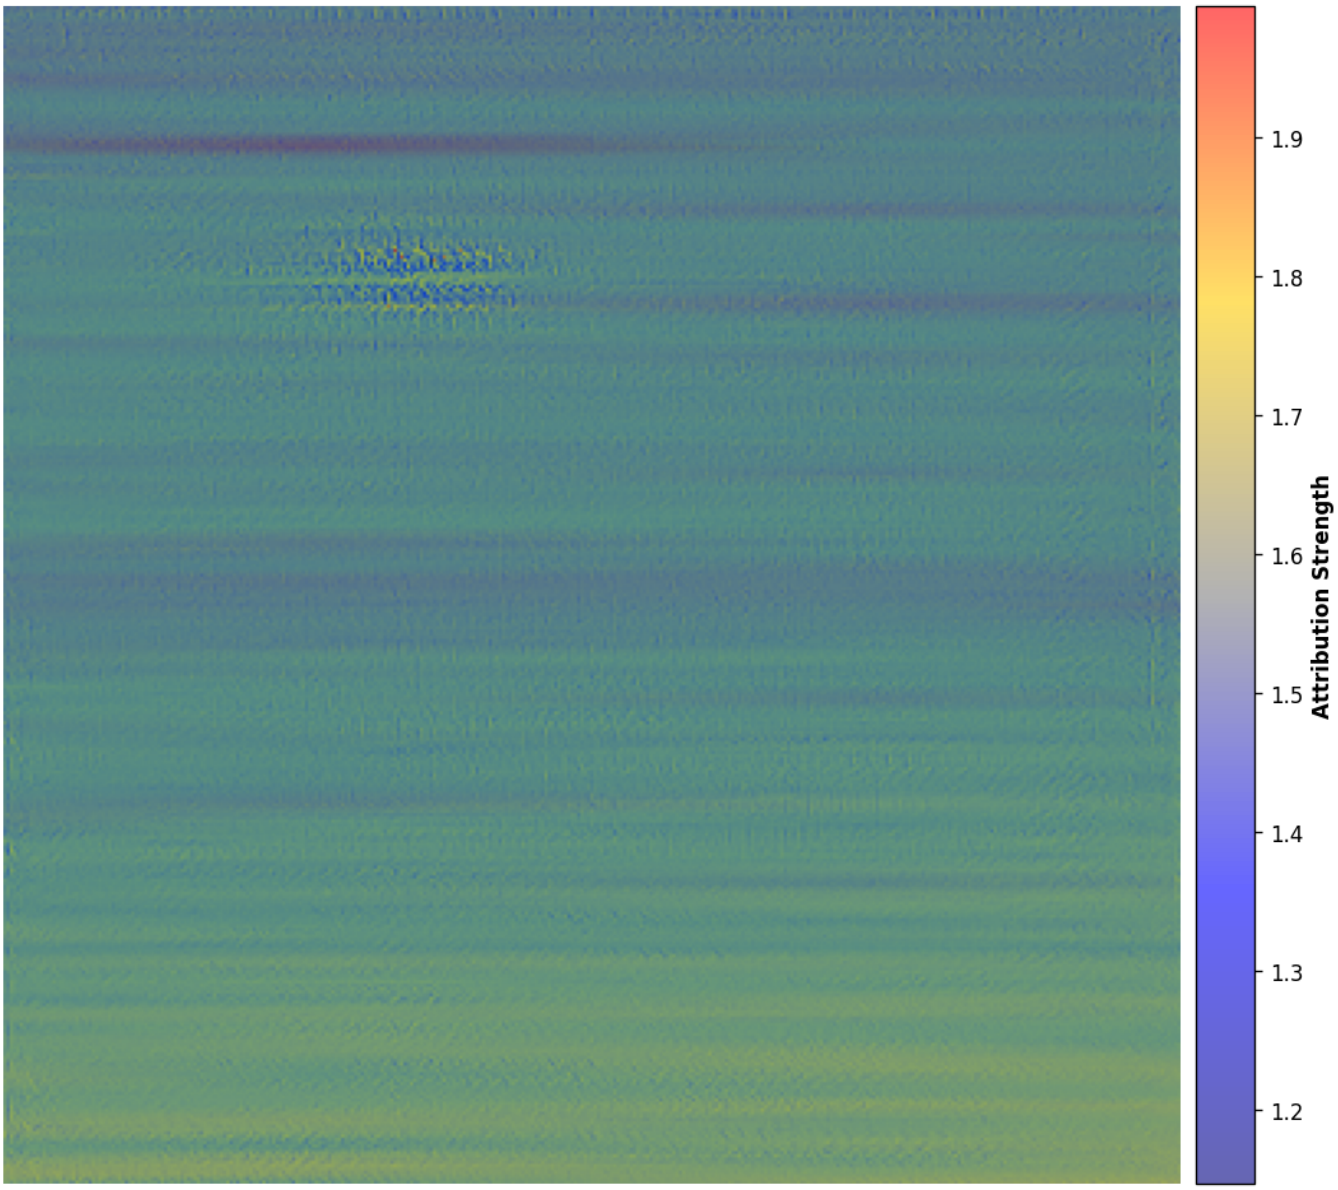

Interpretation:

- Shows differences from reference

- Blue-to-Red = Influence direction

- Handles non-linearities well

Clinical Correlation Notes

1. Correlate findings with patient symptoms and history
2. Review raw EEG traces for verification
3. Consider medication effects
4. Note any technical artifacts

Physician Notes:

_____

_____

This report was automatically generated by EEG XAI System. Clinical correlation required.

PROJECT POSTER

EXPLAINABLE AI FOR EPILEPTIC SEIZURE PREDICTION USING EEG SIGNALS: ENHANCING TRANSPARENCY AND CLINICAL TRUST

THE PROBLEM

- Epileptic seizures are sudden, unpredictable brain events.
- Accurate detection and early prediction are crucial to prevent injury and save lives.
- In critical domains like healthcare, interpretability is not optional — it's essential.

THE CHALLENGE

- Deep learning models (like ResNet) can classify EEG patterns effectively.
- But they work as black boxes — making it hard for doctors to trust or understand their decisions.

FRONTEND INTERFACE

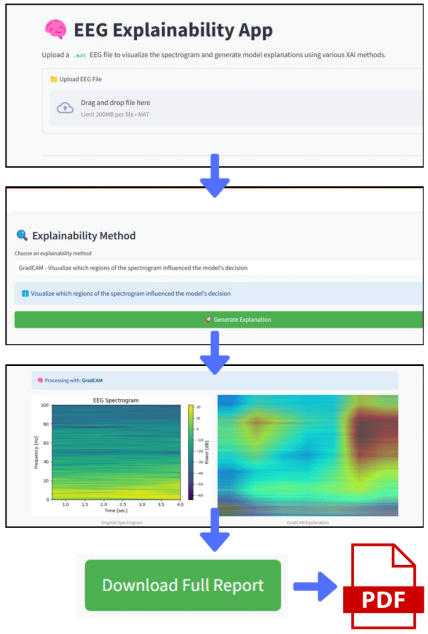

STAGES OF SEIZURE

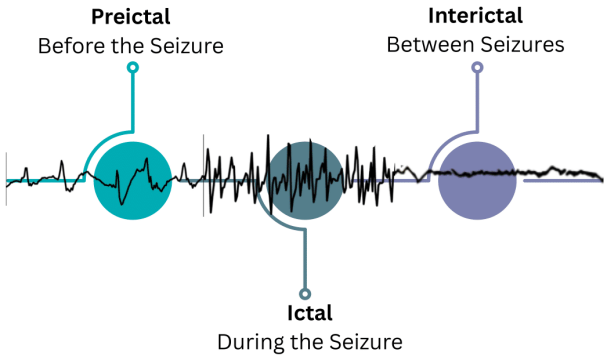

SYSTEM OVERVIEW

1. **EEG Signal Input:** Pre-recorded EEG data categorized into ictal, interictal, and preictal stages. Each signal lasts 5.12 seconds, stored in .mat format.
2. **Preprocessing:** Signals are normalized and optionally transformed into spectrograms to enhance spatial learning. Each signal is labeled for supervised training.
3. **Model Training – ResNet-18:** A modified ResNet-18 convolutional neural network is trained to classify the EEG signals into three classes with high accuracy.
4. **Explainable AI Integration:** Five XAI techniques — SHAP, LIME, Grad-CAM, Integrated Gradients, and DeepLIFT — are applied to visualize and interpret model predictions.
5. **Frontend Deployment:** A Streamlit-based web interface allows users to upload EEG files, get predictions, and interactively view explanation heatmaps and scores.

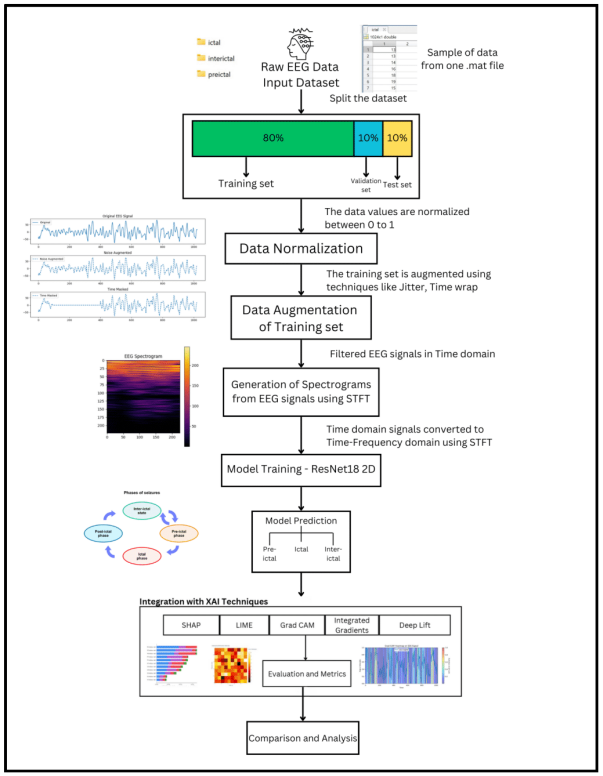

OUR SOLUTION

- Build a system that predicts seizure states (ictal, interictal, preictal)
- Integrate Explainable AI (XAI) techniques to make model decisions visible and understandable.
- Empower clinicians with trustworthy AI that supports better, faster diagnosis.
- Model used: ResNet-18 (pretrained, modified)
- XAI Techniques used:
  - SHAP
  - LIME
  - Grad-CAM
  - Integrated Gradients
  - DeepLIFT
- Interactive frontend web application for report generation for XAI explanations suitable for clinician and doctors.

SYSTEM EVALUATION

FIDELITY
- Higher the better(>0.8)

LIME	0.9205
SHAP	0.7807
Grad-CAM	0.6618

LOCALIZATION
- Optimal range: (0.5 - 0.7)

DeepLIFT	0.7514
Grad-CAM	0.4881
LIME	0.4766

STABILITY
- Higher the better(>0.8)

LIME	0.7894
Integrated G.	0.6250
SHAP	0.5110

Scan the QR for full project repository